

# Week 1: Visual Recognition & Machine Learning

CMPUT 328 – Nilanjan Ray  
University of Alberta


# Part A: Visual Recognition and Deep Learning Context

# What is Visual Recognition?

- In a nutshell:
- Teaching computers to 'see' like humans.
- Covers tasks: classification, detection, segmentation, recognition.
- Applications: autonomous driving, healthcare imaging, surveillance, robotics.

# Visual recognition...

... is teaching computers to see



# Humans see...




# Computers see...

# Teaching computers to “see” like humans




# “see” not just RGB data

- IR (infrared) etc
- ToF camera (Time of Flight)
  - ‘range’ camera gives depth
- Medical
  - ultrasonography
  - MRI
- & more



LIDAR  
(laser radar)



Kinect



# Human vs. Computer Vision


- Humans: rich perception (context, depth, motion, prior knowledge).
- Computers: rely on data (RGB, depth, IR, LIDAR).
- Grand Goal: automated scene understanding comparable to humans.

# Success Stories of Deep Learning

- ImageNet Challenge (2012): AlexNet breakthrough.
- Medical imaging: diabetic retinopathy, pathology, radiology.
- Diffusion models (Stable Diffusion, DALL-E).
- Segment Anything Model (2023, Meta AI).
- Multimodal AI: CLIP, GPT-4V.

# Success stories of deep learning

# Image classification results



Computer vision has surpassed human level performance on this benchmark!

Picture courtesy: <http://cs231n.stanford.edu/index.html>

# Large scale video classification




Figure 4: Predictions on Sports-1M test data. Blue (first row) indicates ground truth label and the bars below show model predictions sorted in decreasing confidence. Green and red distinguish correct and incorrect predictions, respectively.

<http://cs.stanford.edu/people/karpathy/deepvideo/>

# Style Transfer




Figure 1. Example of using the Neural Style Transfer algorithm of Gatys *et al.* to transfer the style of Chinese painting (b) onto The Great Wall photograph (a). The painting that served as style is named “Dwelling in the Fuchun Mountains” by Gongwang Huang.


<https://arxiv.org/abs/1705.04058>

# Deep reinforcement learning




Picture source: <https://deepmind.com/blog/deep-reinforcement-learning/>

# Impressive robotics with deep learning



<https://www.youtube.com/watch?v=B7ZT5oSnRys>

# Diabetic retinopathy using deep learning



<https://www.nature.com/articles/s41467-023-44676-z>

# Photorealistic image generation



From NVIDIA research: <https://arxiv.org/pdf/1710.10196v1.pdf>

# Diffusion-based image and video generation

<https://stability.ai/>

# Segment anything model

<https://segment-anything.com/>

# Deep learning and natural language processing

- ▶ Impressive developments are happening in the NLP space
- ▶ Word embedding
- ▶ Language translation
- ▶ Language modeling
- ▶ ChatGPT!!

<http://ruder.io/nlp-imagenet/>

# What created this revolution?

- Lots and lots of **annotated** data (such as ImageNet)
- Compute power (parallel processing with GPUs)
- Strong neural network architectures
- Good old back-prop algorithm + only a few new tweaks! And
- Open-source software platforms: TensorFlow, PyTorch,...

# Challenges of Deep Learning

- Requires massive labeled datasets.
- Bias in training data → unfair outcomes.
- Adversarial vulnerability (images misclassified with tiny perturbations).
- Poor interpretability → 'black box' issue.
- High compute and energy costs.

WE DON'T HAVE AUTOMATED SCENE UNDERSTANDING YET

# Today at NYT

- Gary Marcus, The Fever Dream of Imminent ‘Superintelligence’ Is Finally Breaking:  
[https://www.nytimes.com/2025/09/03/opinion/ai-gpt5-rethinking.html?unlocked\\_article\\_code=1.jE8.MKWI.YCS0TSrbPReK&smid=url-share](https://www.nytimes.com/2025/09/03/opinion/ai-gpt5-rethinking.html?unlocked_article_code=1.jE8.MKWI.YCS0TSrbPReK&smid=url-share)

# Video: Yann LeCun on AI (Historical Context)

- Watch here:  
[https://www.youtube.com/watch?v=4\\_\\_gg83s  
\\_Do](https://www.youtube.com/watch?v=4__gg83s_Do)
- Pioneer of convolutional neural networks (CNNs).
- Perspective on AI and AGI.

# Video: MIT Economist on AI and Jobs (Societal/Economic Impact)

- Watch here:  
<https://www.youtube.com/watch?v=-zF1mkBpyf4>
- Economic disruption/hype from AI adoption

# Video: Andrej Karpathy on Deep Learning (Modern Breakthroughs)

- Watch here:  
<https://www.youtube.com/watch?v=LCEmiRjPEtQ>
- Former Tesla/Stanford researcher.
- Discusses modern breakthroughs and future directions.

# Part B: Machine Learning Foundations

# What is Machine Learning?

- Arthur Samuel (1959): 'Field of study that gives computers the ability to learn without explicit programming.'
- Tom Mitchell (1997): Learning improves performance on a task with experience.
- Core idea: Learn functions/mapping using data.

# Why do we need ML?

- Hard-coded rules fail in complex domains.
  - Quite relevant for visual recognition: You simply cannot describe a cat or dog or an object with full-proof, explicit description.
- ML adapts to new data automatically.
- Applications: spam filters, recommendation systems, speech recognition, visual recognition,...


# Types of ML

- Supervised: learns from labeled data.
- Unsupervised: finds structure in unlabeled data.
- Semi-supervised: mix of labeled + unlabeled.
- Reinforcement learning: learns by trial and error with rewards.
- Batch vs Online learning.


# Supervised Learning

## Classification:

The spam filter is a good example of this: it is trained with many example emails along with their *class* (spam or ham), and it must learn how to classify new emails.




**Regression:** Another typical task is to predict a *target* numeric value, such as the price of a car, given a set of *features* (mileage, age, brand, etc.) called *predictors*. To train the system, you need to give it many examples of cars, including both their predictors and their labels (i.e., their prices).




# Instance-based supervised learning

- Remember all training examples
- When a test email comes, compare it with its “neighbors” from the training examples and classify accordingly
- Requires a measure of similarity
- Example: k-nearest neighbor (knn) method



# Model-based supervised machine learning



- From all the training examples, build a model for the learner
- When a test example comes, apply the model
- Don't need to remember all training examples, after training
- Examples: neural net, support vector machine, linear regression, etc.



# Linear Regression

- Model:  $y = \theta_0 + \theta_1x$ .
- Fits a straight line to data.
- Used for trend prediction (housing prices, stock prices).


# Linear-model based supervised learning



Linear model:  $\text{life\_satisfaction} = \theta_0 + \theta_1 \times \text{GDP\_per\_capita}$


Parameters of the model:  $\theta_0, \theta_1$

# Linear-model based supervised learning



# Logistic Regression

- Used for binary classification tasks.
- Outputs probability values using sigmoid function.
- Example: predicting if an image is a cat or not.




# Training & Evaluation Metrics

- Accuracy = correct predictions / total predictions.
- Precision, Recall, F1-score (important in imbalanced data).
- Confusion matrix for visualization.
- ROC curve and AUC as performance measures.

# Overfitting vs. Underfitting

- Overfitting: model memorizes training data → poor generalization.
- Underfitting: model too simple → fails to capture patterns.
- Goal: balance bias and variance.




# Regularization

- Penalty on large parameter values (L1, L2 regularization).
- Helps prevent overfitting.
- Encourages simpler models.

# Challenge 1


- Insufficiency of annotated training data






The importance of data versus algorithms: by Peter Norvig

# Challenge 2

- Non-representative training data



# Challenge 3




What is regularization?

# Part C: Instance-based Learning with k-NN

# ML as Function Mapping ( $x \rightarrow y$ )

- Machine learning learns a function  $f(x) \approx y$ .
- $x$  = features (inputs),  $y$  = annotations (outputs).
- Goal: generalize well to unseen test data.

# Supervised machine learning: the tabular view




# k-NN Algorithm (intuition, steps)

- Store all training examples.
- For a new point, compute distance to all training points.
- Pick the  $k$  (could be 1, 2, 3, etc.) closest neighbors.
- Predict class by majority vote (classification) or average (regression).

# Toy Example (2D features, distance calculations)

- Example: classify a point (1,1) with k=3.
- Compute distances to training points.
- Select nearest neighbors → assign majority label.



# K-nn: A toy numerical example...

Training data,  $m = 5$

$x_1$	$x_2$	$y$
2	-1	0
3	2	1
0	4	0
-2	5	0
2	0	1
1	1	?

Test data point

For this problem, note that the feature vector dimension,  $d=2$

Let's assume  $k = 3$

To find out  $k=3$  nearest neighbors, compute distances:

$$D_1([1, 1], [2, -1]) = |1-2| + |1+1| = 3$$

$$D_2([1, 1], [3, 2]) = |1-3| + |1-2| = 3$$

$$D_3([1, 1], [0, 4]) = |1-0| + |1-4| = 4$$

$$D_4([1, 1], [-2, 5]) = |1+2| + |1-5| = 7$$

$$D_5([1, 1], [2, 0]) = |1-2| + |1-0| = 2$$

So,  $k=3$  nearest neighbors are  
 $N_3([1, 1]) = \{1, 2, 5\}$

Prediction for test data point:

$$f([1, 1]) = \text{Ave}([y(1), y(2), y(5)])$$

$$= \text{Ave}([0, 1, 1]) = 1$$

Here, we computed “Ave” by taking mode.


# Choosing k (validation sets, bias-variance tradeoff)

- Small  $k \rightarrow$  sensitive to noise (low bias, high variance).
- Large  $k \rightarrow$  smoother decision boundaries (high bias, low variance).
- Use cross-validation to choose optimal  $k$ .

# Choosing $k$ in practice

- Divide training data into two sets: training (90%) and validation (10%).
- For each  $k$  in a range, find out k-nn prediction accuracy on the validation set.
- Choose the  $k$  that has yielded the highest accuracy on the validation set.

# MNIST digit image classification



Small 28 pixels-by-28 pixels images of handwritten digits

The visual recognition problem definition:  
**to recognize the digit from an image**

We can attempt to solve this using k-nn.

Feature dimension,  $d = 28 * 28 = 784$

Training data

Test data

	$x_1$	$x_2$	...	$x_{784}$	Digit
1	0.1	0.3	...	0.0	0
2	0.2	0.1	...	0.5	1
...	...	...	...	...	...
...	...	...	...	...	...
5	0.0	0.98	...	0.8	9
6	0.5	0.25	...	0.36	?
7	0.1	0.95	...	0.1	?

A recommended resource: <https://cs231n.github.io/classification/>

# Efficient Computation of k-NN

- Brute force distance computation is expensive.
- Vectorization with NumPy/PyTorch speeds up calculations.
- KD-trees, ball trees, and approximate nearest neighbor methods scale to large datasets.

# Efficient computation of K-nn: Vectorization

- For loops are slow. How do we avoid for loops in K-nn computation?
- Suppose  $X^{tr}$  is the training data matrix of shape N-by-d and  $X^{tst}$  is the test data matrix of shape M-by-d, d is the dimension of feature vector N training data points and M test data points.
- We want to compute the M-by-N distance matrix D:

$$D_{ij} = \sum_{k=1}^d (X_{ik}^{tr} - X_{jk}^{tst})^2 = \sum_{k=1}^d (X_{ik}^{tr})^2 + \sum_{k=1}^d (X_{jk}^{tst})^2 - 2 \sum_{k=1}^d X_{ik}^{tr} X_{jk}^{tst}$$

- Using Python's broadcast feature, we can compute D as:

```
D  
= sum(Xtr ** 2, dim = 1, keepdim = True)  
+ transpose(sum(Xtst ** 2, dim = 1, keepdim = True) - 2  
* matmul(Xtr, transpose(Xtst)))
```

# MNIST Digit Classification with k-NN

- 28x28 grayscale images (784 features).
- Train k-NN on digits 0–9.
- Works well for small datasets but slow for large scale.
- Let's look at the notebook

# Strengths & Limitations of k-NN

- Strengths: simple, interpretable, no training time.
- Limitations: high memory usage, slow prediction, **poor in high dimensions.**

# Transition to Neural Networks

- k-NN shows limits for complex tasks.
- Neural networks learn abstract features directly.
- Next week: perceptron, MLPs, backpropagation.

# Wrap-Up (10 min)

# Recap: Visual Recognition, ML Foundations, k-NN

- We introduced visual recognition and its challenges.
- Explored ML foundations.
- Learned k-NN as first ML algorithm for vision tasks.

# Looking Ahead: Neural Networks and CNNs (Week 2)

- Next: neural networks → perceptron, MLPs.
- Then: convolutional neural networks (CNNs) for images.
- Transformers in vision coming later in the course.