

Linear, Multiple, and Logistic Regressions

CMPUT 328

Nilanjan Ray

Computing Science, University of Alberta, Canada

Linear regression with PyTorch

- We will start with a linear regression “model”
- Next, we need to understand “loss” function for regression task
- Next, we will estimate the model by minimizing the loss function
- We will use PyTorch

Supervised machine learning: the tabular view

| Independent variables (aka feature vector) | | | | Prediction / dependent variable |
|---|-------|-------|-------|---------------------------------------|
| x_1 | x_2 | x_3 | x_4 | y |
| 1.2 | -3.9 | 4.0 | 0 | 1.6 |
| 2.1 | 2.4 | -0.7 | -0.2 | 1.2 |
| ... | ... | ... | ... | ... |
| ... | ... | ... | ... | ... |
| 3.2 | ... | ... | 1.9 | 0.3 |
| 1.4 | ... | ... | 1.5 | ? |
| 3.1 | ... | ... | 2.1 | ? |



ML learns to map **x to y**

Training data:
complete table

In other words, ML learns
a function, f so that
 $y = f(x)$

Test data:
incomplete table

The function f is called **prediction function**

Linear prediction: formal setup

Linear prediction function:

$$y^p = \mathbf{x}\boldsymbol{\theta} + b$$

or,

$$y^p = \sum_{j=1}^m \theta_j x_j + b$$

vector equation form

scalar equation form

A training set consists of (\mathbf{x}, y) pairs: $\{(\mathbf{x}_i, y_i)\}_{i=1}^n$

Linear prediction on the training data point i :

$$y_i^p = \mathbf{x}_i \boldsymbol{\theta} + b$$

or,

$$y_i^p = \sum_{j=1}^m \theta_j x_{i,j} + b$$

Loss or cost function (on training data):

$$L = \frac{1}{2} \sum_{i=1}^n (y_i^p - y_i)^2 = \frac{1}{2} \sum_{i=1}^n (\mathbf{x}_i \boldsymbol{\theta} + b - y_i)^2$$

Learning the linear model: find out $\boldsymbol{\theta}$ and b to minimize loss function

Linear regression: A toy example

Let's take a toy example:

| x_1 | x_2 | y |
|-------|-------|-----|
| 1 | 2 | -1 |
| 3 | -4 | 7 |
| 6 | 2 | 3 |
| -3 | 5 | -4 |
| 7 | -3 | 5 |
| 4 | 3 | ? |

This equation $y_i^p = \sum_{j=1}^m \theta_j x_{i,j} + b$

can be written for the toy training set as

We also have ground truth responses:

$$y_1 = -1, y_2 = 7, y_3 = 3, y_4 = -4, y_5 = 5$$

$$\begin{aligned}y_1^p &= \theta_1(1) + \theta_2(2) + b \\y_2^p &= \theta_1(3) + \theta_2(-4) + b \\y_3^p &= \theta_1(6) + \theta_2(2) + b \\y_4^p &= \theta_1(-3) + \theta_2(5) + b \\y_5^p &= \theta_1(7) + \theta_2(-3) + b\end{aligned}$$

So, the loss is $L = \frac{1}{2} \sum_{i=1}^n (y_i^p - y_i)^2 = \frac{1}{2} [(y_1^p + 1)^2 + (y_2^p - 7)^2 + (y_3^p - 3)^2 + (y_4^p + 4)^2 + (y_5^p - 5)^2]$

$$= \frac{1}{2} [(\theta_1 + 2\theta_2 + b + 1)^2 + (3\theta_1 - 4\theta_2 + b - 7)^2 + (6\theta_1 + 2\theta_2 + b - 3)^2 + (-3\theta_1 + 5\theta_2 + b + 4)^2 + (7\theta_1 - 3\theta_2 + b - 5)^2]$$

Learning a linear model

For the convenience of math, let us change our linear model a bit:

$$y_i^p = \sum_{j=1}^m \theta_j (x_{i,j} - \bar{x}_j) + b \quad \text{where} \quad \bar{x}_j = \frac{1}{n} \sum_{i=1}^n x_{i,j}$$

Subtracting independent variable mean is called “data centering”

And a slightly modified loss function:

$$L = \frac{1}{2} \sum_{i=1}^n (y_i^p - y_i)^2 + \frac{\gamma}{2} \sum_{j=1}^m \theta_j^2$$

γ is a hyper parameter



Data fidelity Regularization

Why do we need regularization?

Minimization of linear regression loss function

Regularized loss function: $L = \frac{1}{2} \sum_{i=1}^n (y_i^p - y_i)^2 + \frac{\gamma}{2} \sum_{j=1}^m \theta_j^2$

Taking partial derivative using chain rule: $\frac{\partial L}{\partial b} = \sum_{i=1}^n (y_i^p - y_i) \frac{\partial y_i^p}{\partial b} = \sum_{i=1}^n (y_i^p - y_i)$ because, $\frac{\partial y_i^p}{\partial b} = 1$

(because derivative of $(y_i - \hat{y}_i)^2$ wrt \hat{y}_i is $2(\hat{y}_i - y_i)$, and the $\frac{1}{2}$ cancels the 2).

Using $\bar{x}_j = \frac{1}{n} \sum_{i=1}^n x_{i,j}$ and $y_i^p = \sum_{j=1}^m \theta_j (x_{i,j} - \bar{x}_j) + b$ we get: $\frac{\partial L}{\partial b} = nb - \sum_{i=1}^n y_i$

At the minimum of L , $\frac{\partial L}{\partial b} = 0$ So, $b = \frac{1}{n} \sum_{i=1}^n y_i = \bar{y}$

Linear regression: A toy example...continued

Let's take a toy example:

| x_1 | x_2 | y |
|-------|-------|----|
| 1 | 2 | -1 |
| 3 | -4 | 7 |
| 6 | 2 | 3 |
| -3 | 5 | -4 |
| 7 | -3 | 5 |
| 4 | 3 | ? |

$$b = \frac{1}{n} \sum_{i=1}^n y_i = \bar{y} = \frac{1}{5}(-1 + 7 + 3 - 4 + 5) = 2$$

$$\bar{x}_1 = \frac{1}{n} \sum_{i=1}^n x_{i,1} = \frac{1}{5}(1 + 3 + 6 - 3 + 7) = 2.8$$

$$\bar{x}_2 = \frac{1}{n} \sum_{i=1}^n x_{i,2} = \frac{1}{5}(2 - 4 + 2 + 5 - 3) = 0.4$$

So, using centered data, the prediction equation becomes:

$$y_i^p = \sum_{j=1}^m \theta_j (x_{i,j} - \bar{x}_j) + b = \theta_1(x_{i,1} - 2.8) + \theta_2(x_{i,2} - 0.4) + 2$$

So, the loss is

$$\begin{aligned}
L &= \frac{1}{2} \sum_{i=1}^n (y_i^p - y_i)^2 = \frac{1}{2} [(y_1^p + 1)^2 + (y_2^p - 7)^2 + (y_3^p - 3)^2 + (y_4^p + 4)^2 + (y_5^p - 5)^2] \\
&= \frac{1}{2} [(\theta_1(1 - 2.8) + \theta_2(2 - 0.4) + 2 + 1)^2 + (\theta_1(3 - 2.8) + \theta_2(-4 - 0.4) + 2 - 7)^2 \\
&\quad + (\theta_1(6 - 2.8) + \theta_2(2 - 0.4) + 2 - 3)^2 + (\theta_1(-3 - 2.8) + \theta_2(5 - 0.4) + 2 + 4)^2 + (\theta_1(7 - 2.8) + \theta_2(-3 - 0.4) + 2 - 5)^2]
\end{aligned}$$

Minimization of linear regression loss function...

Regularized loss function:

$$L = \frac{1}{2} \sum_{i=1}^n (y_i^p - y_i)^2 + \frac{\gamma}{2} \sum_{j=1}^m \theta_j^2$$

Taking partial derivative of L using chain rule:

$$\frac{\partial L}{\partial \theta_j} = \sum_{i=1}^n (y_i^p - y_i) \frac{\partial y_i^p}{\partial \theta_j} + \gamma \theta_j$$

Using

$$y_i^p = \sum_{k=1}^m \theta_k (x_{i,k} - \bar{x}_k) + b,$$

$$b = \bar{y}$$

and

$$\frac{\partial y_i^p}{\partial \theta_j} = x_{i,j} - \bar{x}_j$$

We get: $\frac{\partial L}{\partial \theta_j} = \sum_{i=1}^n \left(\sum_{k=1}^m \theta_k (x_{i,k} - \bar{x}_k) + \bar{y} - y_i \right) (x_{i,j} - \bar{x}_j) + \gamma \theta_j$

Linear regression: A toy example...continued

Let's take a toy example:

| x_1 | x_2 | y |
|-------|-------|----|
| 1 | 2 | -1 |
| 3 | -4 | 7 |
| 6 | 2 | 3 |
| -3 | 5 | -4 |
| 7 | -3 | 5 |
| 4 | 3 | ? |

Note: For this toy problem, I assumed $\gamma = 0$ for convenience

$$\frac{\partial L}{\partial \theta_j} = \sum_{i=1}^n \left(\sum_{k=1}^m \theta_k (x_{i,k} - \bar{x}_k) + \bar{y} - y_i \right) (x_{i,j} - \bar{x}_j) + \gamma \theta_j$$

$$\begin{aligned}\frac{\partial L}{\partial \theta_1} &= (\theta_1(1 - 2.8) + \theta_2(2 - 0.4) + 2 + 1)(1 - 2.8) \\ &+ (\theta_1(3 - 2.8) + \theta_2(-4 - 0.4) + 2 - 7)(3 - 2.8) \\ &+ (\theta_1(6 - 2.8) + \theta_2(2 - 0.4) + 2 - 3)(6 - 2.8) \\ &+ (\theta_1(-3 - 2.8) + \theta_2(5 - 0.4) + 2 + 4)(-3 - 2.8) \\ &+ (\theta_1(7 - 2.8) + \theta_2(-3 - 0.4) + 2 - 5)(7 - 2.8) \\ &= (64.8)\theta_1 - (39.6)\theta_2 - 57\end{aligned}$$

$$\begin{aligned}\frac{\partial L}{\partial \theta_2} &= (\theta_1(1 - 2.8) + \theta_2(2 - 0.4) + 2 + 1)(2 - 0.4) \\ &+ (\theta_1(3 - 2.8) + \theta_2(-4 - 0.4) + 2 - 7)(-4 - 0.4) \\ &+ (\theta_1(6 - 2.8) + \theta_2(2 - 0.4) + 2 - 3)(2 - 0.4) \\ &+ (\theta_1(-3 - 2.8) + \theta_2(5 - 0.4) + 2 + 4)(5 - 0.4) \\ &+ (\theta_1(7 - 2.8) + \theta_2(-3 - 0.4) + 2 - 5)(-3 - 0.4) \\ &= (-39.6)\theta_1 + (57.2)\theta_2 + 63\end{aligned}$$

Minimization of linear regression loss function...

$$\frac{\partial L}{\partial \theta_j} = \sum_{i=1}^n \left(\sum_{k=1}^m \theta_k (x_{i,k} - \bar{x}_k) + \bar{y} - y_i \right) (x_{i,j} - \bar{x}_j) + \gamma \theta_j$$

↓
simplification

Gradient of L with
resp. to $\boldsymbol{\theta}$:

$$\nabla_{\boldsymbol{\theta}} L = \left[\sum_{i=1}^n (\mathbf{x}_i - \bar{\mathbf{x}})^T (\mathbf{x}_i - \bar{\mathbf{x}}) \right] \boldsymbol{\theta} + \gamma \boldsymbol{\theta} - \sum_{i=1}^n (y_i - \bar{y}) (\mathbf{x}_i - \bar{\mathbf{x}})^T$$

where $\mathbf{x}_i = [x_{i,1} \dots x_{i,m}]$, $\bar{\mathbf{x}} = [\bar{x}_1 \dots \bar{x}_m]$ and $\boldsymbol{\theta} = [\theta_1 \dots \theta_m]^T$

More simplified form: $\nabla_{\boldsymbol{\theta}} L = (X^T X + \gamma I) \boldsymbol{\theta} - X^T \mathbf{y}$

where matrix X is defined as: $X = \begin{bmatrix} \mathbf{x}_1 - \bar{\mathbf{x}} \\ \vdots \\ \mathbf{x}_n - \bar{\mathbf{x}} \end{bmatrix}$ and vector \mathbf{y} is defined as: $\mathbf{y} = \begin{bmatrix} y_1 - \bar{y} \\ \vdots \\ y_n - \bar{y} \end{bmatrix}$

and I is an identity matrix of size m -by- m

Equating gradient of L to zero vector and solving for $\boldsymbol{\theta}$ gives us:

$$\boldsymbol{\theta} = (X^T X + \gamma I)^{-1} X^T \mathbf{y}$$

Quick review: Gradient of a function

Consider a function of two variables as an example:


$$f(x, y) = -(\cos^2 x + \cos^2 y)^2$$

$$\nabla f(x, y) = \begin{bmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \end{bmatrix} = \begin{bmatrix} 4(\cos^2(x) + \cos^2(y)) \cos(x) \sin(x) \\ 4(\cos^2(x) + \cos^2(y)) \cos(y) \sin(y) \end{bmatrix}$$

Note 1: f is a function of **two variables**,
so gradient (partial derivatives collected as a vector)
of f is a **two-dimensional vector**

Note 2: Gradient (vector) of f points toward the
steepest ascent for f

Note 3: At a (local) minimum of f its gradient
becomes a **zero vector**



Example source: Wikipedia

Gradient in our toy example...

Toy example:

| x_1 | x_2 | y |
|-------|-------|----|
| 1 | 2 | -1 |
| 3 | -4 | 7 |
| 6 | 2 | 3 |
| -3 | 5 | -4 |
| 7 | -3 | 5 |
| 4 | 3 | ? |

$$\frac{\partial L}{\partial \theta_1} = (64.8)\theta_1 - (39.6)\theta_2 - 57$$

$$\frac{\partial L}{\partial \theta_2} = (-39.6)\theta_1 + (57.2)\theta_2 + 63$$

So, gradient is

$$\nabla L_{\theta} = \begin{bmatrix} (64.8)\theta_1 - (39.6)\theta_2 - 57 \\ (-39.6)\theta_1 + (57.2)\theta_2 + 63 \end{bmatrix}$$

Note: For this toy problem, I assumed $\gamma = 0$ for convenience

Equating gradient (vector) to 0 (vector) and solving, we get:

$$\theta = \begin{bmatrix} \theta_1 \\ \theta_2 \end{bmatrix} = \begin{bmatrix} 0.3580 \\ -0.8535 \end{bmatrix}$$

Equivalently, using direct formula...

Let's take a toy example:

| x_1 | x_2 | y |
|-------|-------|----|
| 1 | 2 | -1 |
| 3 | -4 | 7 |
| 6 | 2 | 3 |
| -3 | 5 | -4 |
| 7 | -3 | 5 |
| 4 | 3 | ? |

$$X = \begin{bmatrix} \mathbf{x}_1 - \bar{\mathbf{x}} \\ \vdots \\ \mathbf{x}_n - \bar{\mathbf{x}} \end{bmatrix} = \begin{bmatrix} 1 - 2.8 & 2 - 0.4 \\ 3 - 2.8 & -4 - 0.4 \\ 6 - 2.8 & 2 - 0.4 \\ -3 - 2.8 & 5 - 0.4 \\ 7 - 2.8 & -3 - 0.4 \end{bmatrix} \quad \mathbf{y} = \begin{bmatrix} y_1 - \bar{y} \\ \vdots \\ y_n - \bar{y} \end{bmatrix} = \begin{bmatrix} -1 - 2 \\ 7 - 2 \\ 3 - 2 \\ -4 - 2 \\ 5 - 2 \end{bmatrix}$$

$$\boldsymbol{\theta} = (X^T X)^{-1} X^T \mathbf{y} = \begin{bmatrix} 0.3580 \\ -0.8535 \end{bmatrix}$$

So, finally the prediction for the test data point

$$? = \sum_{j=1}^m \theta_j (x_j - \bar{x}_j) + b = 0.3580(4 - 2.8) - 0.8535(3 - 0.4) + 2 = 0.2105$$

Note: For this problem I did not add any regularization, i.e., I assumed $\gamma = 0$

MNIST Dataset



Classify images into digits

Each image is **28x28**

10 labels

55,000 training images

5,000 validation images

10,000 test images.

Linear regression on MNIST dataset



Small 28 pixels-by-28 pixels images of hand written digits

The visual recognition problem definition:
to recognize the digit from an image

Our very first line of attack would be to
use linear regression.

Feature dimension, $m = 28 * 28 = 784$

Let's look at our PyTorch implementations:
Called direct method because we will use math formula
to find θ and b.

| Pixel values (feature) | | | | | Digit |
|------------------------|-------|-----|-----------|-----|-------|
| x_1 | x_2 | ... | x_{784} | y | |
| 0.1 | 0.3 | ... | 0.0 | 0 | |
| 0.2 | 0.1 | ... | 0.5 | 1 | |
| ... | ... | ... | ... | ... | |
| ... | ... | ... | ... | ... | |
| 0.0 | 0.98 | ... | 0.8 | 9 | |
| 0.5 | 0.25 | ... | 0.36 | ? | |
| 0.1 | 0.95 | ... | 0.1 | ? | |

Linear regression

So, far we have seen:



[28x28]

Array of real numbers
(784 numbers in total)

Image

(1x784)

Parameters

(784x1)

$$\mathbf{y}^p = (\mathbf{x} - \bar{\mathbf{x}}) \boldsymbol{\theta} + \bar{y}$$

Mean vector
of training
images

Mean of
training
labels (digits)

1 number,
indicating digit

See notebook: MNIST_Linear_Regression_Direct.ipynb

| Pixel values (feature) | | | | Digit |
|------------------------|-------|-----|-----------|-------|
| x_1 | x_2 | ... | x_{784} | y |
| 0.1 | 0.3 | ... | 0.0 | 0 |
| 0.2 | 0.1 | ... | 0.5 | 1 |
| ... | ... | ... | ... | ... |
| ... | ... | ... | ... | ... |
| 0.0 | 0.98 | ... | 0.8 | 9 |
| 0.5 | 0.25 | ... | 0.36 | ? |
| 0.1 | 0.95 | ... | 0.1 | ? |

Multiple or Vector Linear Regression



[28x28]
Array of real
numbers (784
numbers total)

Image
(1x784)

Parameters
(784x10)

$$\mathbf{y}^p = (\mathbf{x} - \bar{\mathbf{x}}) \mathbf{W} + \bar{\mathbf{y}}$$

Mean
vector
of
training
images

Mean of
1-hot
training
label
vector

10 numbers,
indicating class
scores

Pixel values (feature) Digit: 1-hot vector

| x_1 | x_2 | ... | x_{784} | y_1 | ... | y_{10} |
|-------|-------|-----|-----------|-------|-----|----------|
| 0.1 | 0.3 | ... | 0.0 | 0 | ... | 1 |
| 0.2 | 0.1 | ... | 0.5 | 1 | | 0 |
| ... | ... | ... | ... | ... | ... | ... |
| ... | ... | ... | ... | ... | ... | ... |
| 0.0 | 0.98 | ... | 0.8 | 0 | ... | 1 |
| 0.5 | 0.25 | ... | 0.36 | ? | ... | ? |
| 0.1 | 0.75 | ... | 0.1 | ? | ... | ? |

Multiple Linear Regression: PyTorch Implementation

See notebook: MNIST_Multiple_Linear_Regression_Direct.ipynb

Prediction model: $\mathbf{y}^p = (\mathbf{x} - \bar{\mathbf{x}})W + \bar{\mathbf{y}}$

https://en.wikipedia.org/wiki/Matrix_calculus

Regularized loss function: $L = \frac{1}{2} \sum_{i=1}^n \|\mathbf{y}_i^p - \mathbf{y}_i\|^2 + \frac{\gamma}{2} \|W\|^2$

Gradient of loss function:

$$\nabla_W L = (X^T X + \gamma I)W - X^T Y$$

where matrix X is defined as:


$$X = \begin{bmatrix} \mathbf{x}_1 - \bar{\mathbf{x}} \\ \vdots \\ \mathbf{x}_n - \bar{\mathbf{x}} \end{bmatrix}$$

and matrix Y is defined as: $Y = \begin{bmatrix} \mathbf{y}_1 - \bar{\mathbf{y}} \\ \vdots \\ \mathbf{y}_n - \bar{\mathbf{y}} \end{bmatrix}$

and I is an identity matrix of size 784-by-784

We will “minimally” modify our linear regression scripts into multiple linear regression implementations!

This derivation requires
matrix-vector
differentiation



Equating gradient of L to zero
matrix and solving for W gives us:

$$W = (X^T X + \gamma I)^{-1} X^T Y$$

What happened to our learning algorithm?

- Step 1: Create training image set (example set):

```
0000000000000000  
1111111111111111  
2222222222222222  
3333333333333333  
4444444444444444  
5555555555555555  
6666666666666666  
7777777777777777  
8888888888888888  
9999999999999999
```

Repeat steps 2, 3 and 4

- Step 2: Show these examples to the machine learner
- Step 3: Measure mistakes made by the machine learner
- Step 4: Tune parameters of the machine learner to minimize its mistakes

Can we apply this learning algorithm to linear or multiple linear regression using PyTorch?

Iterate:

(Load Data): Get a training data batch (also called mini batch)

(Predict): Apply linear model to training feature vector and compute predictions

(Compute loss): Measure discrepancy between predictions and ground truths

(Optimize): Ask PyTorch to reduce loss value by tuning the parameters θ (or W), and b

(Diagnostics): Check if loss is decreasing

See notebook: MNIST_Linear_Regression.ipynb, MNIST_Multiple_Linear_Regression.ipynb

What are the pros and cons of this optimization-based method over the direct formula-based method?



Gradient descent optimization

Start at an initial guess for the optimization variable: \mathbf{x}_0

Iterate until gradient magnitude becomes too small: $\mathbf{x}^{t+1} = \mathbf{x}^t - \alpha \nabla f(\mathbf{x}^t)$

} Gradient descent algorithm

α is called the step-length.



Gradient descent creates a zig-zag path leading to a local minimum of f

Picture source: Wikipedia



Gradient descent visualization

Find x such that $f(x)$ is minimized: $f(x) = \sin(x) + x + x * \sin(x)$

$$\nabla f(x) = \cos(x) + 1 + \sin(x) + x * \cos(x)$$

$$\mathbf{x}_0 = ???$$

$$\mathbf{x}^{t+1} = \mathbf{x}^t - \alpha \nabla f(\mathbf{x}^t)$$



PyTorch optimizer uses GD

Let's try our own gradient descent for multiple linear regression

Gradient of loss function for multiple linear regression: $\nabla_W L = (X^T X + \gamma I)W - X^T Y$

$$\frac{\partial L}{\partial b} = \sum_{i=1}^n (y_i^p - y_i)$$

Exercise: write GD for MNIST multiple linear regression

Look at `MNIST_Multiple_Linear_Regression.ipynb`

Logistic Regression

Can we modify scores from multiple regression function to output probabilities?

What is a suitable loss function for classification?

Logistic Regression

Would it not be nice if we can predict **class probabilities** instead of scores?



[28x28]

Array of real numbers
(784 numbers total)

image parameters

$$\mathbf{y}^p = f(\mathbf{x}, \mathbf{W})$$

prediction function
For logistic regression

10 numbers,
indicating class
probabilities

| Pixel values (feature) | | | | | | Digit: 1-hot vector |
|------------------------|-------|-----|-----------|-------|-----|---------------------|
| x_1 | x_2 | ... | x_{784} | y_1 | ... | y_{10} |
| 0.1 | 0.3 | ... | 0.0 | 0 | ... | 1 |
| 0.2 | 0.1 | ... | 0.5 | 1 | ... | 0 |
| ... | ... | ... | ... | ... | ... | ... |
| ... | ... | ... | ... | ... | ... | ... |
| 0.0 | 0.98 | ... | 0.8 | 0 | ... | 1 |
| 0.5 | 0.25 | ... | 0.36 | ? | ... | ? |
| 0.4 | 0.95 | ... | 0.1 | ? | ... | ? |

Logistic regression: from multiple linear regression

Scores from multiple linear regression:

$$\mathbf{s}_i = (\mathbf{x}_i - \bar{\mathbf{x}})W + \bar{\mathbf{y}}$$

$$\mathbf{s}_{i,k} = (\mathbf{x}_i - \bar{\mathbf{x}})W_{:,k} + \bar{\mathbf{y}}_k$$


Score for k^{th} class, $k = 0, \dots, 9$

Predicted probability for k^{th} class:

$$y_{i,k}^p = \frac{\exp(s_{i,k})}{\sum_{c=0}^9 \exp(s_{i,c})}$$

“Softmax” function

Computational graphs



$$z = [-1, -3, +10, \dots, 2]$$

$$y_i^p = \frac{\exp(z_i)}{\sum_{k=0}^9 \exp(z_k)}$$

$$y_1^p = \frac{\exp(-1)}{\exp(-1) + \exp(-3) + \dots + \exp(2)}, \quad y_2^p = \frac{\exp(-3)}{\dots}$$

y_i , "ground truths"

$$\sum_{i=0}^q y^p_i = 1$$

Logistic regression: loss function

Cross entropy loss:
$$loss(\mathbf{y}^p, \mathbf{y}) = - \sum_{k=0}^9 \mathbf{y}_k \log(\mathbf{y}_k^p)$$

Why this loss function? What does it mean? Why not use Euclidean loss as in MLR?

Do we have a direct formula to compute parameters like MLR?