

# Introduction to Neural Networks and Backpropagation

Computing Science

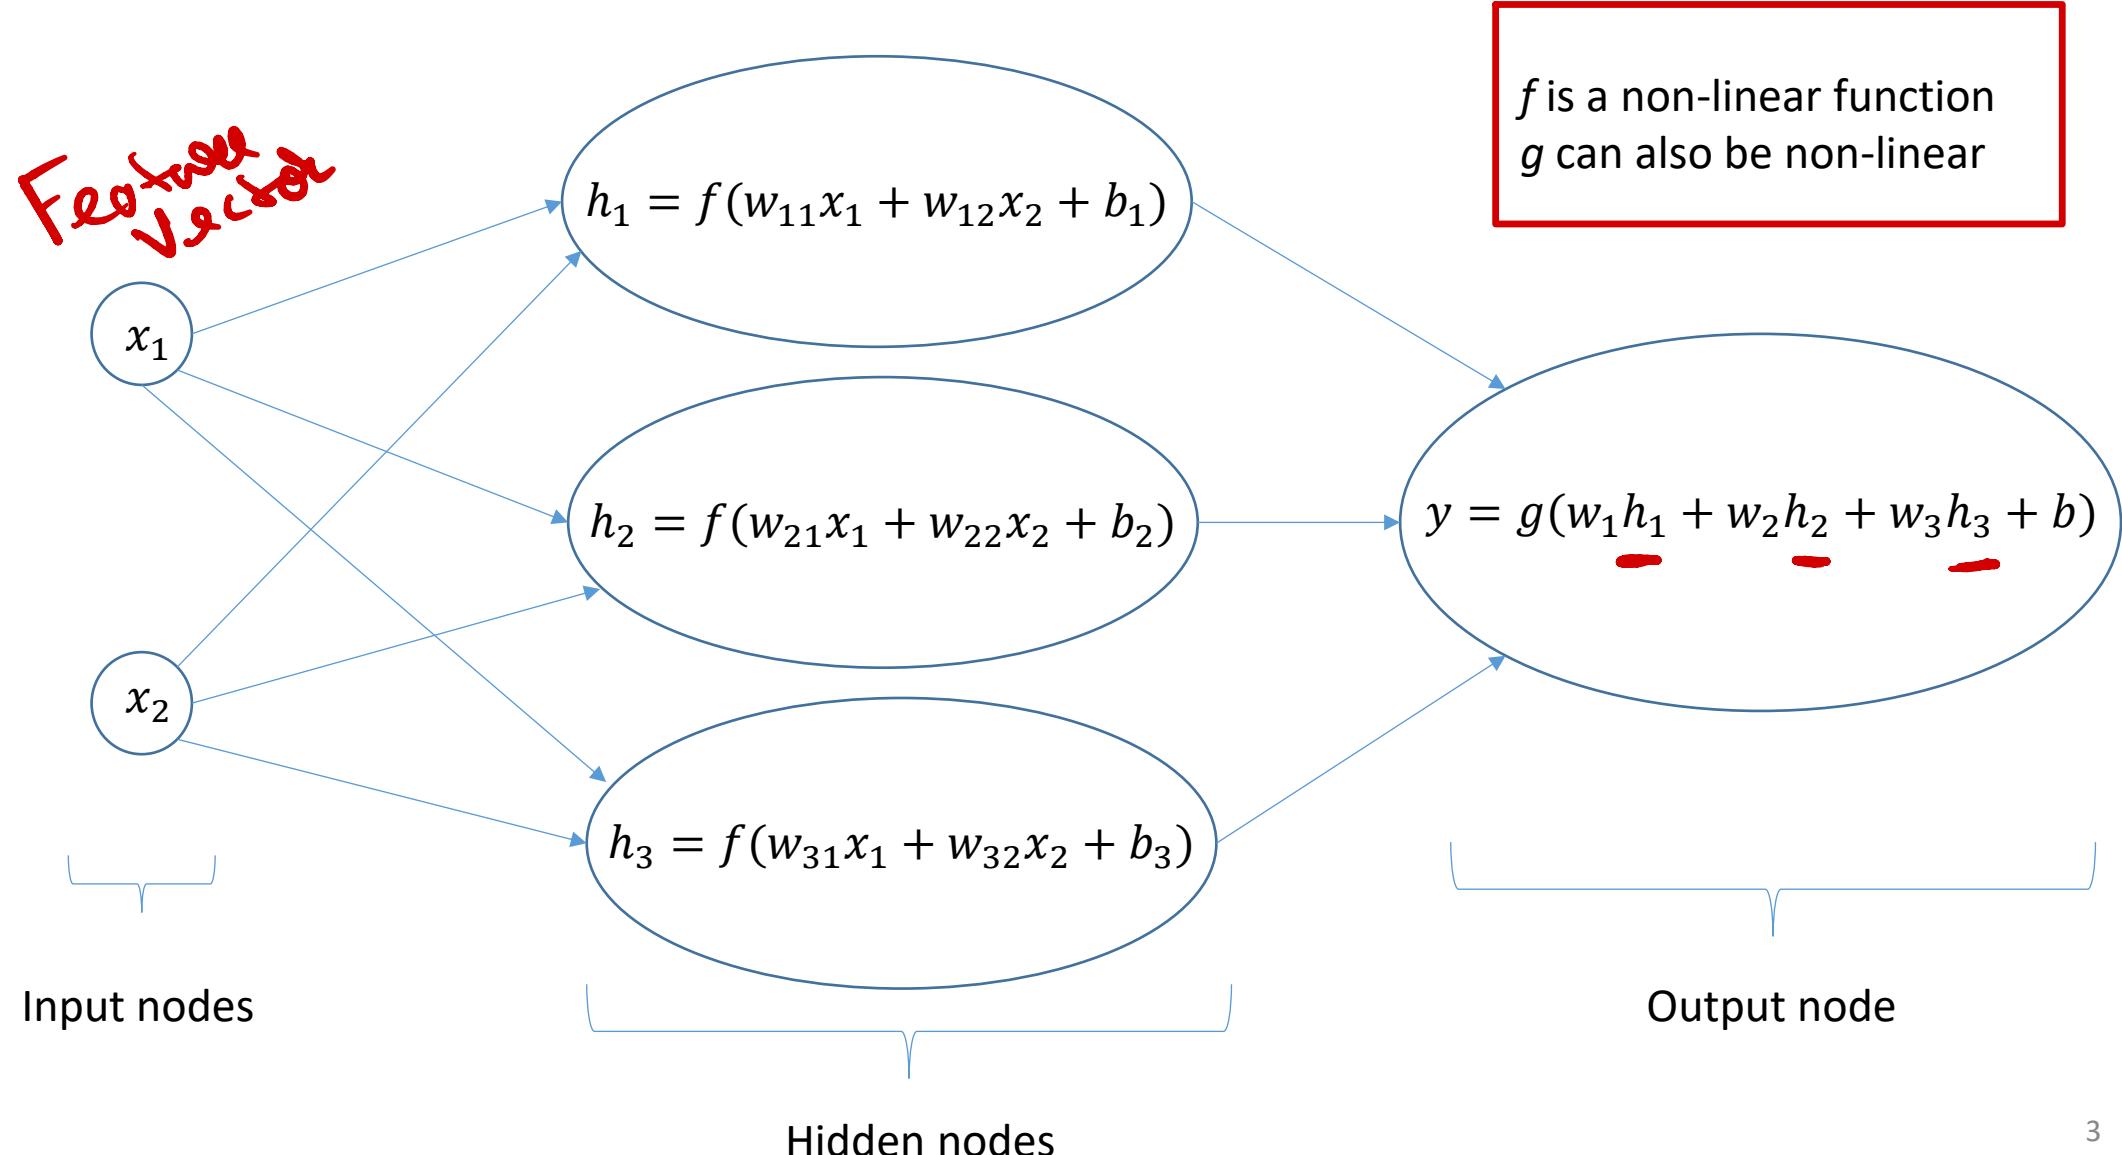
University of Alberta

Nilanjan Ray

# Agenda

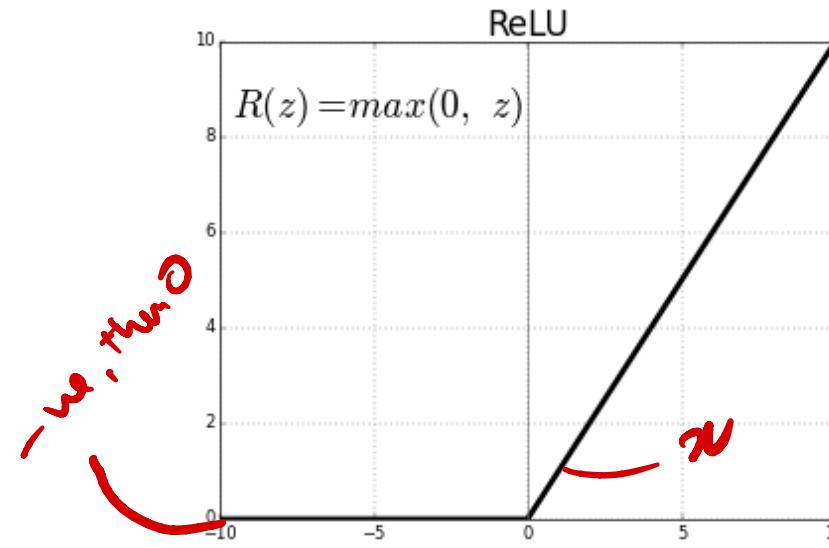
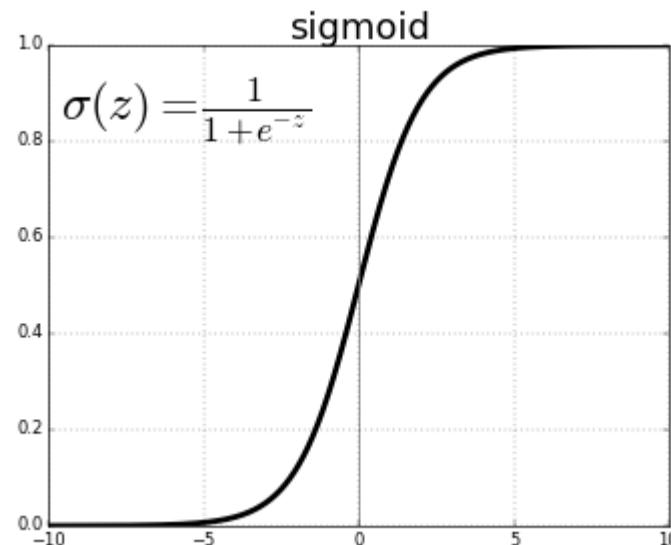
- What is a Neural Net?
  - Neural net as a computational graph
- Approximating “XOR” function with neural net
- Applying a neural net to classify MNIST
- Universal function approximation by a neural net
- (re)Introduction to gradient descent optimization
- Chain rule of derivatives
- Understanding backpropagation algorithm

# Feed forward neural network



# Feed forward net: non-linear functions

- Non-linear functions at hidden nodes are known as “activation function”
  - Sigmoid, ReLU, ELU, ....

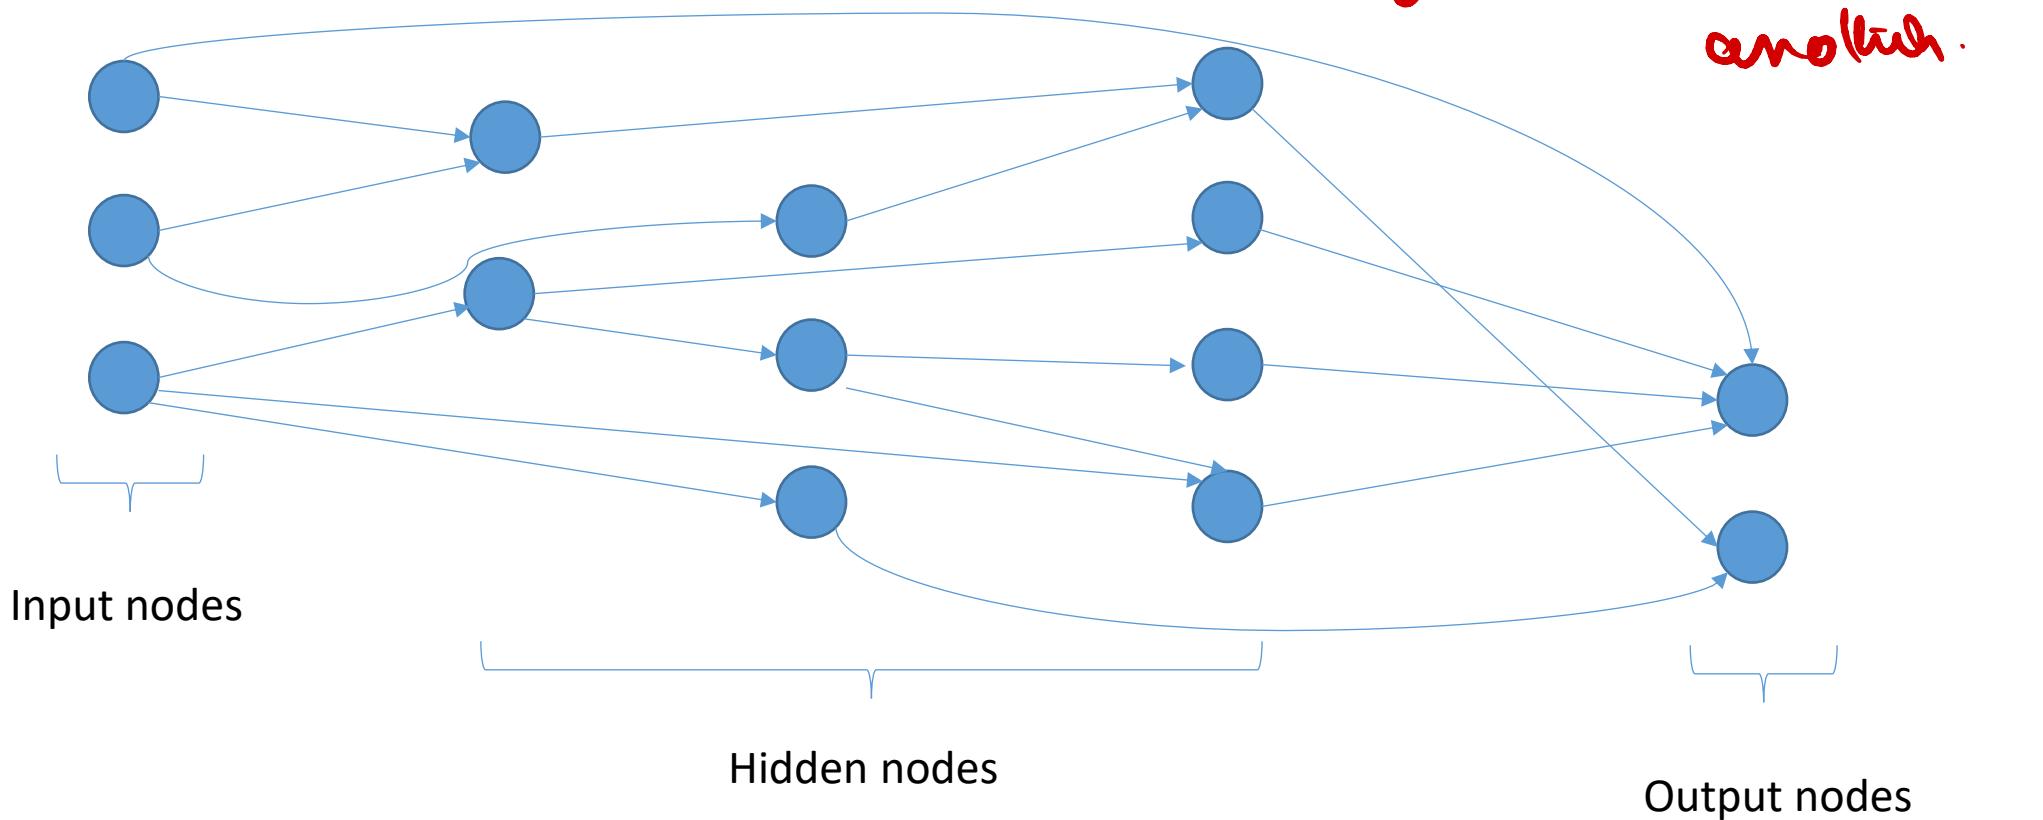


Why activation functions are non-linear?

# Feedforward net in general: Directed acyclic graph

one side to  
another side

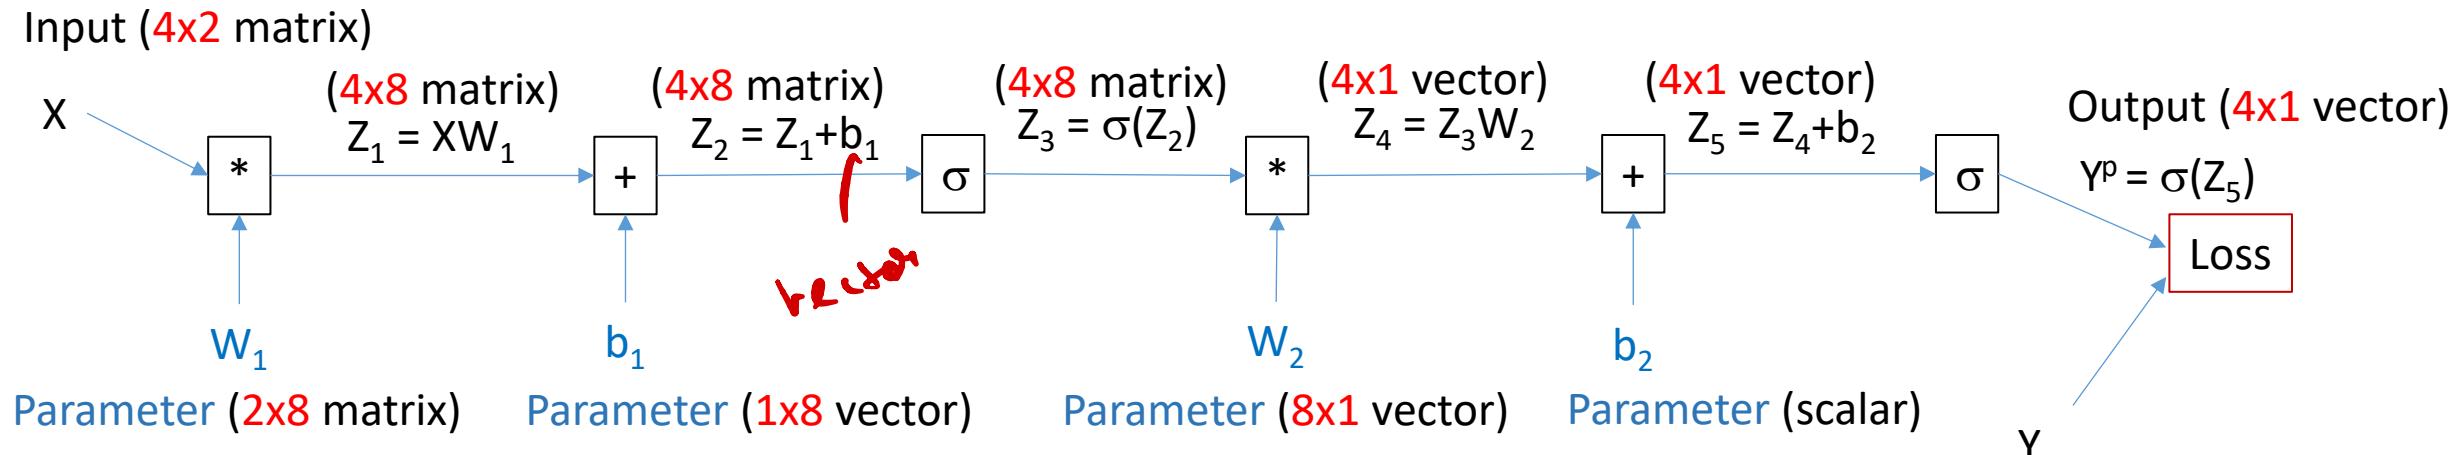
No loop  
flow of info from one side to  
another.



# What's the big deal about neural net?

- Mathematically rich: it can approximate any function
- It is biologically inspired: (loosely) resembles brain connections
- Computationally:
  - Simple: matrix-vector multiplication and point-wise non-linear function
  - Highly parallelizable: cuBLAS, GEMM, Batched GEMM!
- Excellent **empirical** results on “generalization capability” over variety of applications!

# Neural network as a computational graph

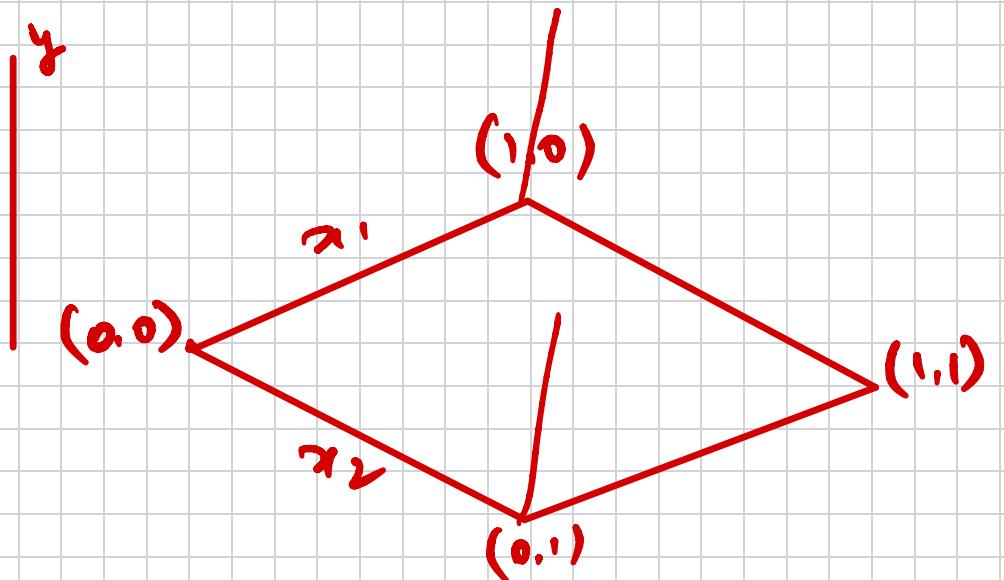


Sigmoid function;  
applied **pointwise**  
to a vector or  
matrix input

See Learn\_XOR.ipynb

This network is  
trying to learn  
XOR function

X	Y
1	0
0	0
0	1
1	1



# How does PyTorch optimize parameters?

- By **gradient descent** PyTorch adjusts network parameters to reduce the value of the loss function.
- But how?
  - Answer: Backpropagation
  - We will learn to do backpropagation on a computational graph later

# Using PyTorch to Learn XOR

Learning algorithm has this basic structure as we have already seen in logistic regression

- Define an architecture for the neural network and instantiate it
- Instantiate an optimizer to adjust the parameters of the neural net
- Iterate
  - Load data  $(X, y)$
  - Do a **forward pass**, i.e., compute output of neural net:  $f(X; \theta)$
  - Do a **backward pass**:
    - Compute loss  $L(f(X; \theta), y)$ . The function  $L$  (loss) measures discrepancy between ground truth annotation  $y$  and the output of the neural net  $f(X; \theta)$
    - Adjust parameters  $\theta$  of the neural network to reduce loss value
  - **Diagnostic**: from time to time print loss value

# MNIST classification problem



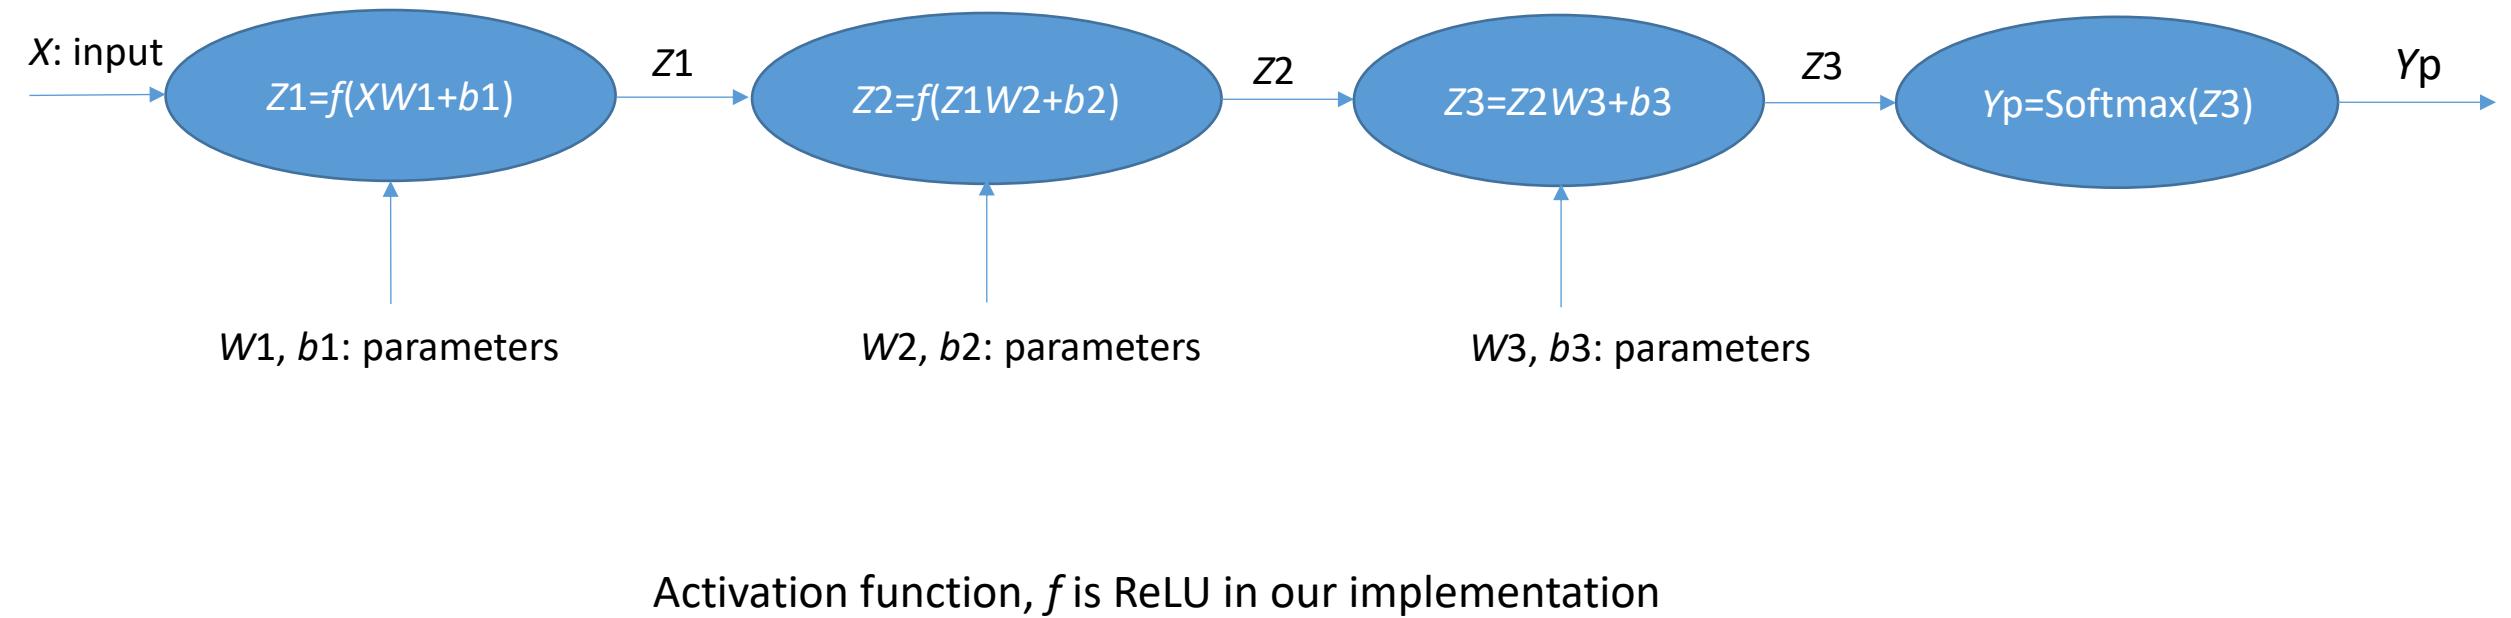
Small 28 pixels-by-28 pixels images of hand written digits

The visual recognition problem definition:  
to recognize the digit from an image



$x_1$	$x_2$	...	$x_{784}$	$y_1$	...	$y_{10}$
0.1	0.3	...	0.0	0	...	1
0.2	0.1	...	0.5	1	...	0
...	...	...	...	...	...	...
...	...	...	...	...	...	...
0.0	0.98	...	0.8	0	...	1
0.5	0.25	...	0.36	?	...	?
0.1	0.95	...	0.1	?	...	?

# NN Architecture for MNIST Classification



# Learning MNIST NN with Backprop and SGD

Initialize all parameters of the neural network

Initialize learning rate variable  $lr$

Iterate:

(Load Data): Get training data batch  $X$

(Forward pass): Compute  $Z1, Z2, Z3, Yp$

(Compute loss): Compute a suitable loss between ideal output  $Y$  and output of NN  $Yp$

(Backward pass): Ask PyTorch optimizer to adjust neural network parameters

(Diagnostics): Compute loss on training and validation sets

# Neural Net as Universal Function Approximator

- <http://neuralnetworksanddeeplearning.com/index.html>

# Gradient Descent: PyTorch under the hood

- How does PyTorch optimizes parameters of a model to reduce loss value?
  - Using gradient descent
- We will apply GD to multiple linear regression
- Then we will move on to using it for a neural net
  - We must learn how to use the chain rule of differentiation

# Gradient of a function

Example:

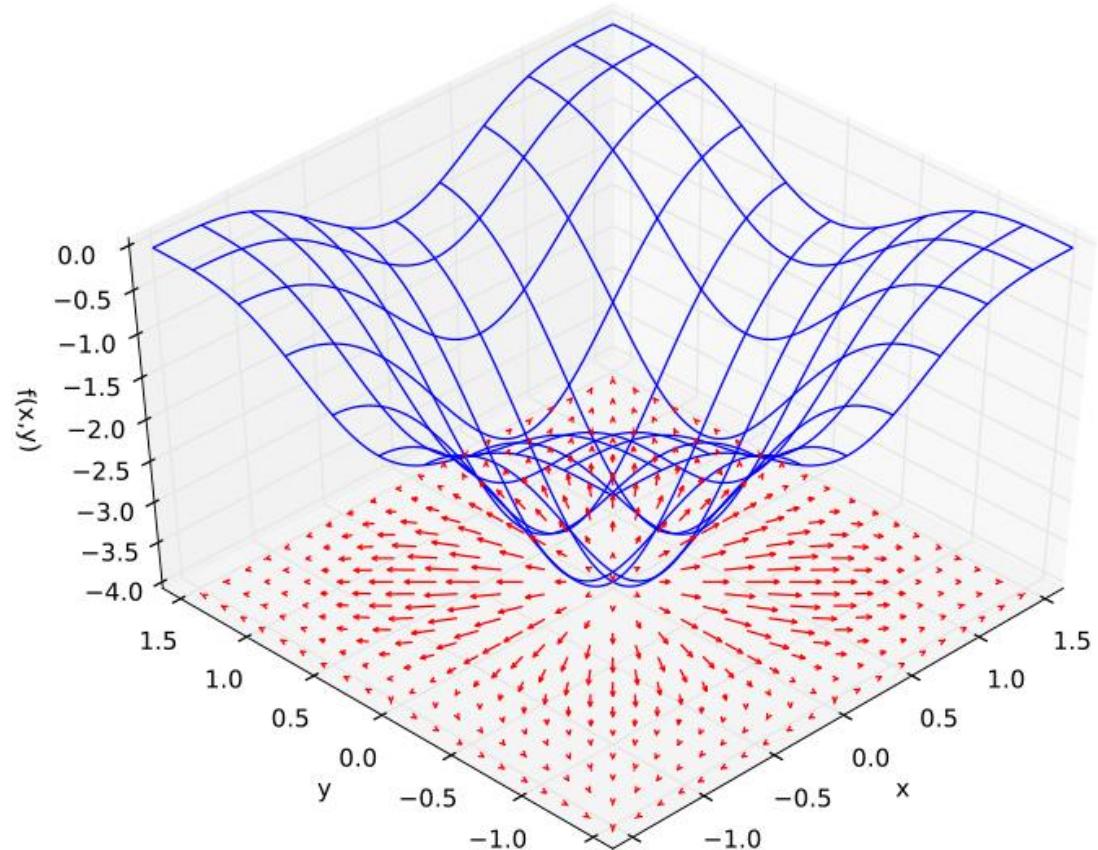
$$f(x, y) = -(\cos^2 x + \cos^2 y)^2$$

$$\nabla f(x, y) = \begin{bmatrix} \frac{\partial f}{\partial x} \\ \frac{\partial f}{\partial y} \end{bmatrix} = \begin{bmatrix} 4(\cos^2(x) + \cos^2(y)) \cos(x) \sin(x) \\ 4(\cos^2(x) + \cos^2(y)) \cos(y) \sin(y) \end{bmatrix}$$

**Note 1:**  $f$  is a function of **two variables**,  
so gradient of  $f$  is a **two dimensional vector**

**Note 2:** Gradient (vector) of  $f$  points toward the  
**steepest ascent for  $f$**

**Note 3:** At a (local) minimum of  $f$  its gradient  
becomes a **zero vector**



Example source: Wikipedia

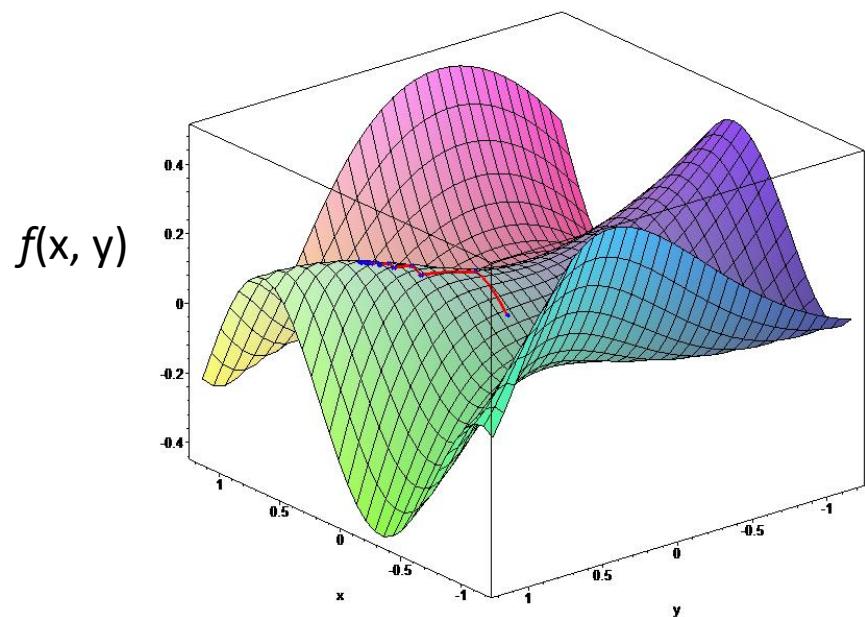
# Gradient descent optimization

Start at an initial guess for the optimization variable:  $\mathbf{x}_0$

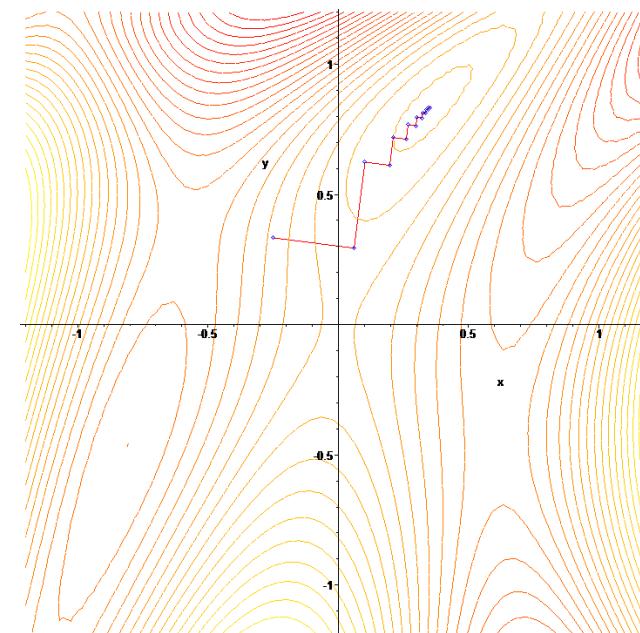
Iterate until gradient magnitude becomes too small:  $\mathbf{x}^{t+1} = \mathbf{x}^t - \alpha \nabla f(\mathbf{x}^t)$

$\alpha$  is called the step-length.

} Gradient descent algorithm



Picture source: Wikipedia



Gradient descent creates a zig-zag path leading to a local minimum of  $f$

Look at  
`GradientDescentDemo.ipynb`

# PyTorch optimizer uses GD

Let's try our own gradient descent for multiple linear regression

Gradient of loss function for multiple linear regression:

$$\nabla_W L = (X^T X + \gamma I)W - X^T Y$$

$$\nabla_b L = \sum_{i=1}^n (y_i^p - y_i)$$

Exercise: write GD for MNIST multiple linear regression

For implementation of this GD, look at  
[MNIST\\_Multiple\\_Linear\\_Regression.ipynb](#)

# How do we apply GD for learning a neural net?

We need to compute gradient of the loss function with respect to all parameters in a neural net:

$$\delta\theta_i \equiv \nabla_{\theta_i} L(y^p, y)$$

Parameter in the  $i^{\text{th}}$  layer

Ground truth/tag

Output (aka prediction) from neural net

Once we have this loss gradient, we can adjust parameters using gradient descent rule:

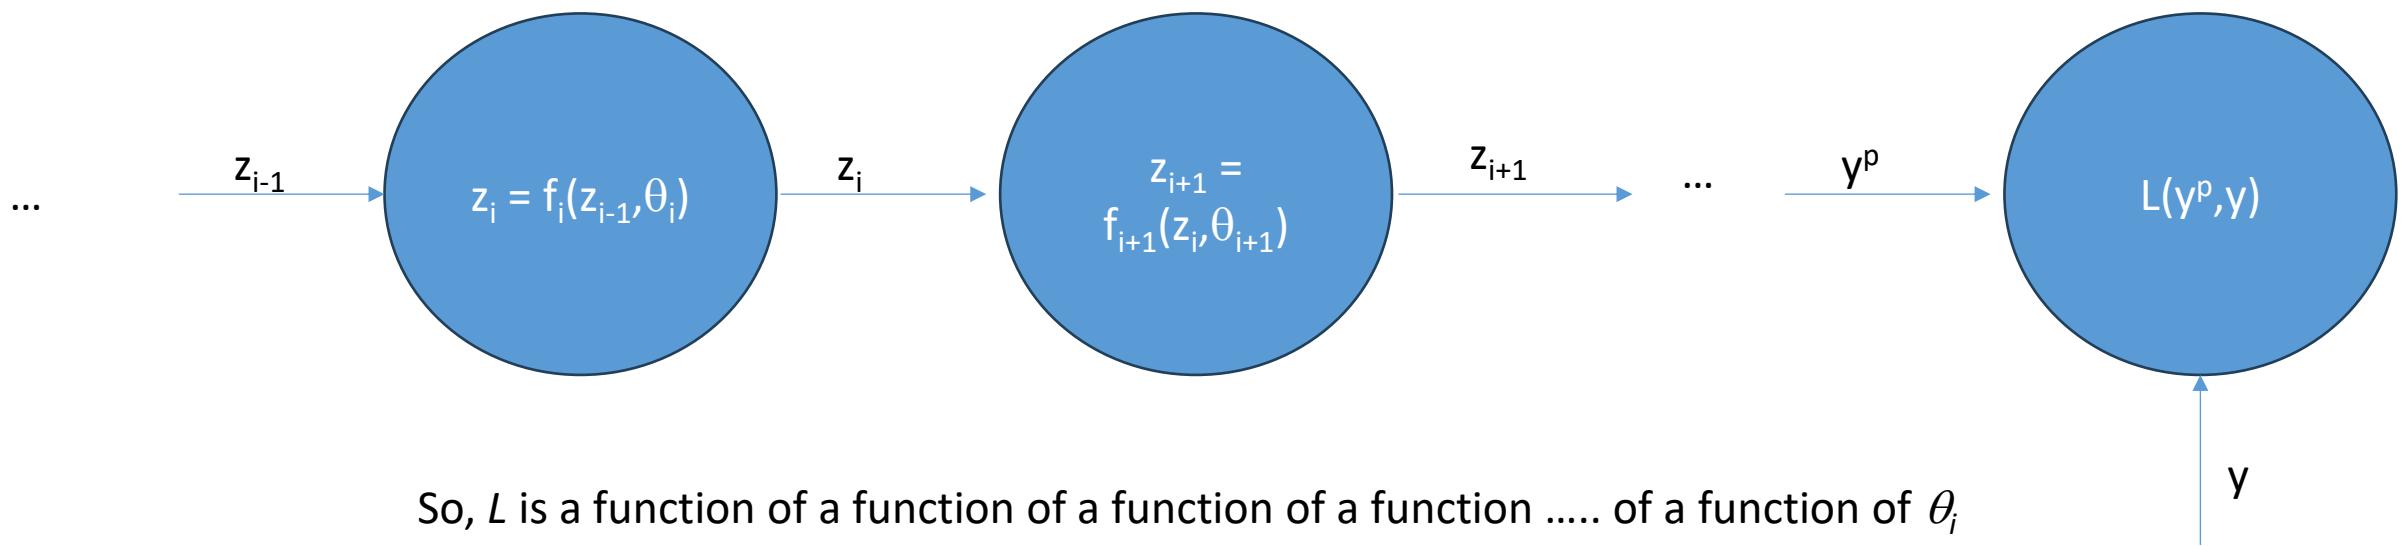
$$\theta_i = \theta_i - \alpha \delta \theta_i$$


Learning rate/step size

# How do we apply GD for learning a neural net?

We need to compute gradient of the loss function with respect to all parameters in a neural net:

$$\delta\theta_i \equiv \nabla_{\theta_i} L(y^p, y)$$



Therefore, we need chain rule of derivative to compute  $\delta\theta_i$

# To apply chain rule of derivative in a neural net...

- We need to understand chain rule of derivative for multivariate functions: Jacobian vector product
- We also need to understand the notion of a computational graph and how to apply Jacobian vector product to a computational graph
- These components will lead us to the well acclaimed **backpropagation** algorithm for learning parameters of a neural net using GD

# Example gradient computations

- Let's consider the following function of four variables:

$$f(x_1, x_2, x_3, x_4) = (x_1 + 10x_2)^2 + 5(x_3 - x_4)^2 + (x_2 + 2x_3)^4 + 10(x_1 - x_4)^4$$

- Let's compute derivative (gradient) of this function at

$$[x_1, x_2, x_3, x_4] = [3, -1, 0, 1]$$

- Cross-verify PyTorch partial derivative computations with math formulas
- Gradient descent optimization

Look into Understanding\_chain\_rule.ipynb

# Chain rule of derivatives

- Let consider the same function as before:

$$f(x_1, x_2, x_3, x_4) = (x_1 + 10x_2)^2 + 5(x_3 - x_4)^2 + (x_2 + 2x_3)^4 + 10(x_1 - x_4)^4$$

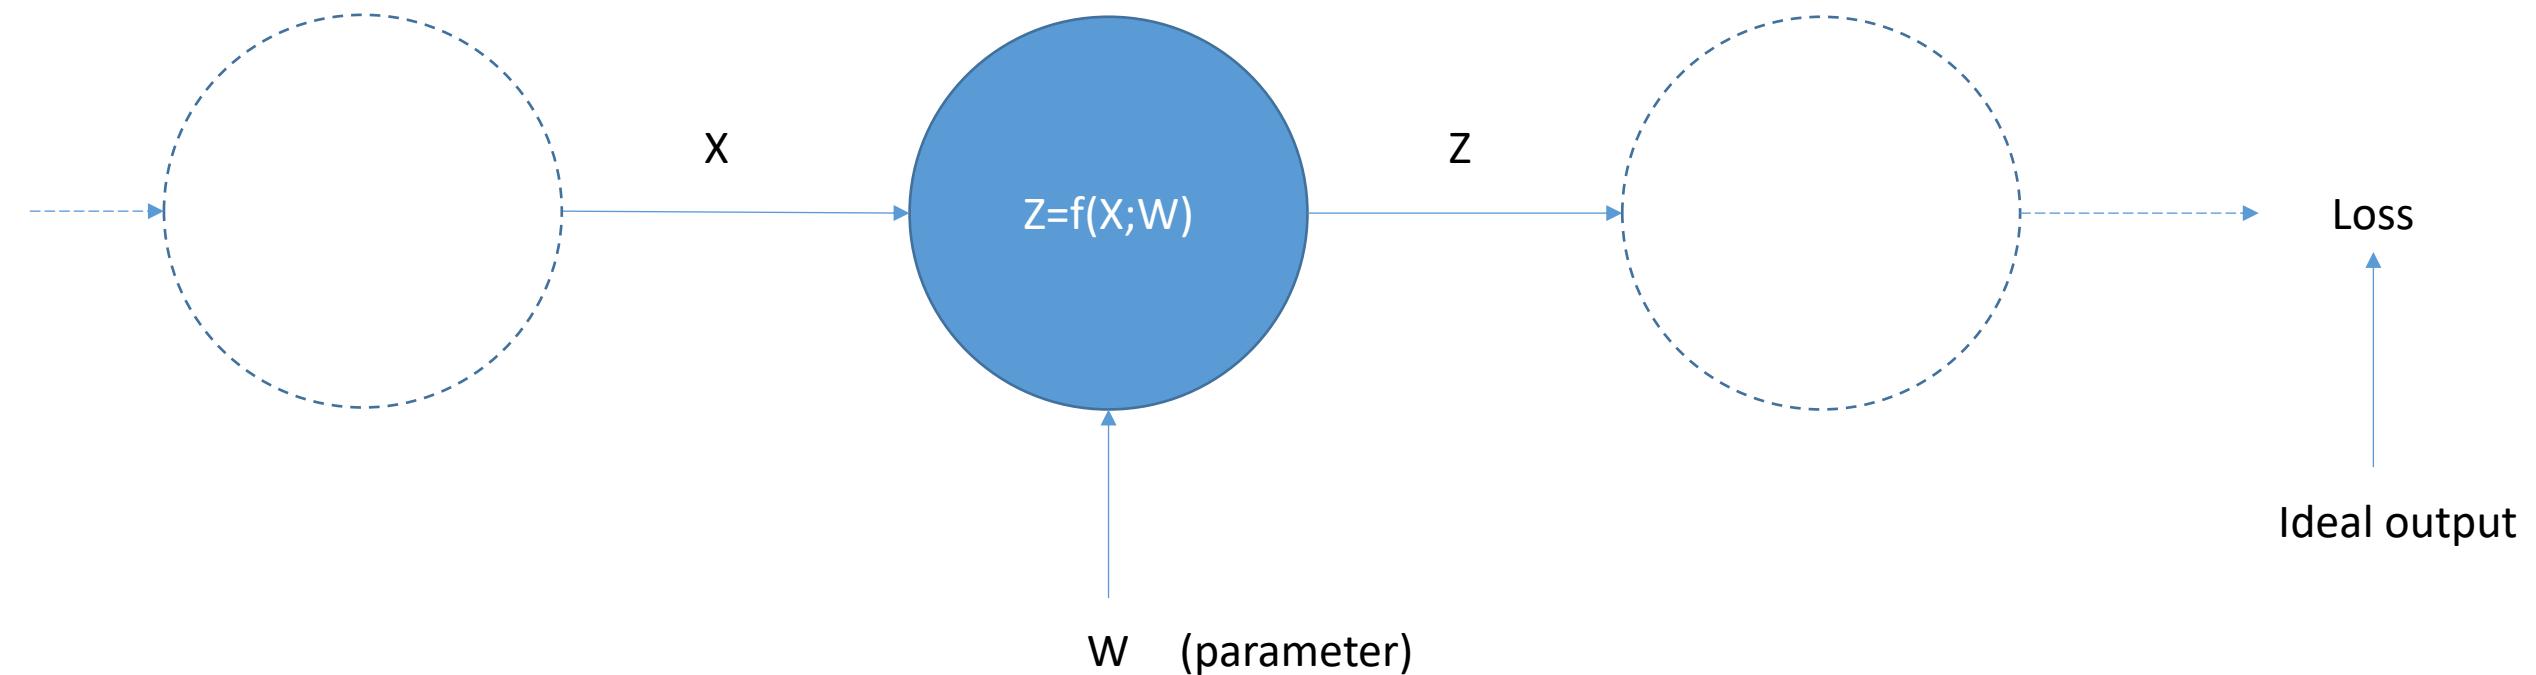
- But this time  $x$  is a (vector-valued) function of two variables  $z_1$  and  $z_2$ :

$$\begin{aligned}x_1 &= z_1 - z_2, \\x_2 &= z_1^2, \\x_3 &= z_2^2, \\x_4 &= z_1^2 + z_1 z_2\end{aligned}$$

- Let's compute gradient of  $f$  with respect to  $z$  using chain rule:  
Jacobian vector product

Look into Understanding\_chain\_rule.ipynb

# Chain rule of derivative for a computational node



If  $X, Z, W$  are all scalars, then usual chain rule of derivative applies:

$$\frac{\partial(\text{Loss})}{\partial X} = \frac{\partial Z}{\partial X} \frac{\partial(\text{Loss})}{\partial Z}$$

$$\frac{\partial(\text{Loss})}{\partial W} = \frac{\partial Z}{\partial W} \frac{\partial(\text{Loss})}{\partial Z}$$

OK, let's apply chain rule to a computational graph where all variables and parameters are scalars

So, our scalar neural net is:

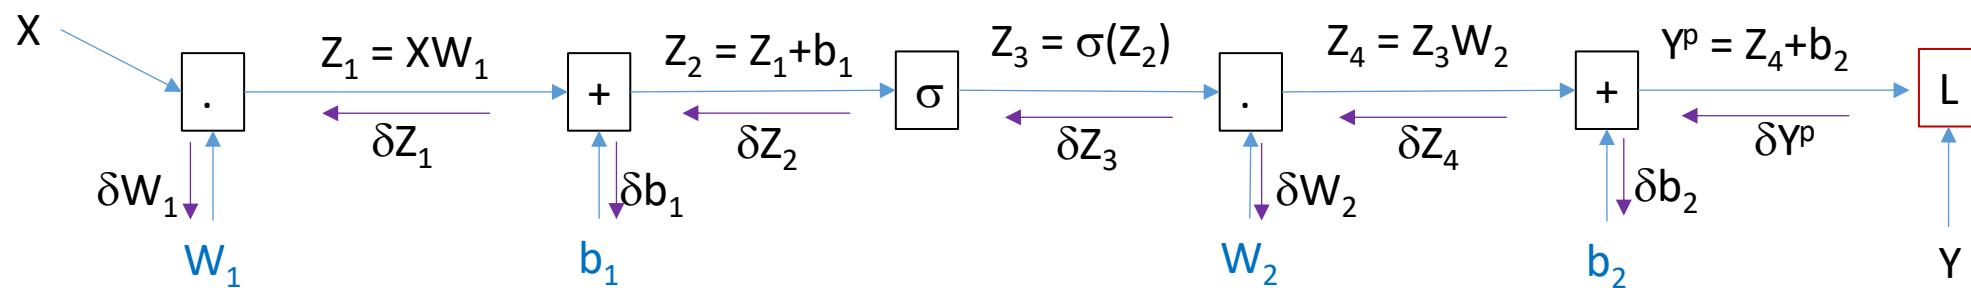
$$Y^p = \sigma(XW_1 + b_1)W_2 + b_2$$

with a square (aka Euclidean) loss function:

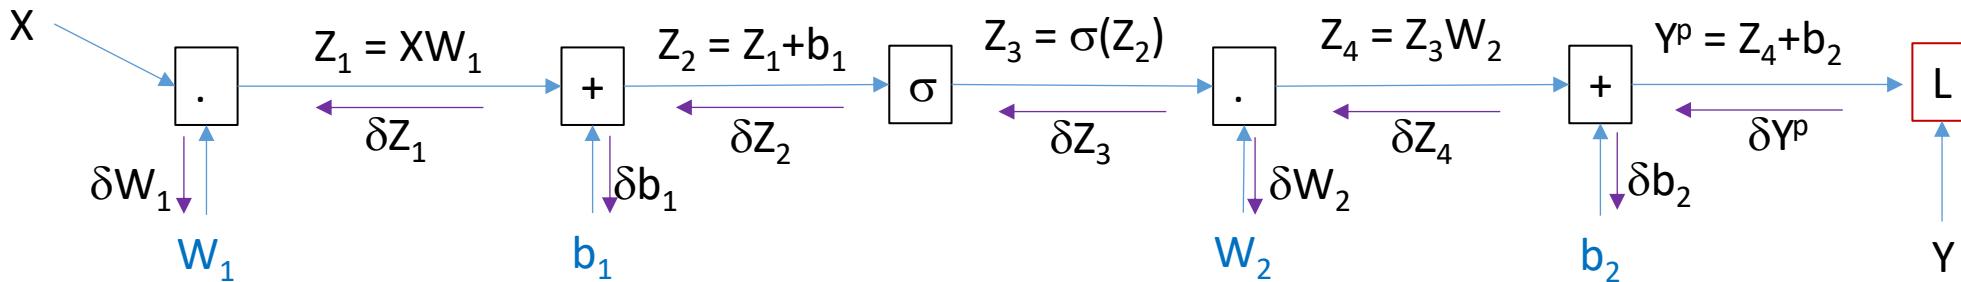
$$L(Y^p, Y) = \frac{1}{2}(Y^p - Y)^2$$

As usual,  $X$  is the input,  $Y^p$  is the output, and  $W_1, b_1, W_2, b_2$  are parameters of the neural net,  $\sigma$  is a non-linear function.

The computational graph for this scalar neural net is (also showing loss gradient symbols):



# Chain rule for a scalar neural net...



$$1 \quad \delta Y^p \equiv \frac{\partial L}{\partial Y^p} = \frac{\partial}{\partial Y^p} \left[ \frac{1}{2}(Y^p - Y)^2 \right] = Y^p - Y$$

$$4 \quad \delta Z_2 \equiv \frac{\partial L}{\partial Z_2} = \frac{\partial Z_3}{\partial Z_2} \frac{\partial L}{\partial Z_3} = \sigma'(Z_2) \delta Z_3$$

$$2 \quad \delta Z_4 \equiv \frac{\partial L}{\partial Z_4} = \frac{\partial Y^p}{\partial Z_4} \frac{\partial L}{\partial Y^p} = \delta Y^p$$

Because  $Z_3 = \sigma(Z_2)$ ,  $\frac{\partial Z_3}{\partial Z_2} = \sigma'(Z_2)$

$$\text{Because } Y^p = Z_4 + b_2, \frac{\partial Y^p}{\partial Z_4} = 1$$

$$5 \quad \delta Z_1 \equiv \frac{\partial L}{\partial Z_1} = \frac{\partial Z_2}{\partial Z_1} \frac{\partial L}{\partial Z_2} = \delta Z_2$$

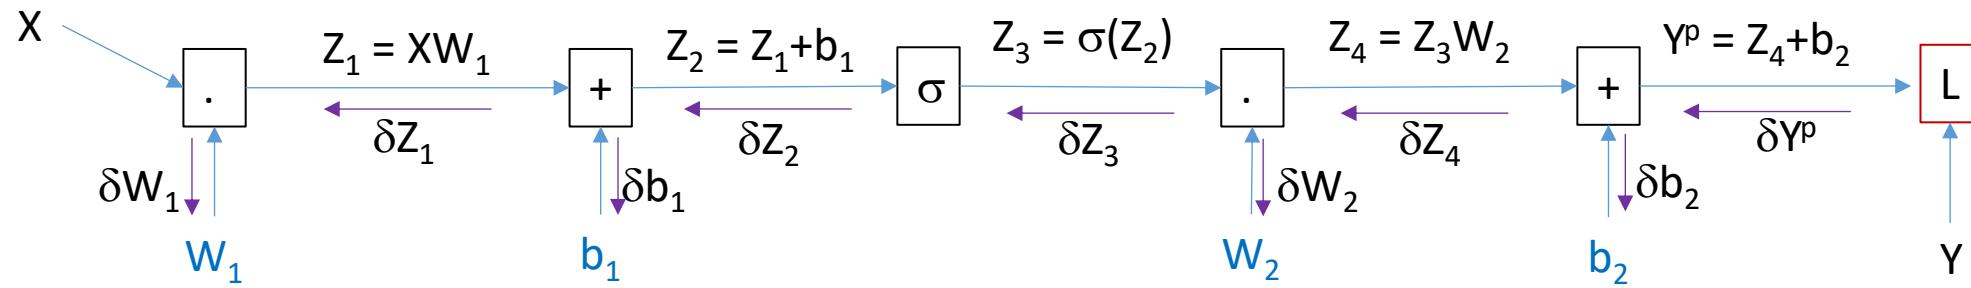
$$3 \quad \delta Z_3 \equiv \frac{\partial L}{\partial Z_3} = \frac{\partial Z_4}{\partial Z_3} \frac{\partial L}{\partial Z_4} = W_2 \delta Z_4$$

$$\text{Because } Z_2 = Z_1 + b_1, \frac{\partial Z_2}{\partial Z_1} = 1$$

$$\text{Because } Z_4 = Z_3 W_2, \frac{\partial Z_4}{\partial Z_3} = W_2$$

But we need loss derivatives with respect to parameters...

# Loss derivatives w.r.t. parameters



$$1 \quad \delta W_1 \equiv \frac{\partial L}{\partial W_1} = \frac{\partial Z_1}{\partial W_1} \frac{\partial L}{\partial Z_1} = X \delta Z_1$$

Because  $Z_1 = XW_1, \frac{\partial Z_1}{\partial W_1} = X$

$$3 \quad \delta W_2 \equiv \frac{\partial L}{\partial W_2} = \frac{\partial Z_4}{\partial W_2} \frac{\partial L}{\partial Z_4} = Z_3 \delta Z_4$$

Because  $Z_4 = Z_3W_2, \frac{\partial Z_4}{\partial W_2} = Z_3$

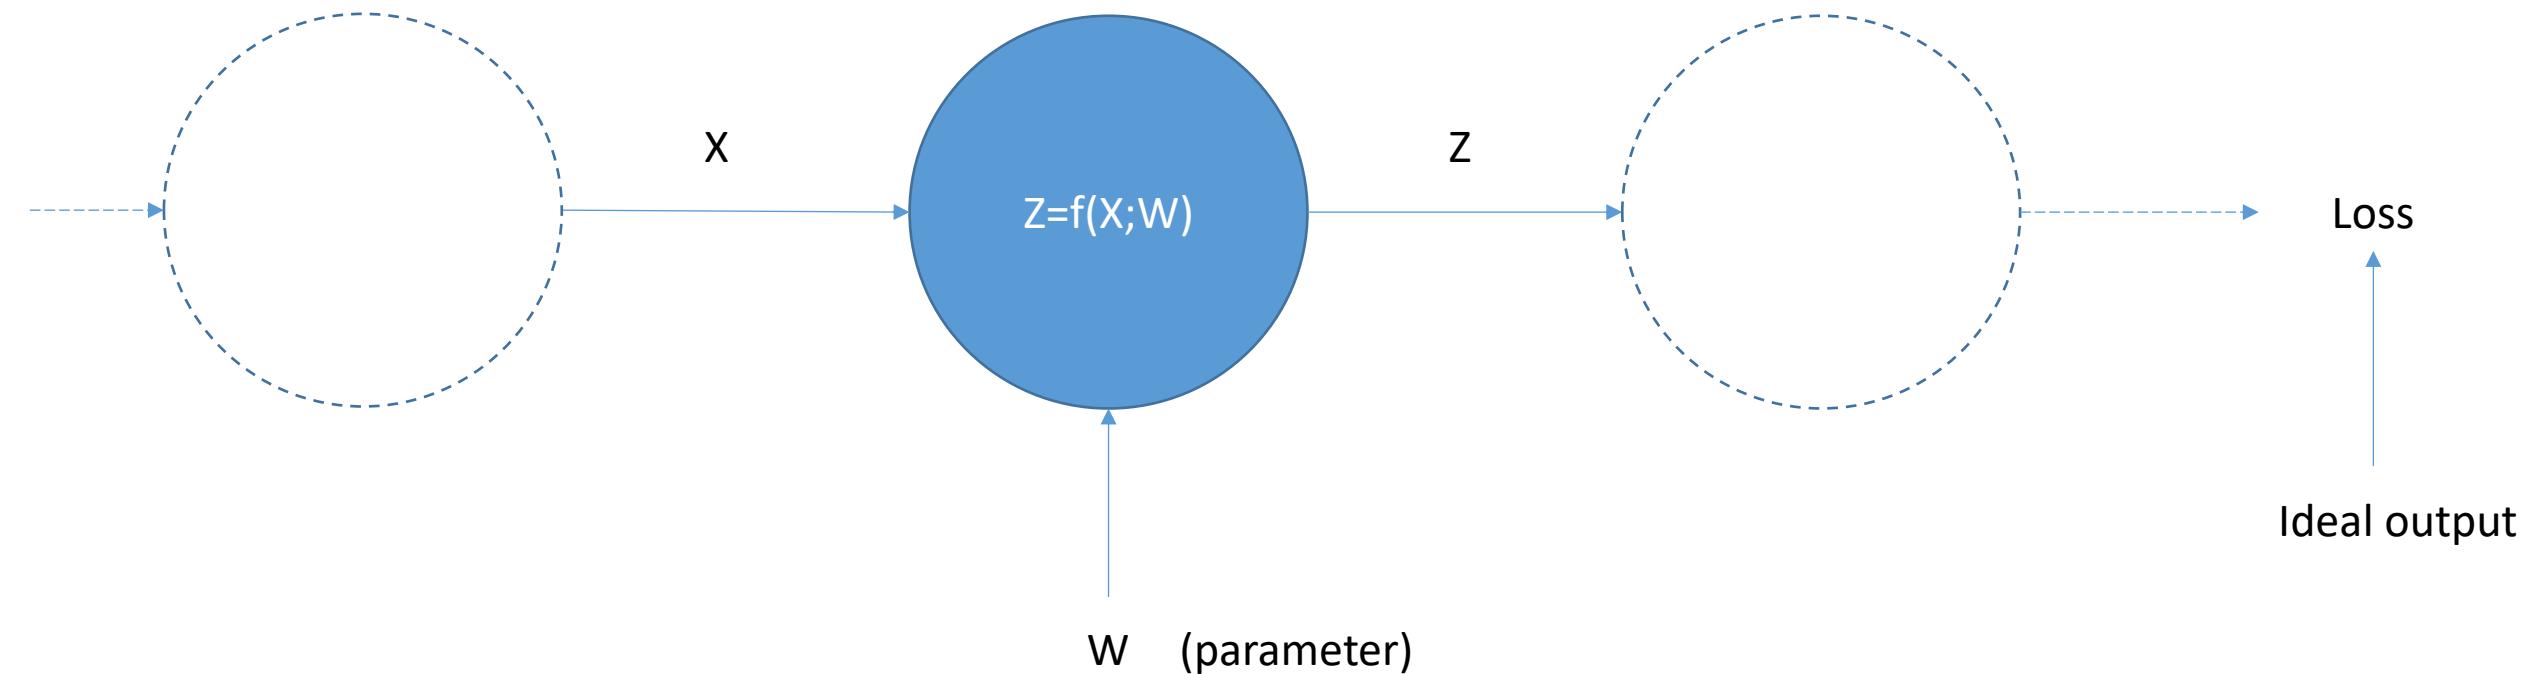
$$2 \quad \delta b_1 \equiv \frac{\partial L}{\partial b_1} = \frac{\partial Z_2}{\partial b_1} \frac{\partial L}{\partial Z_2} = \delta Z_2$$

Because  $Z_2 = Z_1 + b_1, \frac{\partial Z_2}{\partial b_1} = 1$

$$4 \quad \delta b_2 \equiv \frac{\partial L}{\partial b_2} = \frac{\partial Y^p}{\partial b_2} \frac{\partial L}{\partial Y^p} = \delta Y^p$$

Because  $Y^p = Z_4 + b_2, \frac{\partial Y^p}{\partial b_2} = 1$

# Chain rule of derivative...



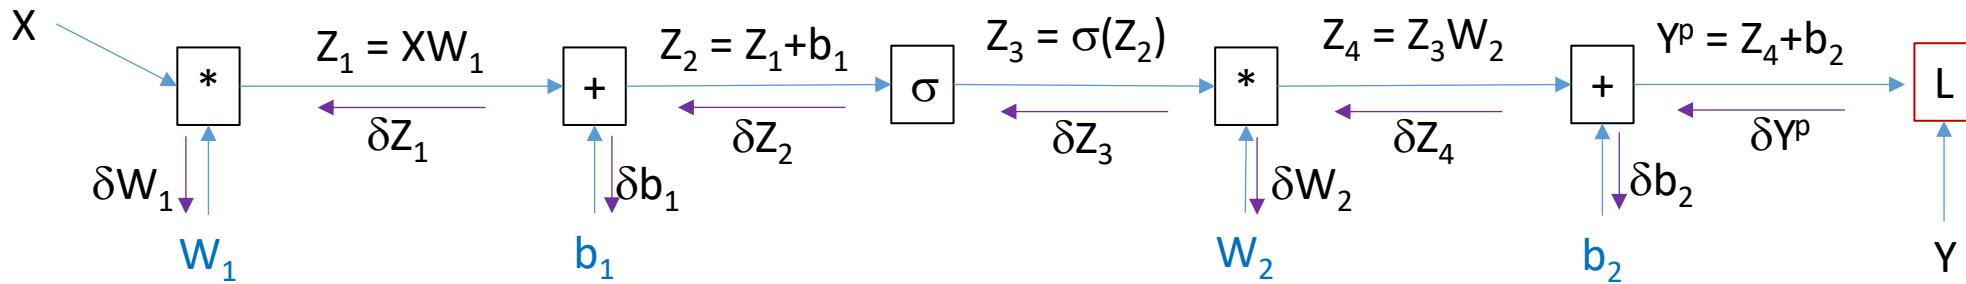
If  $X, Z, W$  are **matrices or vectors**, then :

$$\nabla_X(\text{Loss}) = \left( \frac{\partial Z}{\partial X} \right) * \nabla_Z(\text{Loss})$$

$$\nabla_W(\text{Loss}) = \left( \frac{\partial Z}{\partial W} \right) * \nabla_Z(\text{Loss})$$

“\*” refers to matrix vector multiplication

# Chain rule for a (general) neural net



1  $\delta Y^p \equiv \frac{\partial L}{\partial Y^p} = \frac{\partial}{\partial Y^p} \left[ \frac{1}{2} \|Y^p - Y\|^2 \right] = y^p - y$

4  $\delta Z_2 \equiv \frac{\partial L}{\partial Z_2} = \frac{\partial Z_3}{\partial Z_2} \frac{\partial L}{\partial Z_3} = \sigma'(Z_2) \cdot \delta Z_3$

2  $\delta Z_4 \equiv \frac{\partial L}{\partial Z_4} = \frac{\partial Y^p}{\partial Z_4} \frac{\partial L}{\partial Y^p} = \delta Y^p$

Because  $Z_3 = \sigma(Z_2)$ ,  $\frac{\partial Z_3}{\partial Z_2} = \sigma'(Z_2)$

Because  $Y^p = Z_4 + b_2$ ,  $\frac{\partial Y^p}{\partial Z_4} = 1$

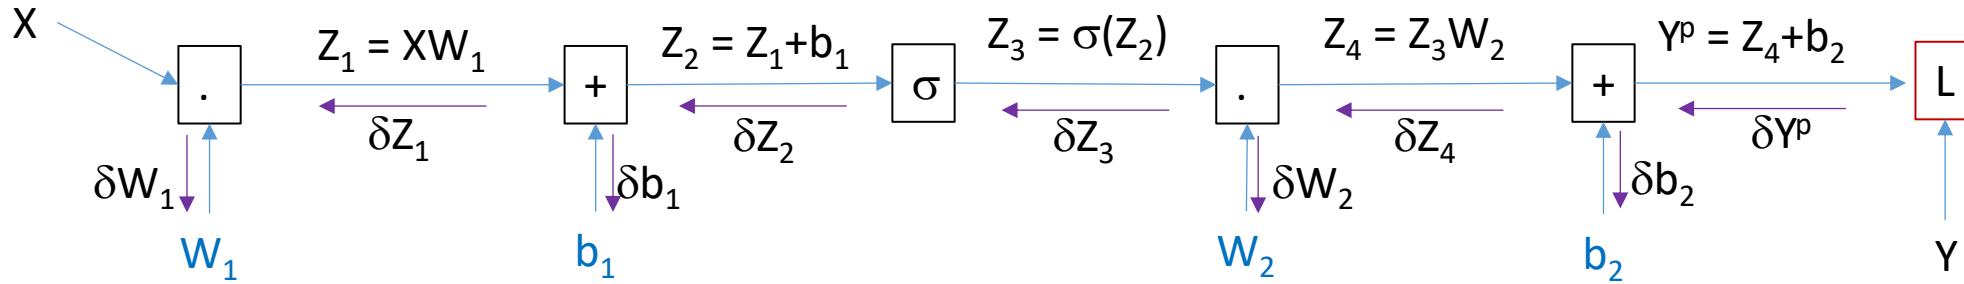
5  $\delta Z_1 \equiv \frac{\partial L}{\partial Z_1} = \frac{\partial Z_2}{\partial Z_1} \frac{\partial L}{\partial Z_2} = \delta Z_2$

3  $\delta Z_3 \equiv \frac{\partial L}{\partial Z_3} = \frac{\partial Z_4}{\partial Z_3} \frac{\partial L}{\partial Z_4} = \delta Z_4 W_2^T$

Because  $Z_2 = Z_1 + b_1$ ,  $\frac{\partial Z_2}{\partial Z_1} = 1$

Because  $Z_4 = Z_3 W_2$ ,  $\frac{\partial Z_4}{\partial Z_3} = W_2^T$

# Loss derivatives w.r.t. matrix or vector parameters



$$1 \quad \delta W_1 \equiv \frac{\partial L}{\partial W_1} = \frac{\partial Z_1}{\partial W_1} \frac{\partial L}{\partial Z_1} = X^T \delta Z_1$$

Because  $Z_1 = XW_1, \frac{\partial Z_1}{\partial W_1} = X^T$

$$3 \quad \delta W_2 \equiv \frac{\partial L}{\partial W_2} = \frac{\partial Z_4}{\partial W_2} \frac{\partial L}{\partial Z_4} = Z_3^T \delta Z_4$$

Because  $Z_4 = Z_3W_2, \frac{\partial Z_4}{\partial W_2} = Z_3^T$

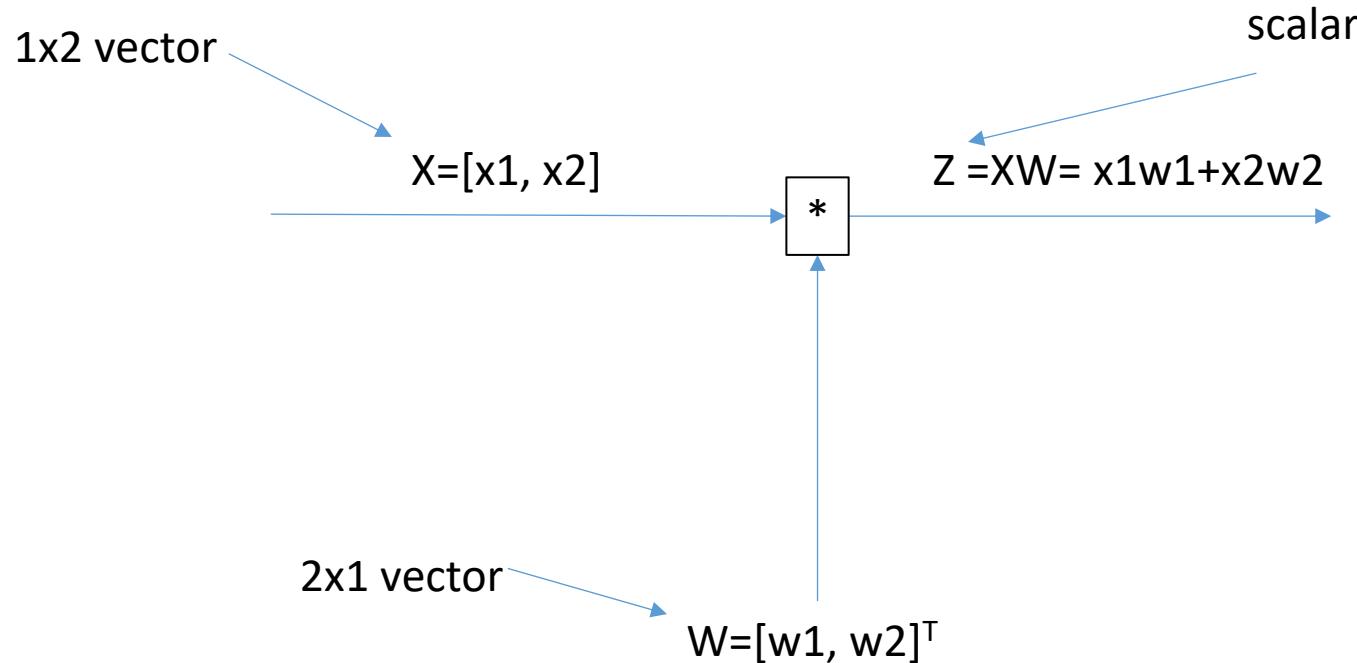
$$2 \quad \delta b_1 \equiv \frac{\partial L}{\partial b_1} = \frac{\partial Z_2}{\partial b_1} \frac{\partial L}{\partial Z_2} = \sum_k (\delta Z_2)_{k,:}$$

Because  $Z_2 = Z_1 + b_1, \frac{\partial Z_2}{\partial b_1} = [1, \dots, 1]$

$$4 \quad \delta b_2 \equiv \frac{\partial L}{\partial b_2} = \frac{\partial Y^p}{\partial b_2} \frac{\partial L}{\partial Y^p} = \sum_k (\delta Y^p)_{k,:}$$

Because  $Y^p = Z_4 + b_2, \frac{\partial Y^p}{\partial b_2} = [1, \dots, 1]$

# Example 1

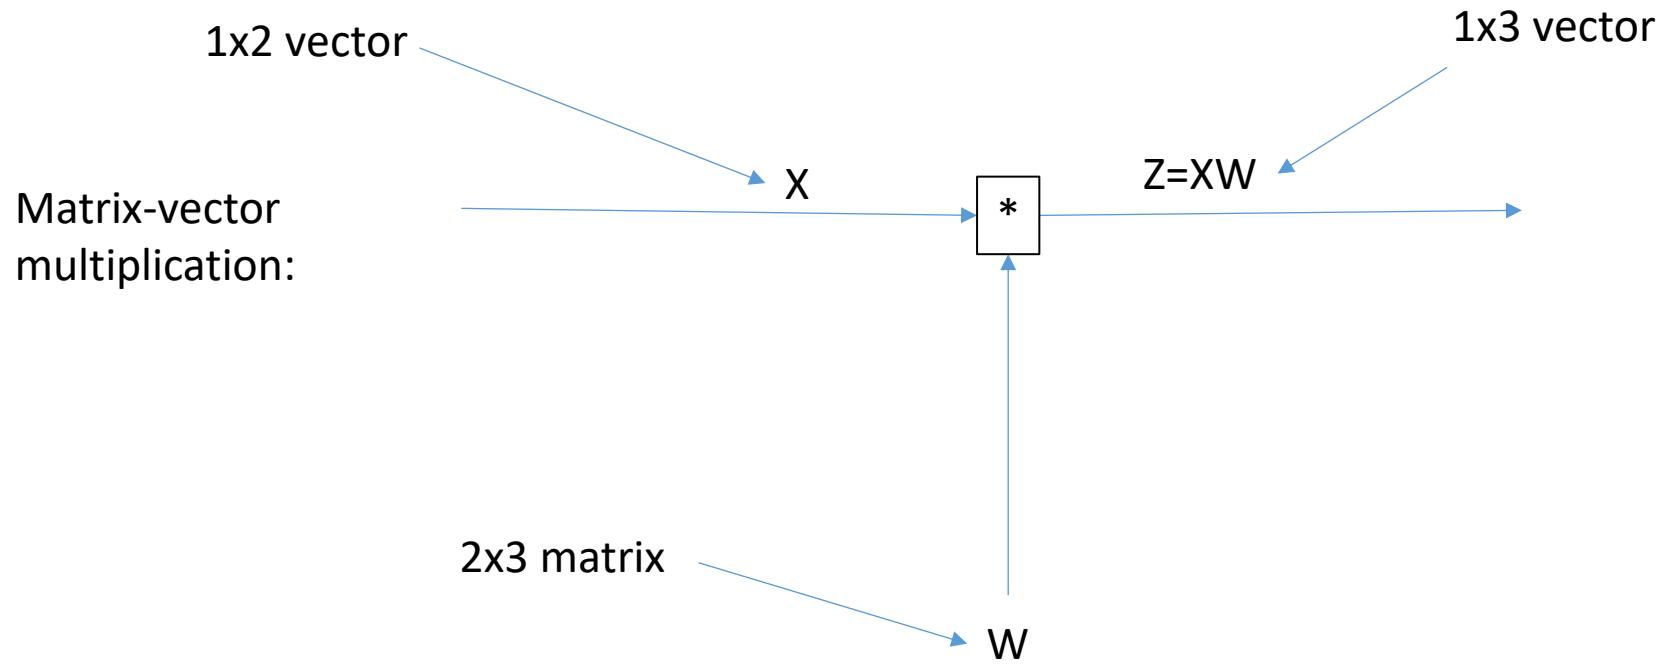


Chain rules:  $\nabla_X(\text{Loss}) = W^T \frac{\partial(\text{Loss})}{\partial Z} = [w_1 \quad w_2] \frac{\partial(\text{Loss})}{\partial Z}$

Why?

$$\nabla_W(\text{Loss}) = X^T \frac{\partial(\text{Loss})}{\partial Z} = [x_1 \quad x_2] \frac{\partial(\text{Loss})}{\partial Z}$$

# Example 2



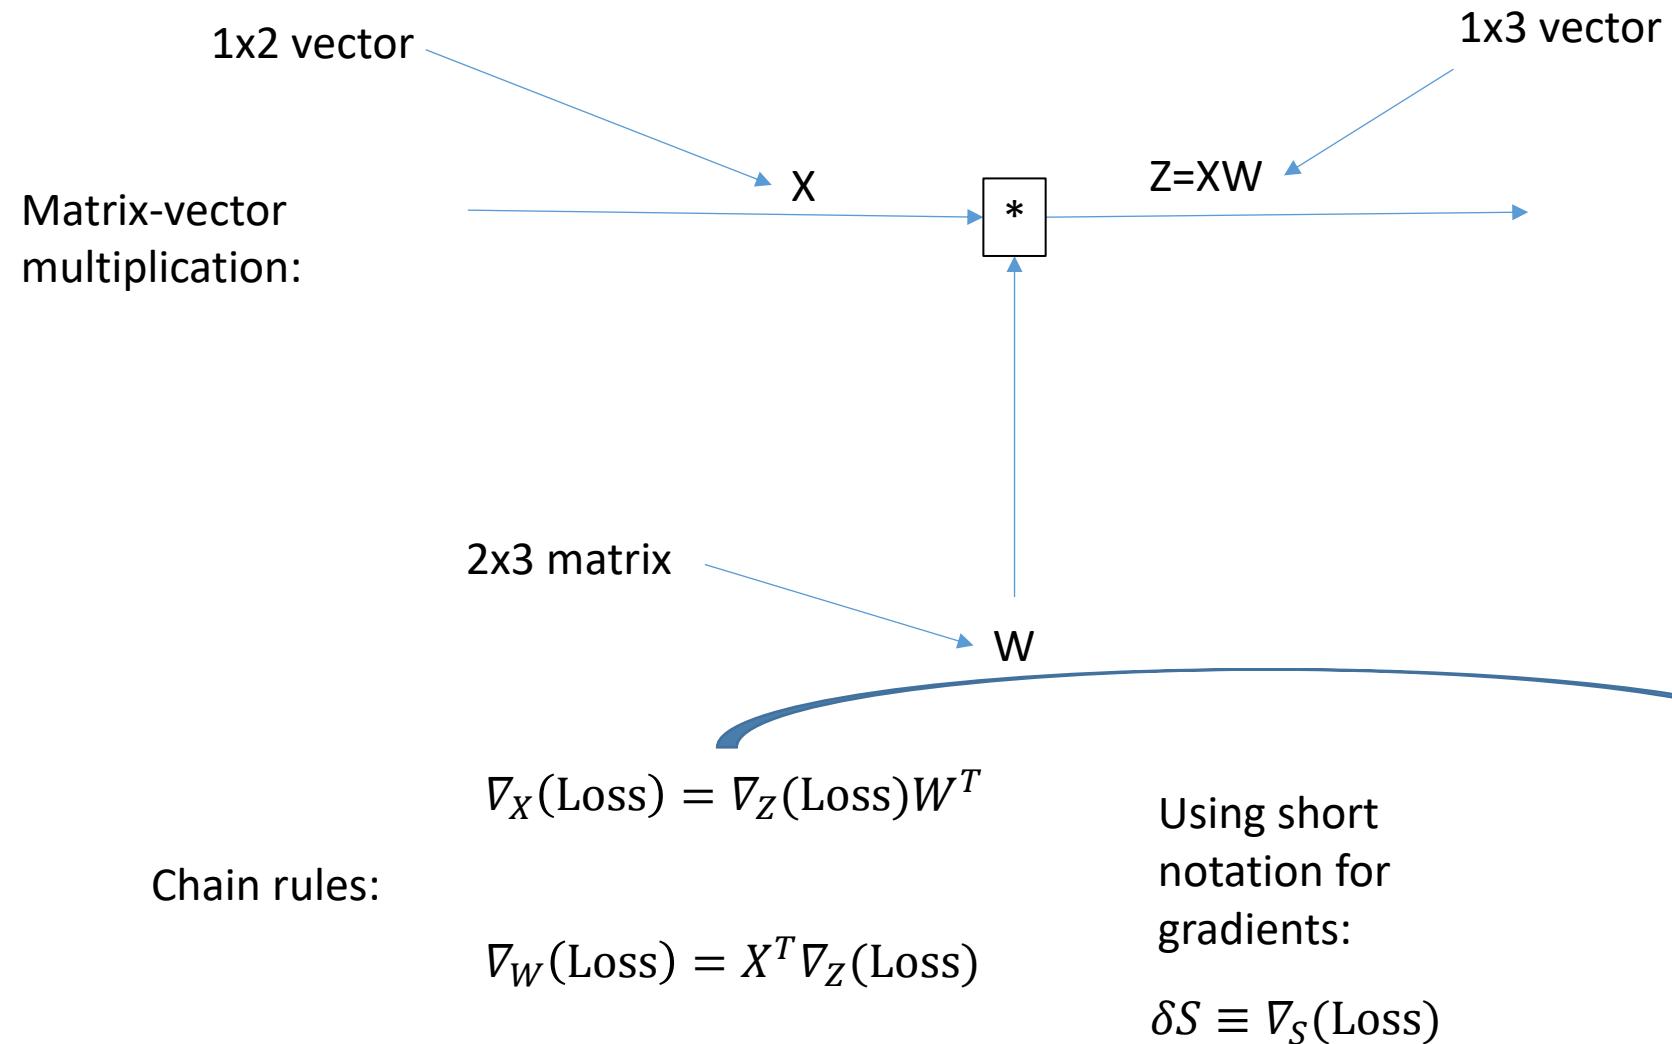
$$\nabla_X(\text{Loss}) = \nabla_Z(\text{Loss})W^T$$

Chain rules:

$$\nabla_W(\text{Loss}) = X^T \nabla_Z(\text{Loss})$$

Why?

# Backprop derivation



# Backprop derivation...

$$\delta X_i = \sum_k \frac{\partial Z_k}{\partial X_i} \frac{\partial (\text{Loss})}{\partial Z_k} = \sum_k \frac{\partial}{\partial X_i} \left[ \sum_j X_j W_{jk} \right] \delta Z_k = \sum_k W_{ik} \delta Z_k$$

➡  $\delta X = \delta Z W^T$

*i<sup>th</sup> component of  $\delta X$  vector*

*Chain rule of derivative*

*Substitute  $Z_k$*

*k<sup>th</sup> component of  $\delta Z$  vector*

Because,

$$\frac{\partial}{\partial X_i} \left[ \sum_j X_j W_{jk} \right] = W_{ik}$$

*Writing in matrix-vector multiplication form*

# Backprop derivation...

$$\delta W_{ij} = \sum_k \frac{\partial Z_k}{\partial W_{ij}} \frac{\partial (\text{Loss})}{\partial Z_k} = \sum_k \frac{\partial}{\partial W_{ij}} \left[ \sum_m X_m W_{mk} \right] \delta Z_k = X_i \delta Z_j \longrightarrow \delta W = X^T \delta Z$$

(i,j)<sup>th</sup> component of  $\delta W$  matrix

Chain rule of derivative

Substitute  $Z_k$

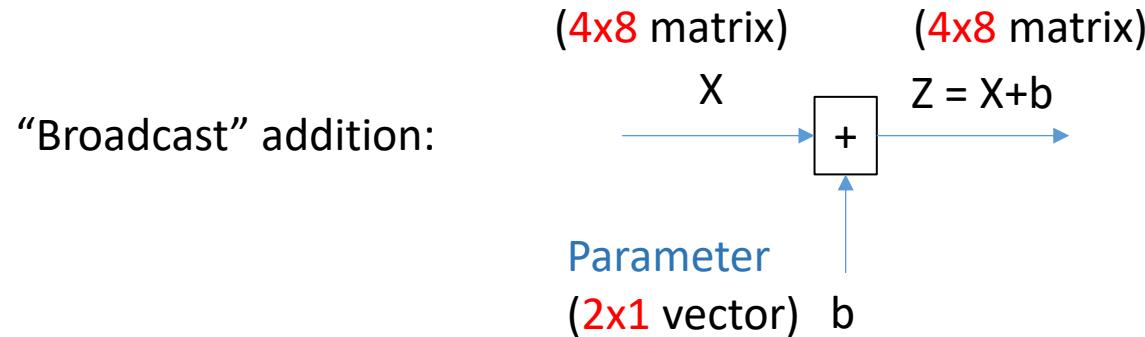
$k^{\text{th}}$  component of  $\delta Z$  vector

Because,

$$\frac{\partial}{\partial W_{ij}} \left[ \sum_m X_m W_{mk} \right] = \begin{cases} X_i, & \text{if } i = m \text{ and } j = k, \\ 0, & \text{otherwise.} \end{cases}$$

Writing in matrix-vector multiplication form

# Backprop derivation...



$$\delta X_{i,j} = \sum_k \sum_l \frac{\partial Z_{k,l}}{\partial X_{i,j}} \delta Z_{k,l} = \sum_k \sum_l \frac{\partial}{\partial X_{i,j}} [X_{k,l} + b_k] \delta Z_{k,l} = \delta Z_{i,j}$$

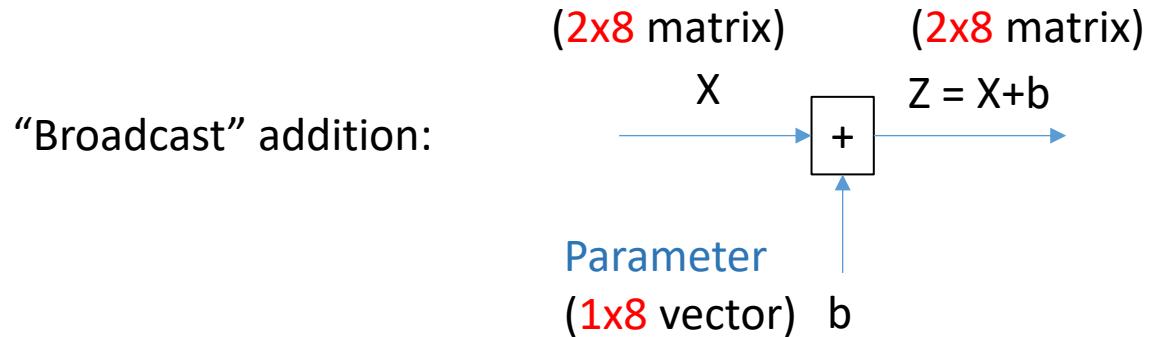
Chain rule

Substitute  $Z_{k,l}$

Because,

$$\frac{\partial}{\partial X_{i,j}} [X_{k,l} + b_k] = \begin{cases} 1, & \text{if } i = k \text{ and } j = l, \\ 0, & \text{otherwise.} \end{cases}$$

# Backprop derivation for broadcast addition



$$\delta b_i = \sum_k \sum_l \frac{\partial Z_{k,l}}{\partial b_i} \delta Z_{k,l} = \sum_k \sum_l \frac{\partial}{\partial b_i} [X_{k,l} + b_l] \delta Z_{k,l} = \sum_k \delta Z_{k,i}$$

Because,

$$\frac{\partial}{\partial b_i} [X_{k,l} + b_l] = \begin{cases} 1, & \text{if } i = l, \\ 0, & \text{otherwise.} \end{cases}$$

# Backprop derivation for activation function

Non-linear function:  
(applied **pointwise**)

The diagram shows a blue arrow pointing from input  $X$  to a square box labeled  $\sigma$ . Another blue arrow points from the box to the output  $Z = \sigma(X)$ .

Using chain rule:

$$\delta X_{i,j} = \frac{dZ_{i,j}}{dX_{i,j}} \delta Z_{i,j} = \frac{d\sigma(X_{i,j})}{dX_{i,j}} \delta Z_{i,j}$$

If the non-linear  
function is sigmoid,

$$\sigma(a) = \frac{1}{1 + \exp(-a)}$$

$$\frac{d\sigma}{da} = \frac{\exp(-a)}{(1 + \exp(-a))^2} = \frac{1}{1 + \exp(-a)} \left(1 - \frac{1}{1 + \exp(-a)}\right) = \sigma(a)(1 - \sigma(a))$$

$$\delta X_{i,j} = \sigma(X_{i,j})(1 - \sigma(X_{i,j}))\delta Z_{i,j}$$

# Backprop derivation for loss function

Euclidean loss function:

$$Loss(Y^p, Y) = \frac{1}{2} \|Y^p - Y\|^2 = \frac{1}{2} \sum_i (Y_i^p - Y_i)^2$$

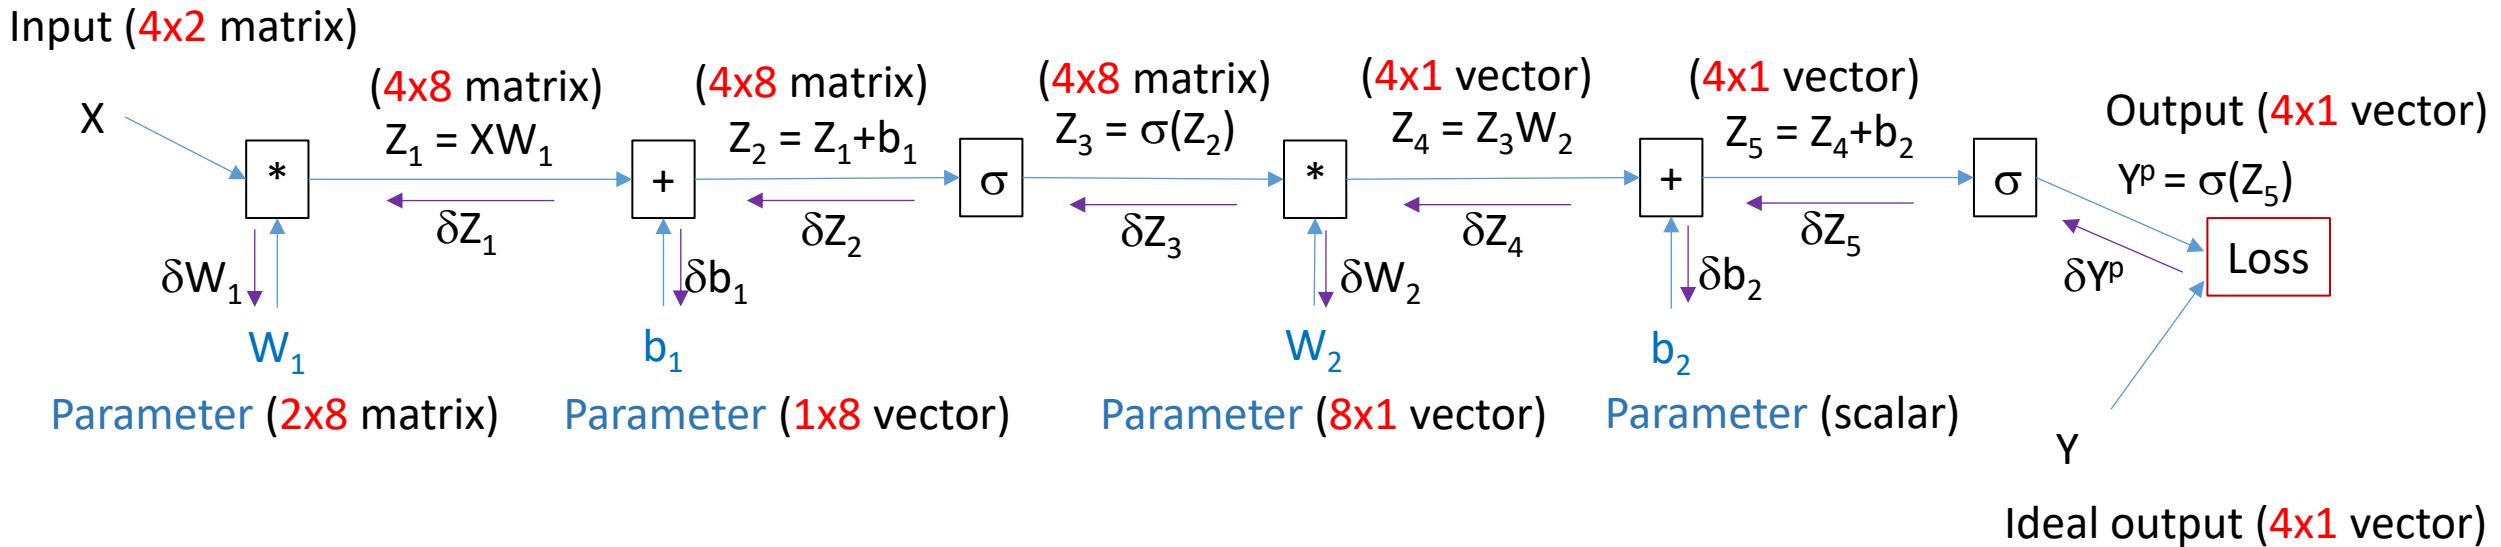
$i^{\text{th}}$  component of  $\delta Y^p$  vector:

$$\delta Y_i^p = \frac{\partial}{\partial Y_i^p} Loss(Y^p, Y) = \frac{\partial}{\partial Y_i^p} \frac{1}{2} \sum_k (Y_k^p - Y_k)^2 = Y_i^p - Y_i$$

Using vector notation:

$$\delta Y^p = Y^p - Y$$

# Apply chain rule to XOR neural network



Chain rule of derivatives:

$$\delta Y^p = Y^p - Y$$

$$\delta Z_5 = \sigma(Z_5) \cdot (1 - \sigma(Z_5)) \cdot \delta Y^p$$

$$\delta Z_4 = \delta Z_5$$

$$\delta Z_3 = \delta Z_4 W_2^T$$

$$\delta Z_2 = \sigma(Z_2) \cdot (1 - \sigma(Z_2)) \cdot \delta Z_3$$

$$\delta Z_1 = \delta Z_2$$

New notation:  
 $\delta S \equiv \nabla_S(\text{Loss})$

Gradient of “Loss” with respect to input signals  
 ↓  
 Propagates backward

$$\delta W_2 = Z_3^T \delta Z_4$$

$$\delta b_2 = \sum_k (\delta Z_5)_k$$

$$\delta W_1 = X^T \delta Z_1$$

$$\delta b_1 = \sum_k (\delta Z_2)_{k,:}$$

Gradient of “Loss” with respect to parameters

# Backprop to train a neural net

Initialize all parameters of the neural network

Initialize learning rate variable  $lr$

Iterate:

(Load Data): Get training data batch

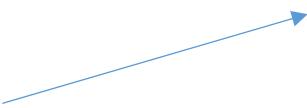
(Forward pass): Compute  $Z_1, Z_2, \dots, Y^p$

(Backward pass): Compute gradients  $\delta Y^p, \delta Z_5, \dots, \delta Z_1, \delta W_2, \delta W_1, \delta b_2, \delta b_1$

(Gradient descent to update parameters):  $W_2 \leftarrow W_2 - lr * \delta W_2, \quad b_2 \leftarrow b_2 - lr * \delta b_2, \dots,$

(Diagnostics): Compute “Loss” from time to time to check if it is decreasing

If loading the whole training data, do it **once** outside the “Iterate” loop, to be efficient



# MNIST classification problem



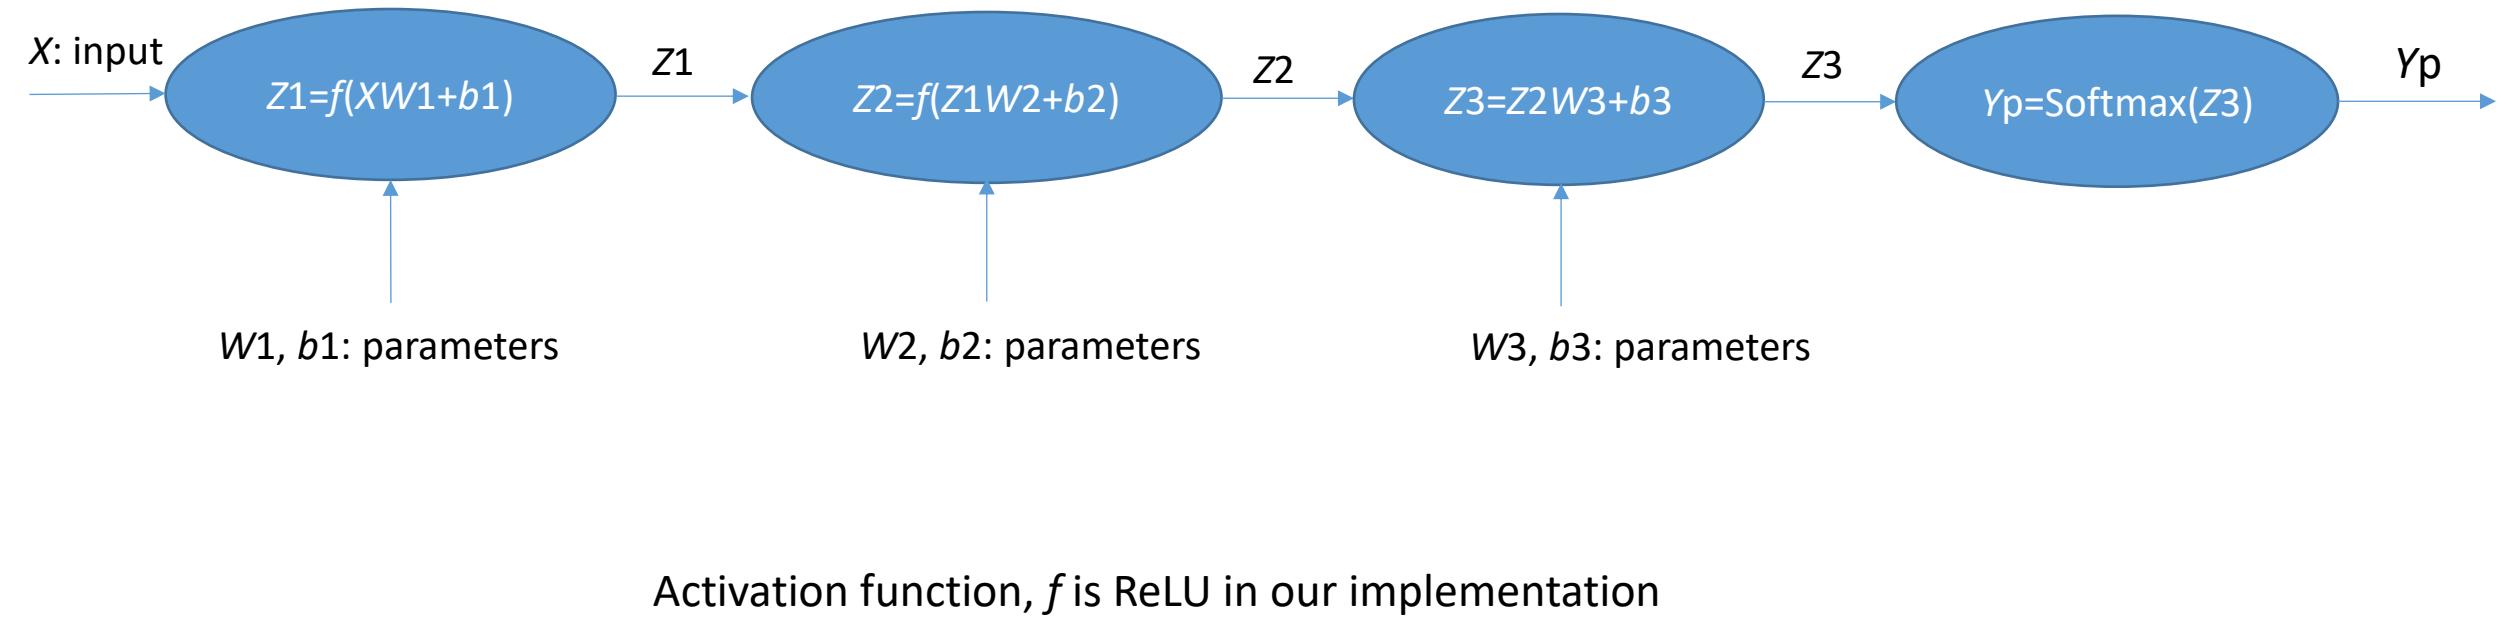
Small 28 pixels-by-28 pixels images of hand written digits

The visual recognition problem definition:  
to recognize the digit from an image

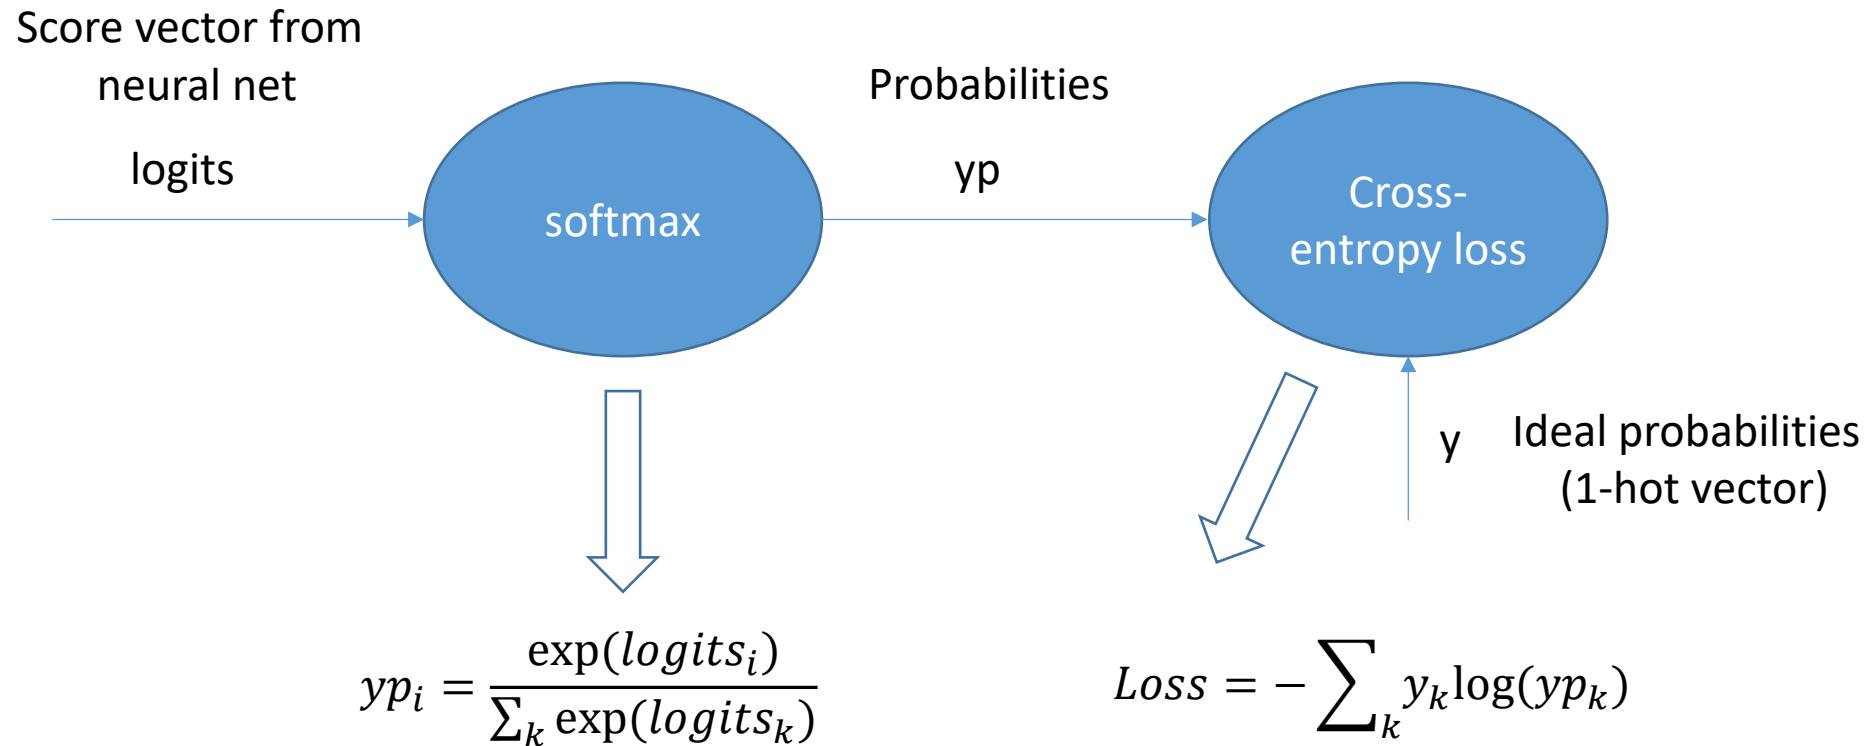


$x_1$	$x_2$	...	$x_{784}$	$y_1$	...	$y_{10}$
0.1	0.3	...	0.0	0	...	1
0.2	0.1	...	0.5	1	...	0
...	...	...	...	...	...	...
...	...	...	...	...	...	...
0.0	0.98	...	0.8	0	...	1
0.5	0.25	...	0.36	?	...	?
0.1	0.95	...	0.1	?	...	?

# NN Architecture for MNIST Classification



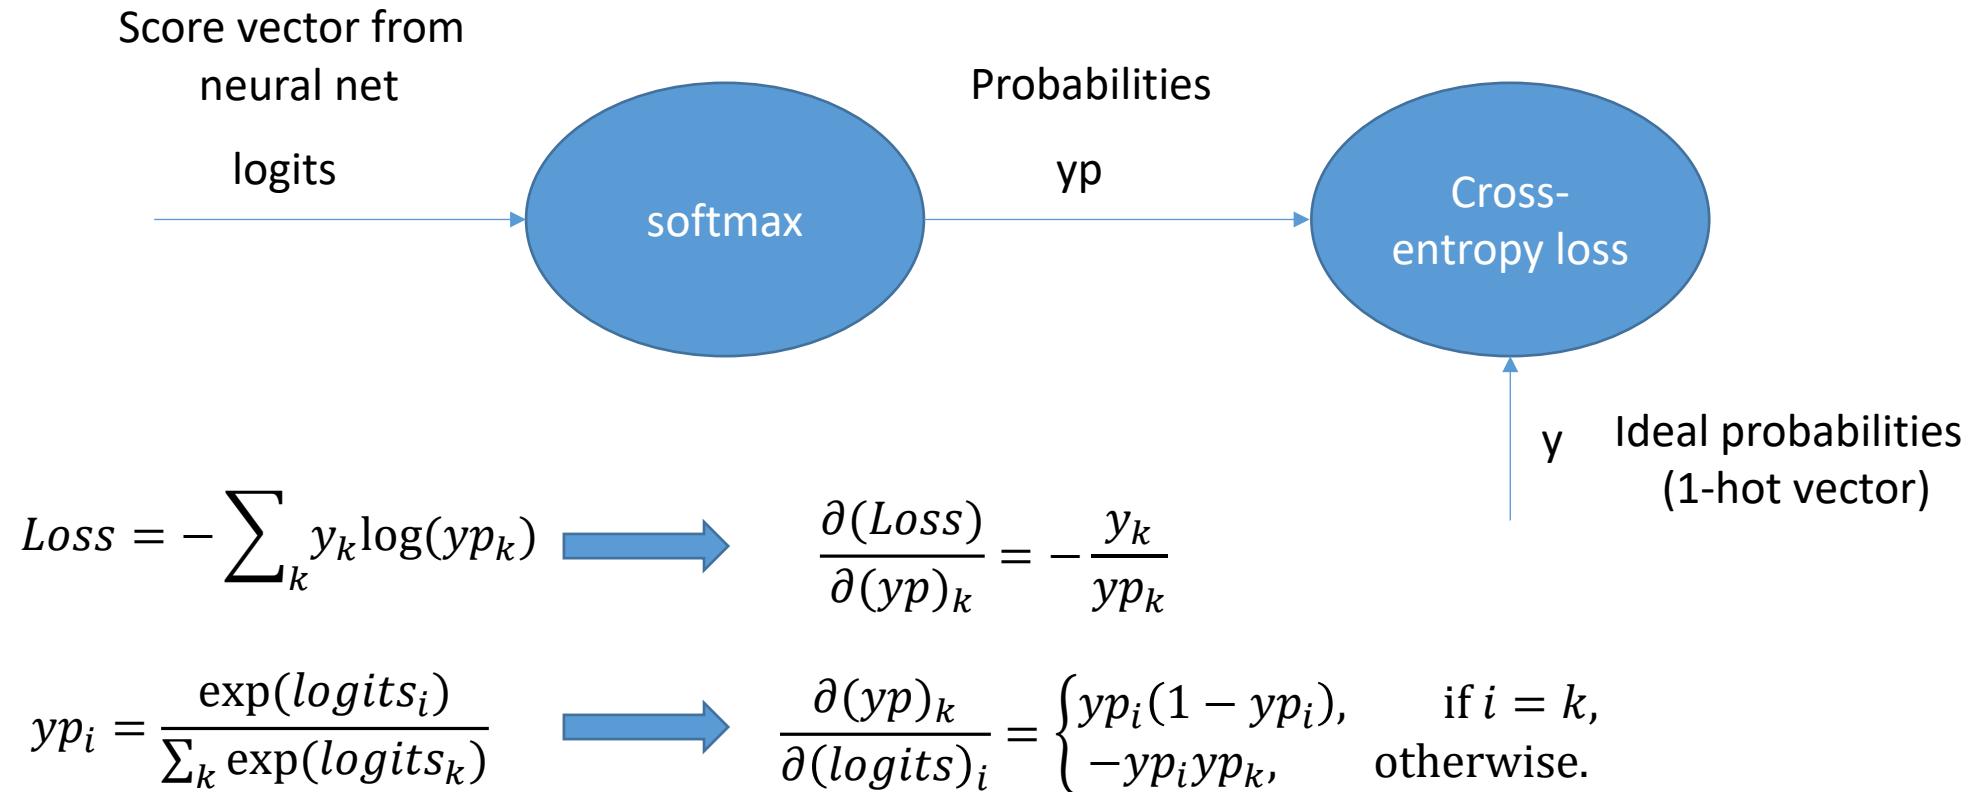
# Softmax and cross-entropy loss



To backpropagate error, we need to compute:

$$\delta(\text{logits})_i \equiv \frac{\partial(\text{Loss})}{\partial(\text{logits})_i}$$

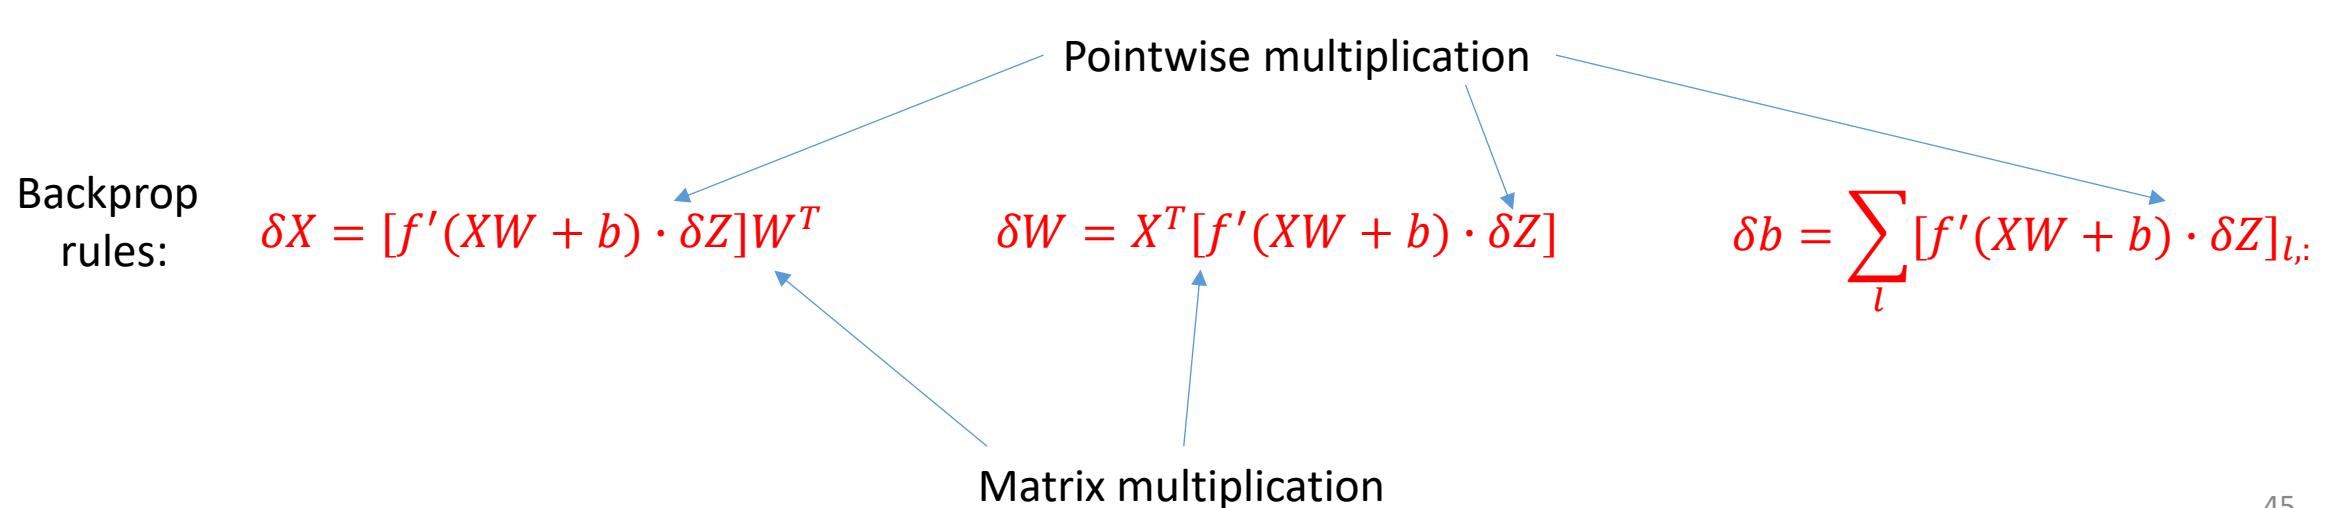
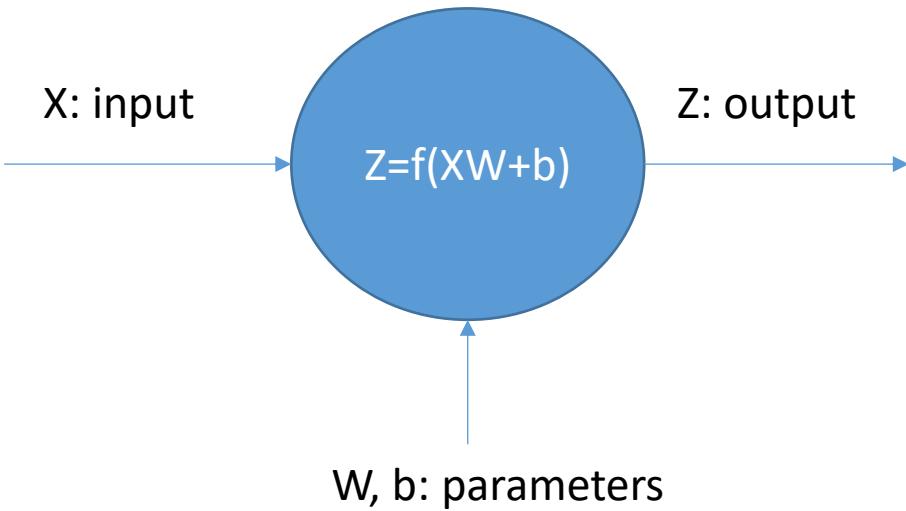
# Softmax and cross-entropy loss: backprop



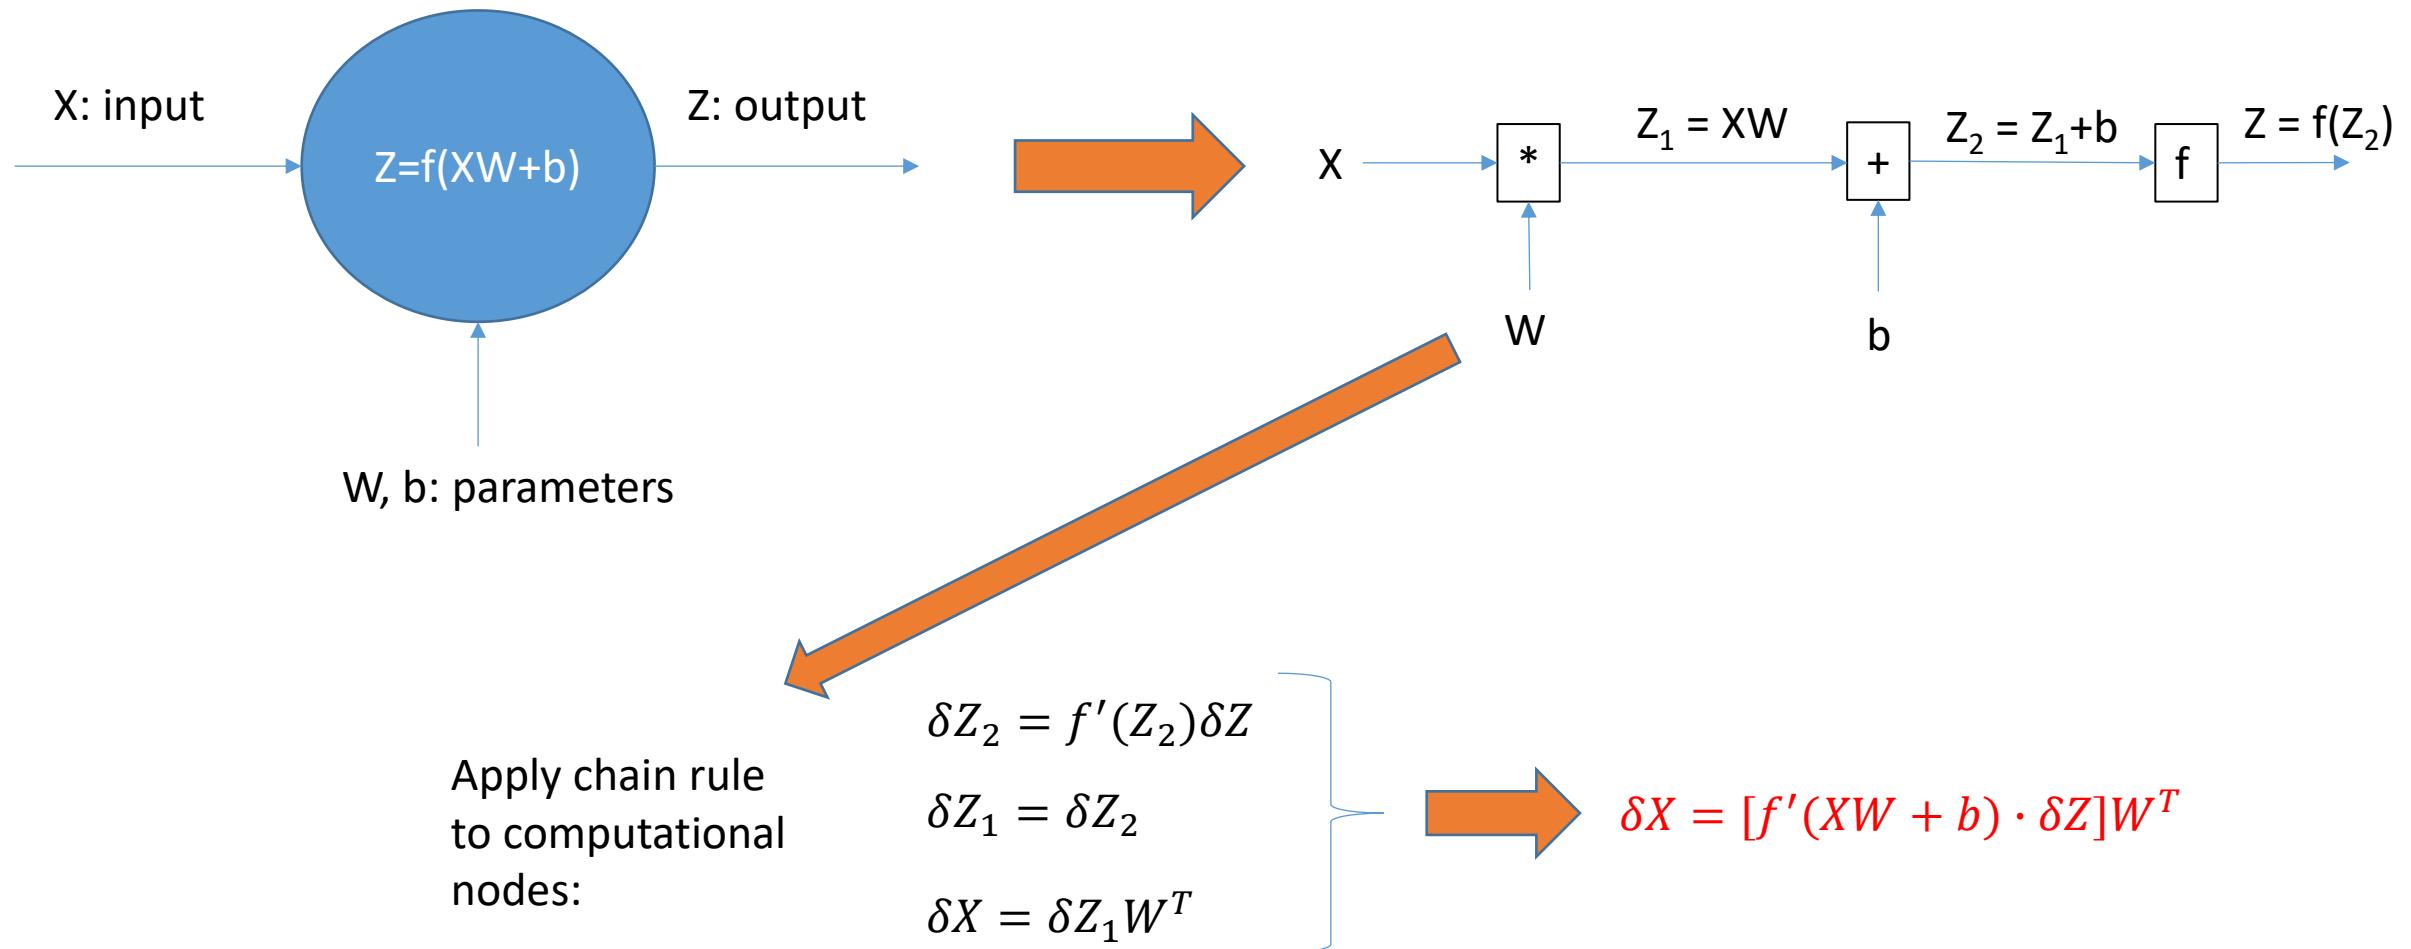
Using the above two results in the chain rule,  $\delta(logits)_i \equiv \frac{\partial(Loss)}{\partial(logits)_i} = \sum_k \frac{\partial(yp)_k}{\partial(logits)_i} \frac{\partial(Loss)}{\partial(yp)_k} = \color{red}{yp_i - y_i}$

What if, instead of cross-entropy, we used L2 loss along with softmax?

# Backprop across a neural net layer

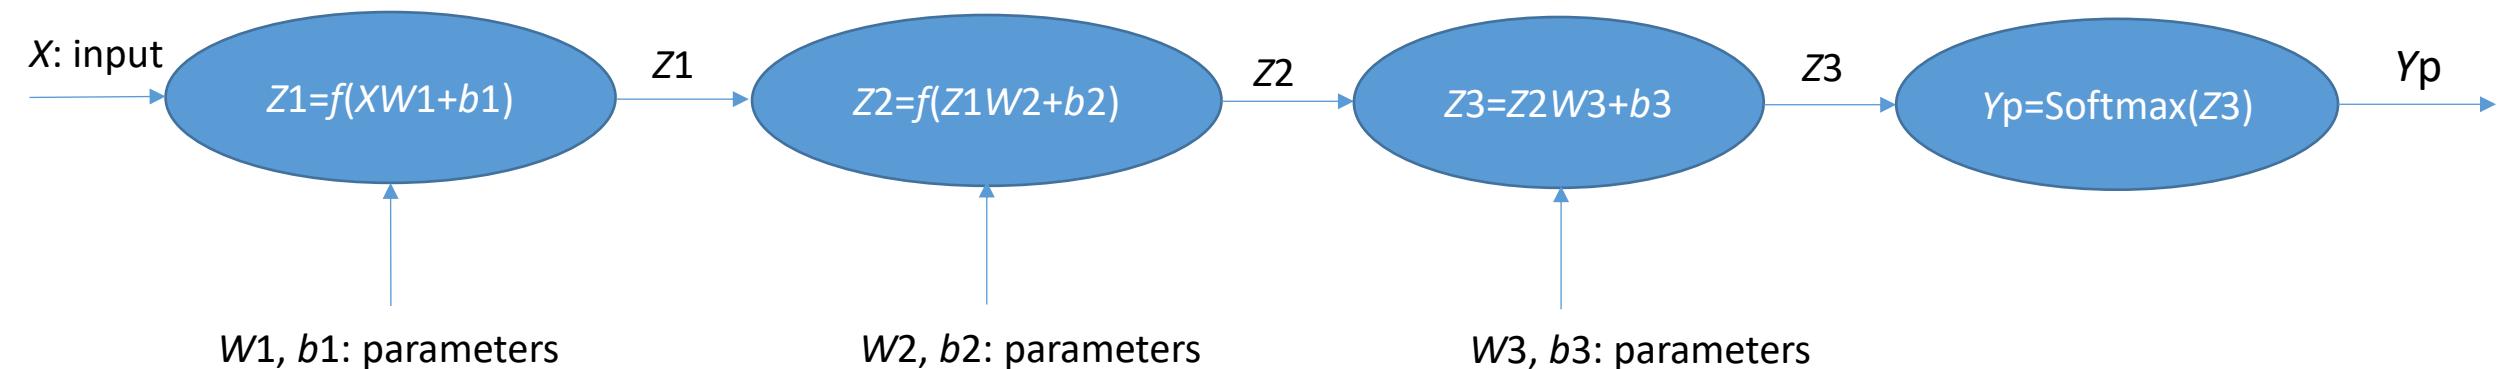


# Backprop across a neural net layer: derivation



Similarly, we can derive backprop rules for  $\delta W$  and  $\delta b$

# NN for MNIST Classification: Gradients and Manual Backprop



Backprop:

$$\delta Z_3 = \gamma_p - Y$$

$$\delta W_3 = (Z_2^T) \delta Z_3$$

$$\delta b_3 = \sum_l [\delta Z_3]_{l,:}$$

$$\delta Z_2 = (\delta Z_3) W_3^T$$

$$\delta W_2 = Z_1^T [f'(Z_1W_2 + b_2) \cdot \delta Z_2]$$

$$\delta b_2 = \sum_l [f'(Z_1W_2 + b_2) \cdot \delta Z_2]_{l,:}$$

$$\delta Z_1 = [f'(Z_1W_2 + b_2) \cdot \delta Z_2] W_2^T$$

$$\delta W_1 = X^T [f'(XW_1 + b_1) \cdot \delta Z_1]$$

$$\delta b_1 = \sum_l [f'(XW_1 + b_1) \cdot \delta Z_1]_{l,:}$$