

# Vision Transformers

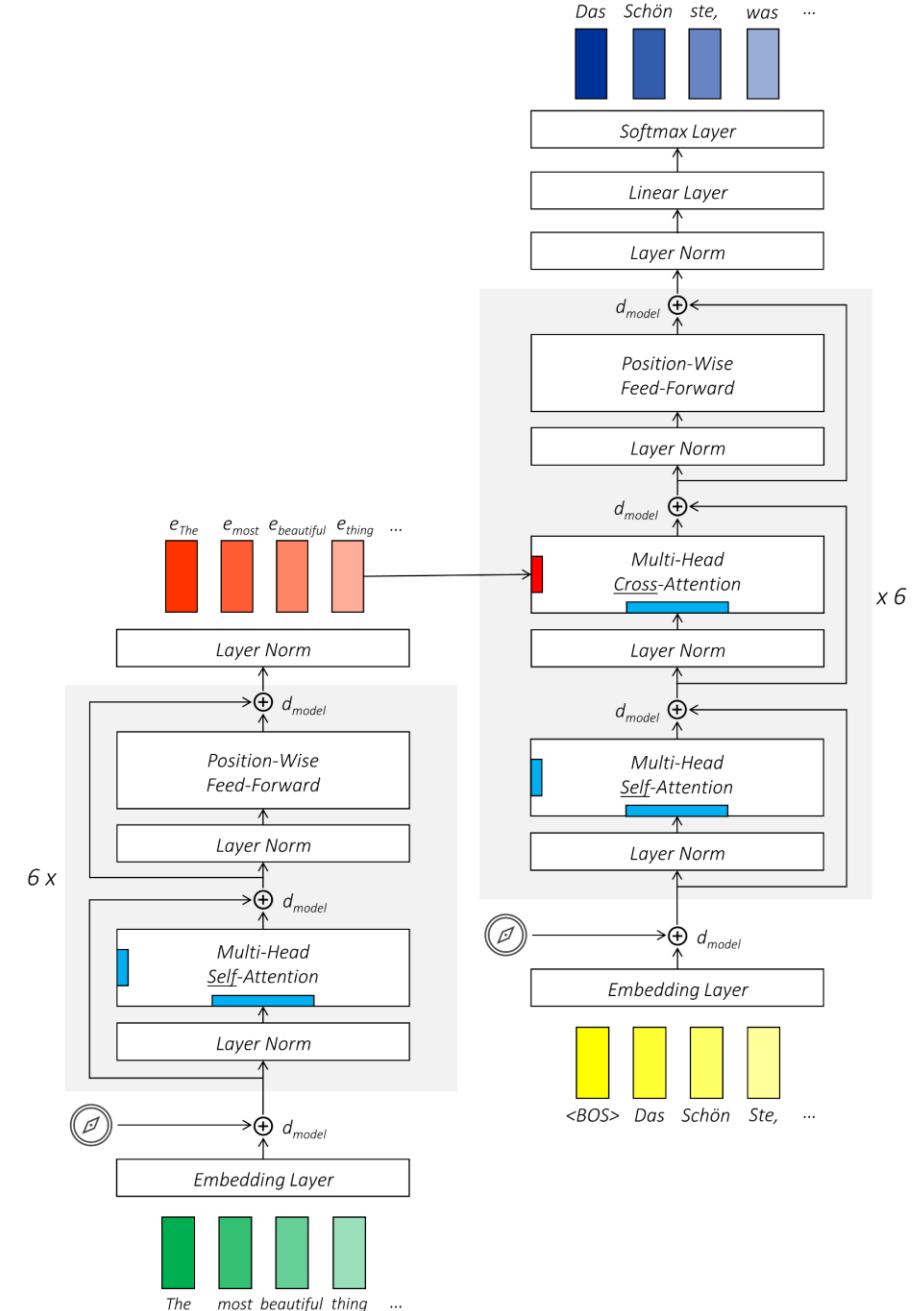
Nilanjan Ray

# Plan

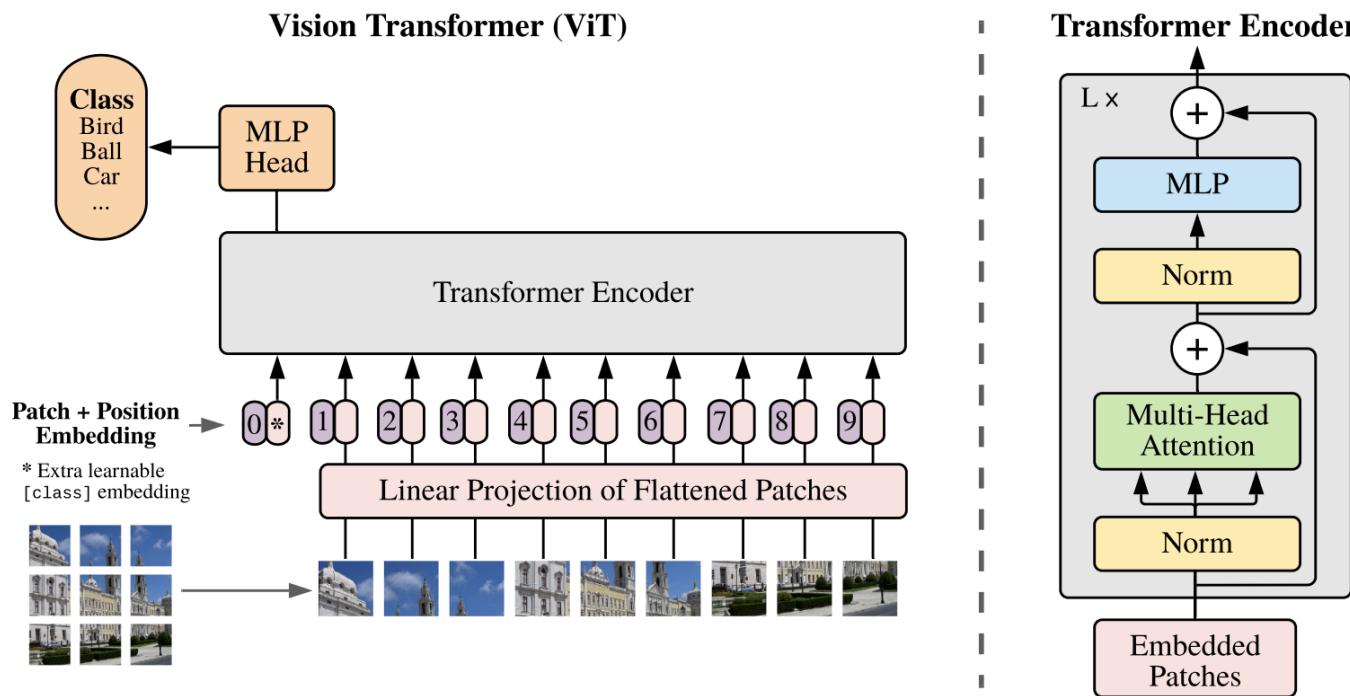
- To understand vision transformers, we need to understand how a transformer works
- There are plenty of tutorials out there. Let's use this one:  
<https://github.com/sgrvinod/a-PyTorch-Tutorial-to-Machine-Translation>
- We will take about 15 mins to go through this page. Then, we will come back to describe a vision transformer (ViT) for image classification.
- After studying ViT, we will study how image captions can be generated using transformers.

# Transformer: Encoder and Decoder

<https://github.com/sgrvinod/a-PyTorch-Tutorial-to-Transformers?tab=readme-ov-file>



# What are ViTs?



Only encoders are used in ViT,  
no decoders are needed  
for image classification

Tutorial: [https://uvadlc-notebooks.readthedocs.io/en/latest/tutorial\\_notebooks/tutorial15/Vision\\_Transformer.html](https://uvadlc-notebooks.readthedocs.io/en/latest/tutorial_notebooks/tutorial15/Vision_Transformer.html)

# Let's dissect “VisionTransformer” module

- Image shape
  - `print(CIFAR_images.shape)`
  - `torch.Size([4, 3, 32, 32])`: batch size, #channels, image height, image width
- Image to patches:
  - `img_patches = img_to_patch(CIFAR_images, patch_size=4, flatten_channels=True)`
  - `print(img_patches.shape)`
  - `torch.Size([4, 64, 48])`: batch size, #patches, (#channels)x(patch height)x(patch width)
- Preprocess patches by linear layer
  - Input shape to linear layer: [4, 64, 48]
  - Output shape from the linear layer: [4, 64, 256]. 256 is the embedding dimension
- CLS tokens shape: [1, 1, 256] but it is repeated for batches, so shape becomes [4, 1, 256]
- Shape after CLS tokens concatenated to output of linear embedding: [4, 65, 256].
- This tensor is added to the positional embedding.
- Positional embedding shape: [1, 65, 256]
- Transformer input and output - note that 0<sup>th</sup> and 1<sup>st</sup> dimensions are exchanged
  - Input shape to transformer: [65, 4, 256]
  - Output shape from transformer: [65, 4, 256]
- MLP head input and output
  - Input shape to MLP head: [4, 256]
  - Output shape from MLP head: [4, 10]

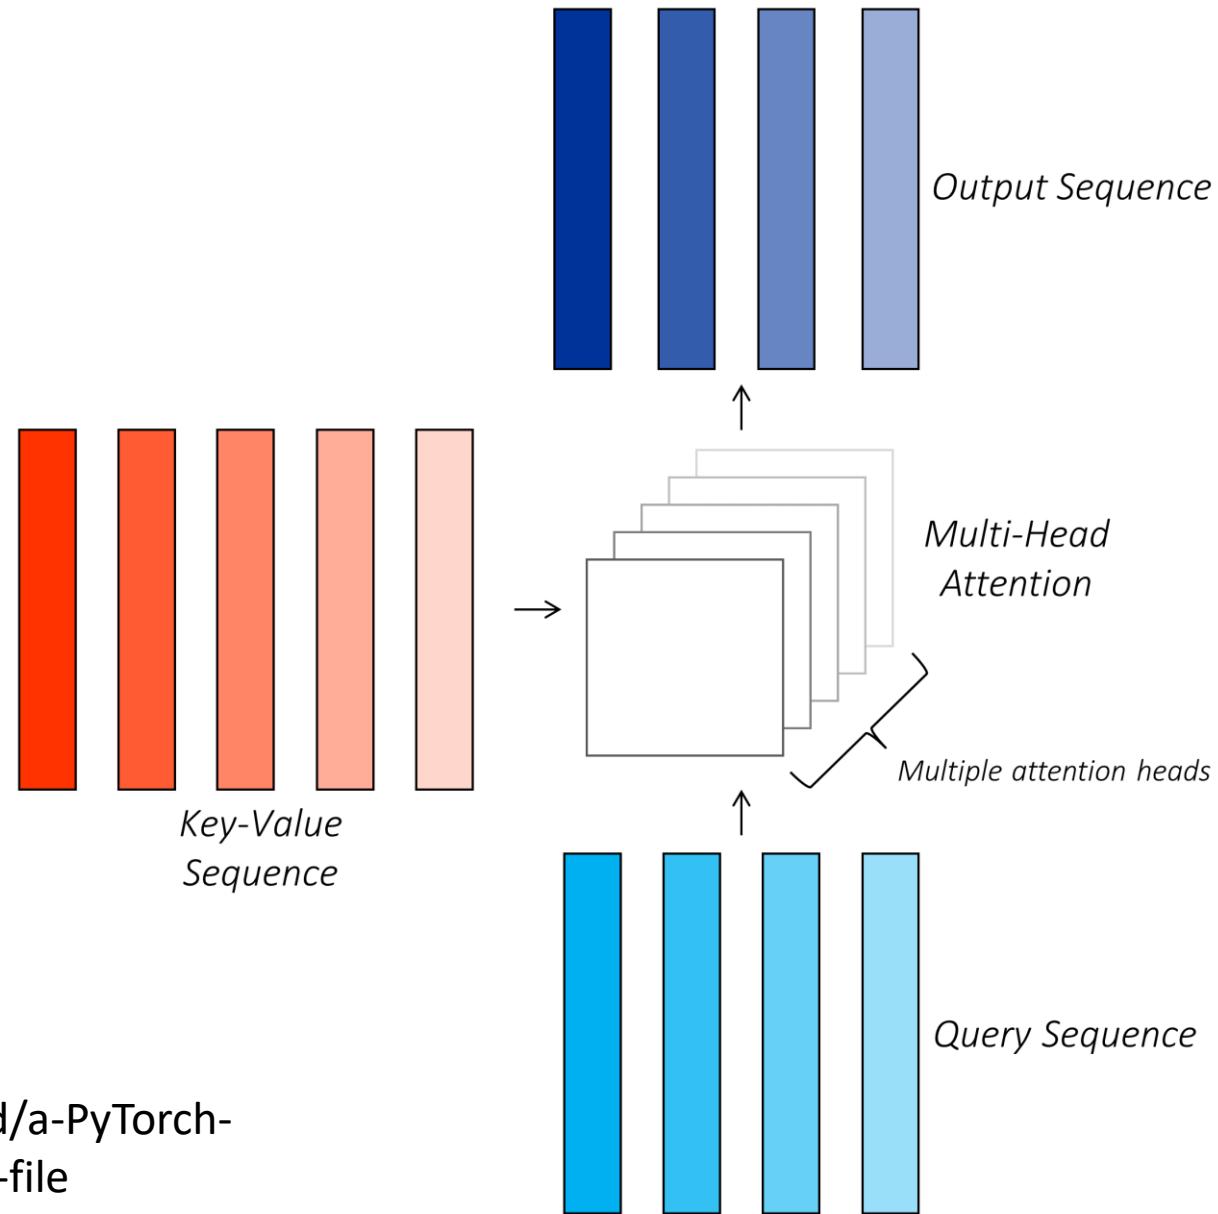
```
def forward(self, x):  
    # Preprocess input  
    x = img_to_patch(x, self.patch_size)  
    B, T, _ = x.shape  
    x = self.input_layer(x)  
  
    # Add CLS token and positional encoding  
    cls_token = self.cls_token.repeat(B, 1, 1)  
    x = torch.cat([cls_token, x], dim=1)  
    x = x + self.pos_embedding[:, :T+1]  
  
    # Apply Transformer  
    x = self.dropout(x)  
    x = x.transpose(0, 1)  
    x = self.transformer(x)  
  
    # Perform classification prediction  
    cls = x[0]  
    out = self.mlp_head(cls)  
    return out
```

# Self.transformer: Sequentially applies a few attention blocks

- Input shape to the Attn block: [65, 4, 256]
- Output shape from the Attn block: [65, 4, 256]
- Residual connections, linear layers, layer normalizations only work on the last dimensions of the tensor [65, 4, 256].
- So where is the mixing happening?
  - It must be happening inside the nn.MultiheadAttention module

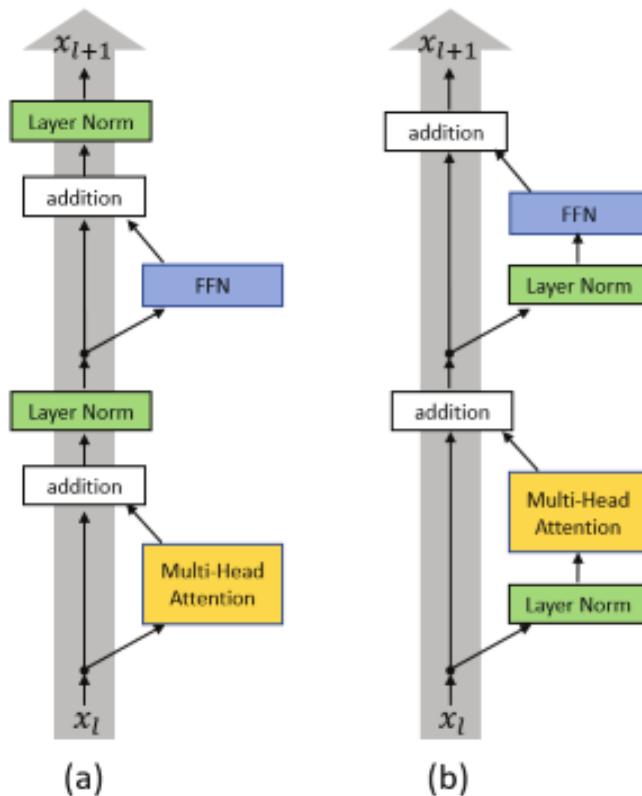
```
class AttentionBlock(nn.Module):  
  
    def __init__(self, embed_dim, hidden_dim, num_heads, dropout=0.0):  
        super().__init__()  
  
        self.layer_norm_1 = nn.LayerNorm(embed_dim)  
        self.attn = nn.MultiheadAttention(embed_dim, num_heads,  
                                         dropout=dropout)  
        self.layer_norm_2 = nn.LayerNorm(embed_dim)  
        self.linear = nn.Sequential(  
            nn.Linear(embed_dim, hidden_dim),  
            nn.GELU(),  
            nn.Dropout(dropout),  
            nn.Linear(hidden_dim, embed_dim),  
            nn.Dropout(dropout)  
        )  
  
    def forward(self, x):  
        print("Input shape to Attn block:", x.shape)  
        inp_x = self.layer_norm_1(x)  
        x = x + self.attn(inp_x, inp_x, inp_x)[0]  
        x = x + self.linear(self.layer_norm_2(x))  
        return x
```

# Multi-head Attention



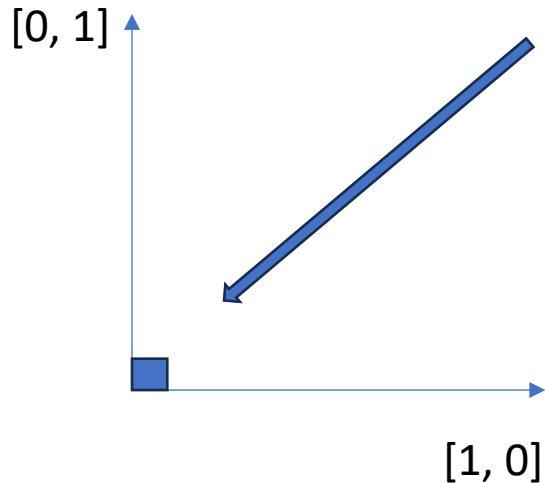
Details here: <https://github.com/sgrvinod/a-PyTorch-Tutorial-to-Transformers?tab=readme-ov-file>

# Transformer encoder



(a) Post-LN Transformer layer; (b) Pre-LN Transformer layer.

# Vectors and their relationships



The relationship between 2 vectors is represented by their angle

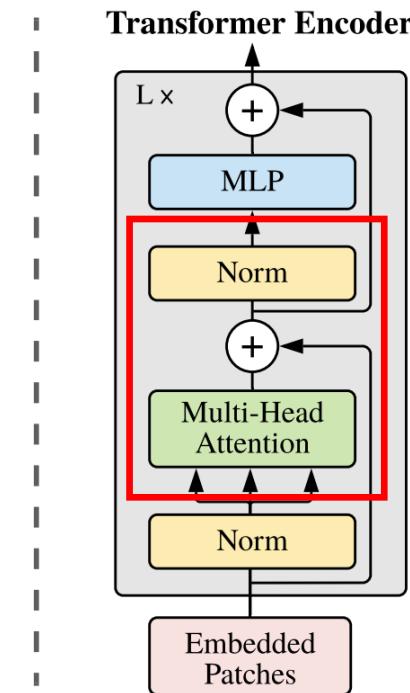
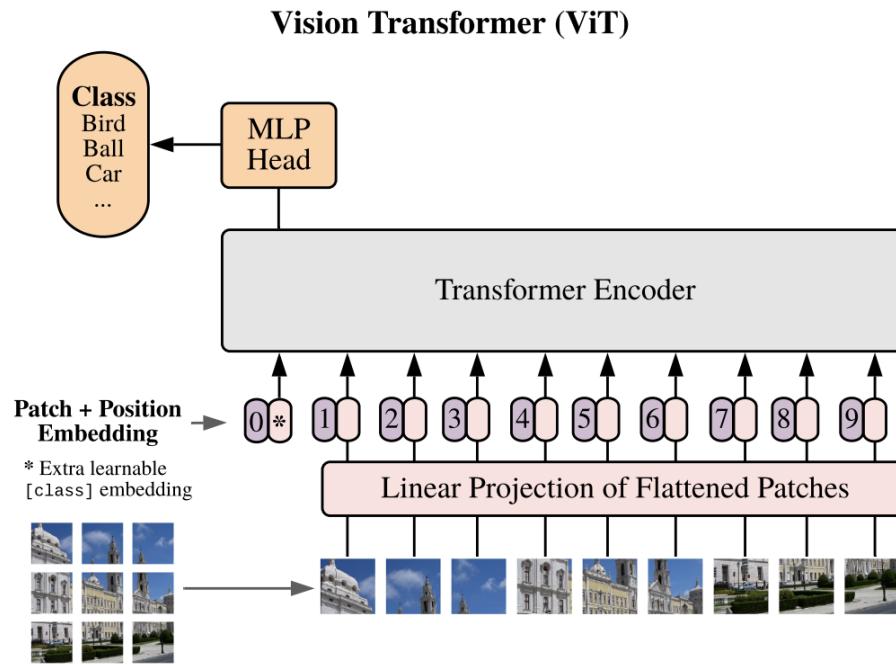
Consine similarity formula

$$\cos \theta = \frac{\vec{a} \cdot \vec{b}}{\|\vec{a}\| \|\vec{b}\|} \quad \frac{\text{dot product}}{\text{Norms}}$$

$$\|\vec{a}\| = \sqrt{a_1^2 + a_2^2 + a_3^2 + \dots + a_n^2}$$

$$\|\vec{b}\| = \sqrt{b_1^2 + b_2^2 + b_3^2 + \dots + b_n^2}$$

# Attention



*dot product*  
*Norms*

# Attention

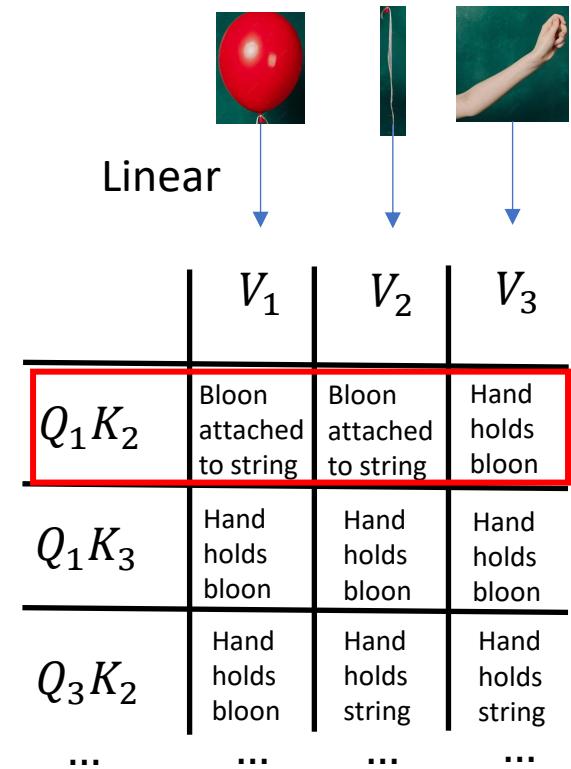
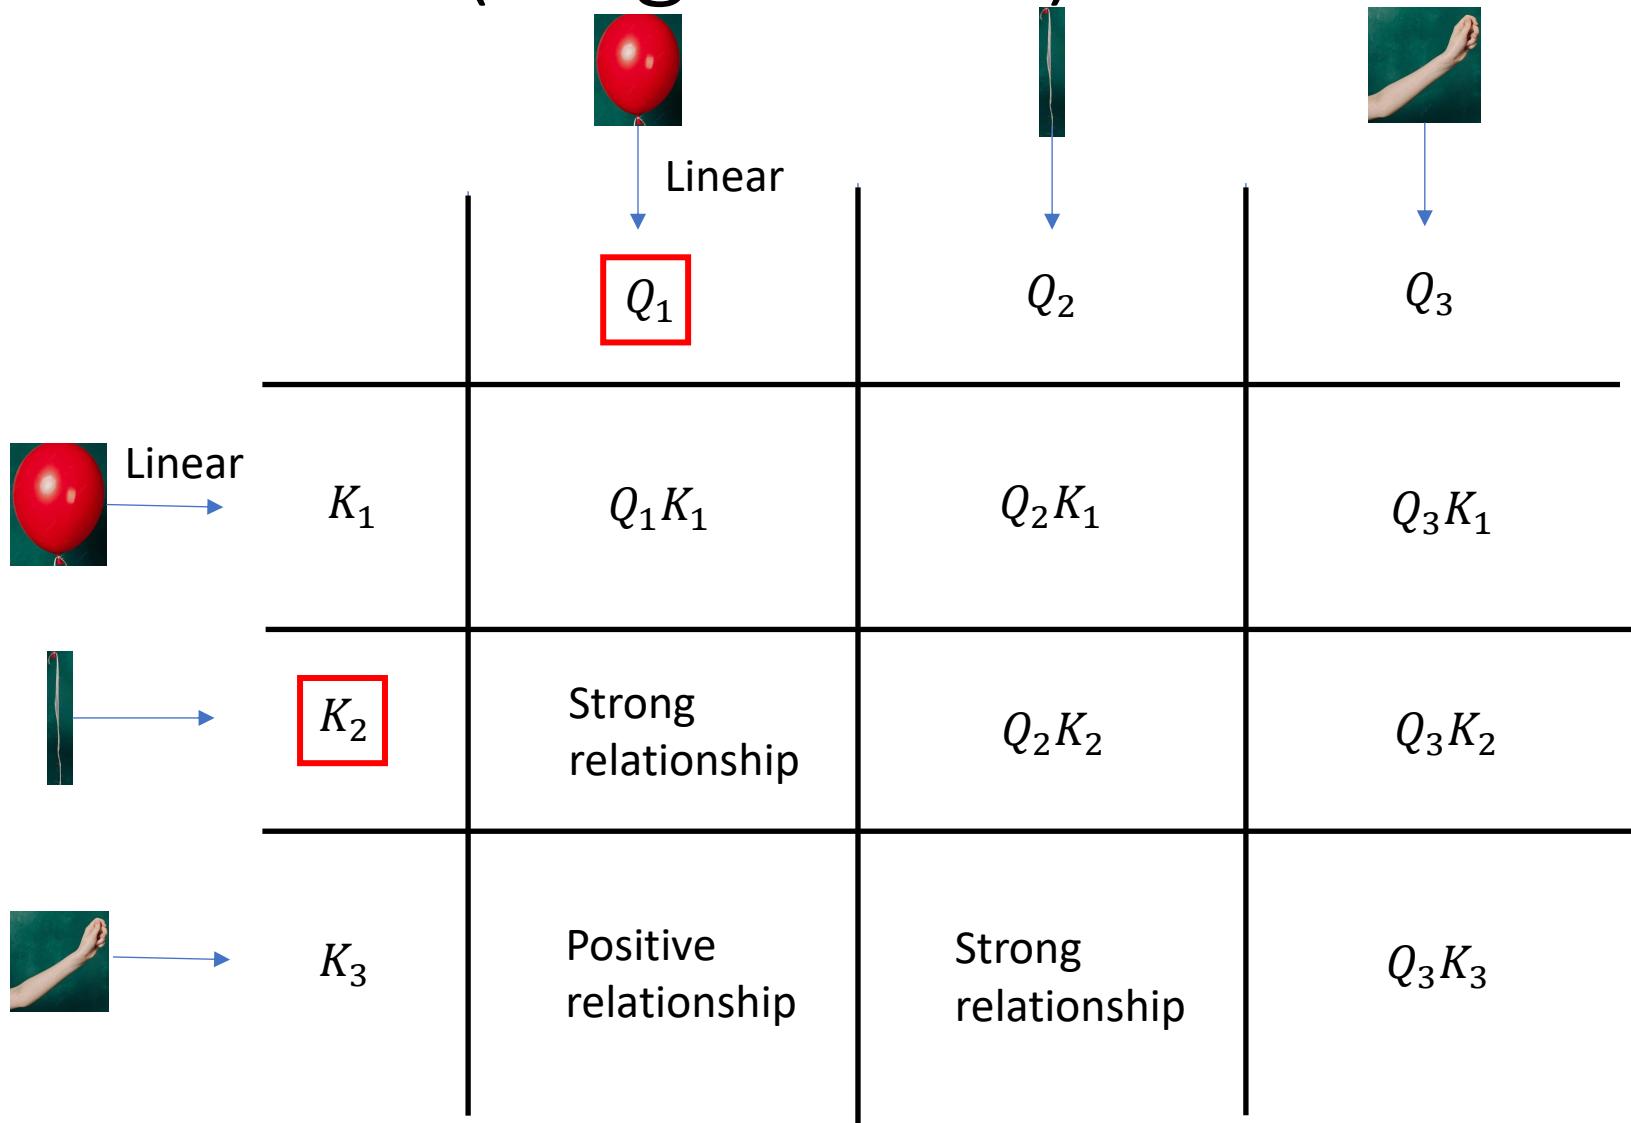


Query  
(Question)

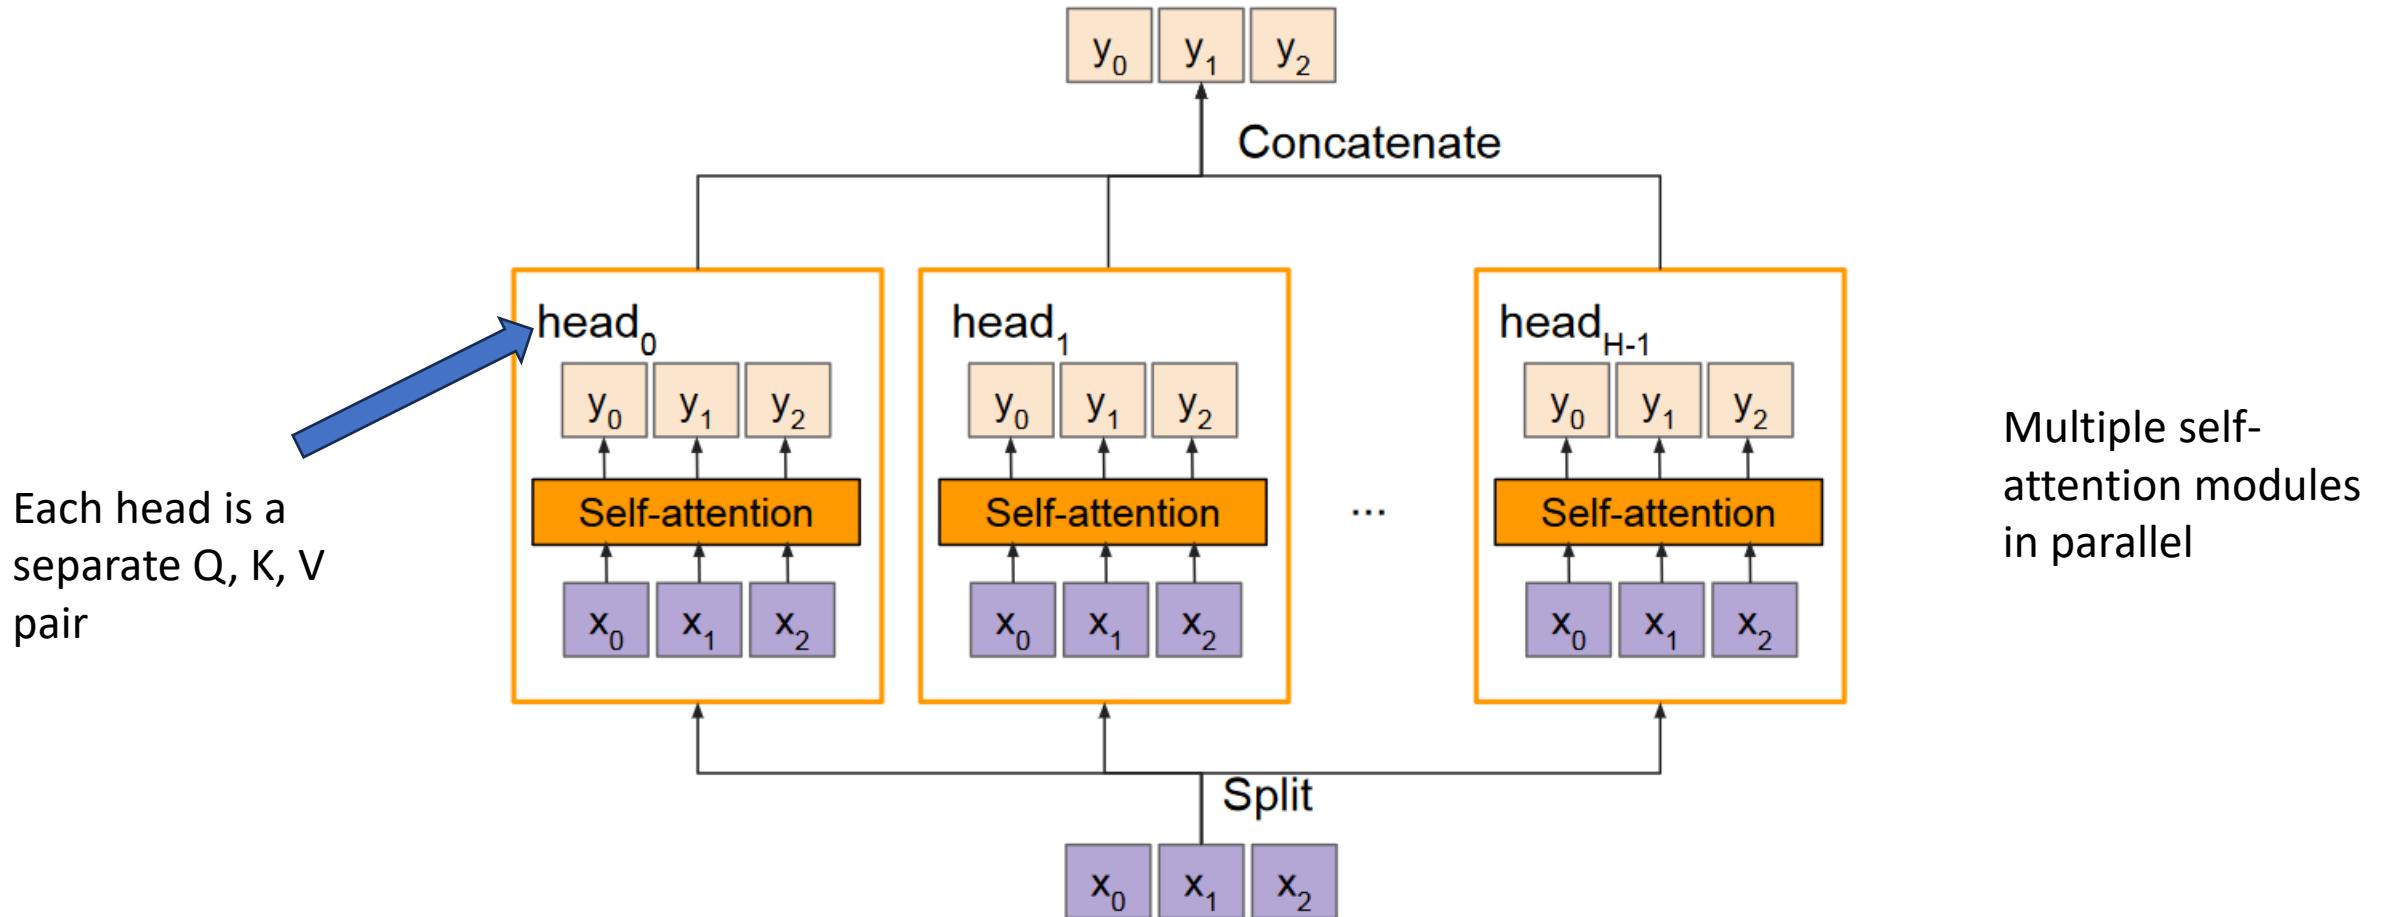
Key  
(Matching)

Value  
(Answer)

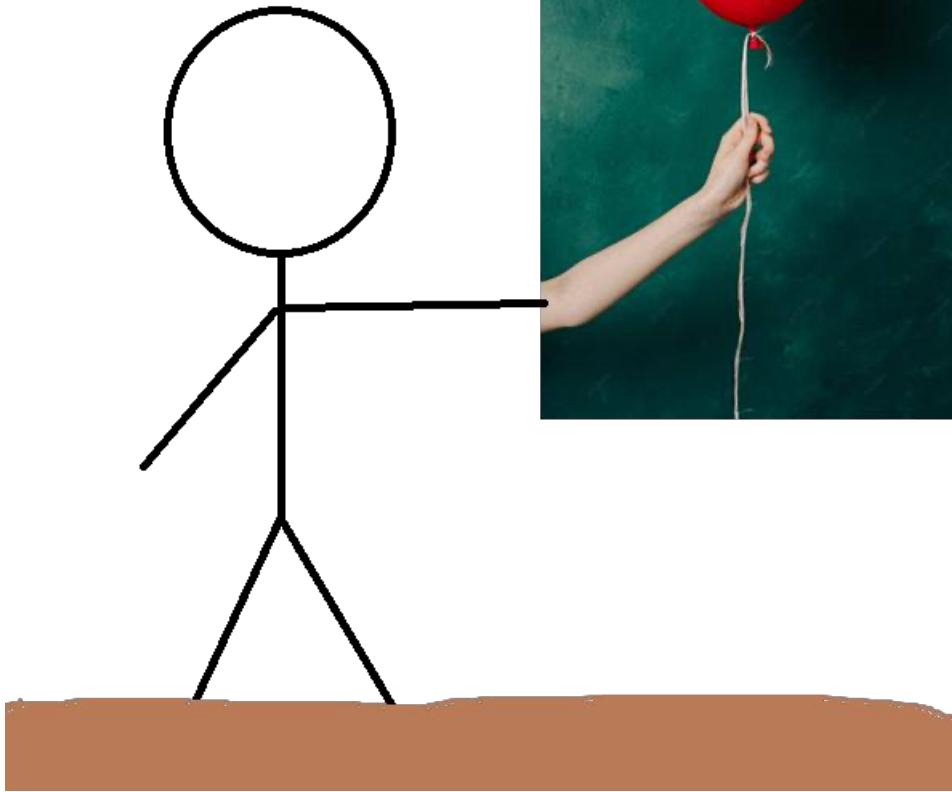
# Attention (Single-head)



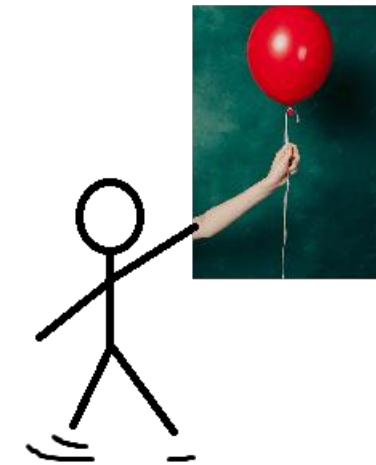
# Multi-head attention



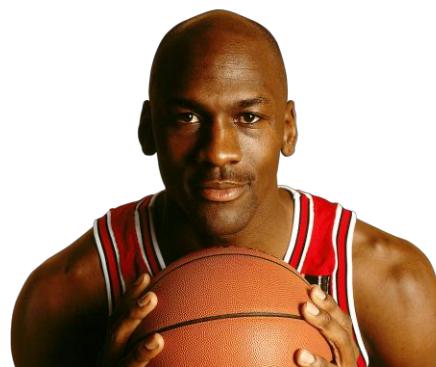
# Heads



Different  
pictures describe  
different  
scenarios for the  
same object



# Heads (Michael, depending on the context)



# nn.MultiheadAttention module

Multihead attention is defined as:

$$\text{Multihead}(Q, K, V) = \text{Concat}(\text{head}_1, \dots, \text{head}_h)W^O$$

where  $\text{head}_i = \text{Attention}(QW_i^Q, KW_i^K, VW_i^V)$

W's are learnable parameters:

$$W_{1\dots h}^Q \in \mathbb{R}^{D \times d_k}, W_{1\dots h}^K \in \mathbb{R}^{D \times d_k}, W_{1\dots h}^V \in \mathbb{R}^{D \times d_v}, \text{ and } W^O \in \mathbb{R}^{h \cdot d_k \times d_{out}}$$

Attention function is defined as:

$$\text{Attention}(Q, K, V) = \text{softmax}\left(\frac{QK^T}{\sqrt{d_k}}\right)V \quad \textit{Scaling factor to resolve curse of dimensionality}$$

where

$$\text{queries } Q \in \mathbb{R}^{T \times d_k}, \text{keys } K \in \mathbb{R}^{T \times d_k} \text{ and values } V \in \mathbb{R}^{T \times d_v}$$

# A simple implementation using PyTorch modules

---

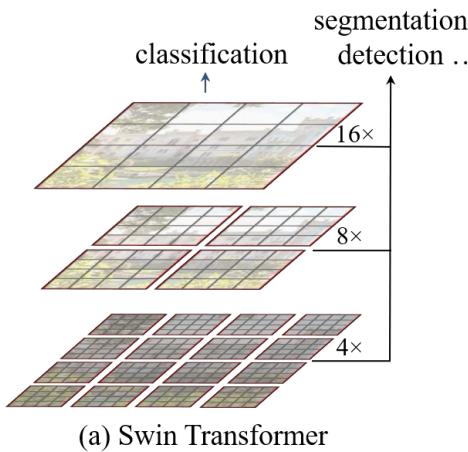
```
import torch, torch.nn as nn

class PatchEmbed(nn.Module):
    def __init__(self, in_ch=3, dim=128, patch=4):
        super().__init__()
        self.proj = nn.Conv2d(in_ch, dim, patch, patch)
    def forward(self, x):
        return self.proj(x).flatten(2).transpose(1,2) # [B, N, D]

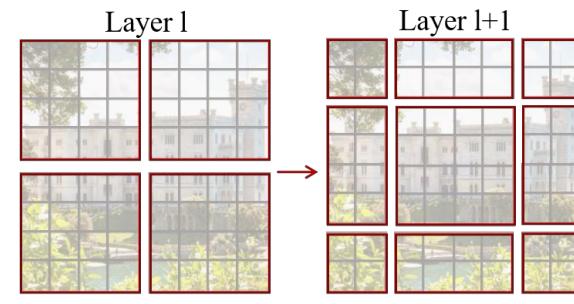
class ViT(nn.Module):
    def __init__(self, img=32, patch=4, dim=128, heads=4, depth=4, n_cls=10):
        super().__init__()
        self.patch_embed = PatchEmbed(3, dim, patch)
        n_p = (img//patch)**2
        self.cls = nn.Parameter(torch.zeros(1,1, dim))
        self.pos = nn.Parameter(torch.randn(1, n_p+1, dim))
        enc = nn.TransformerEncoderLayer(dim, heads, dim*4, batch_first=True)
        self.enc = nn.TransformerEncoder(enc, depth)
        self.head = nn.Linear(dim, n_cls)
    def forward(self, x):
        B = x.size(0)
        x = self.patch_embed(x)
        x = torch.cat([self.cls.expand(B,-1,-1), x], 1) + self.pos
        return self.head(self.enc(x)[:,0])
```



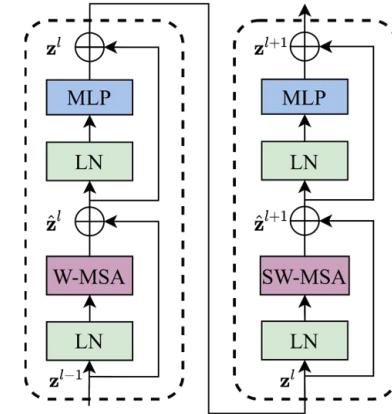
# SOTA architecture: Swin transformers



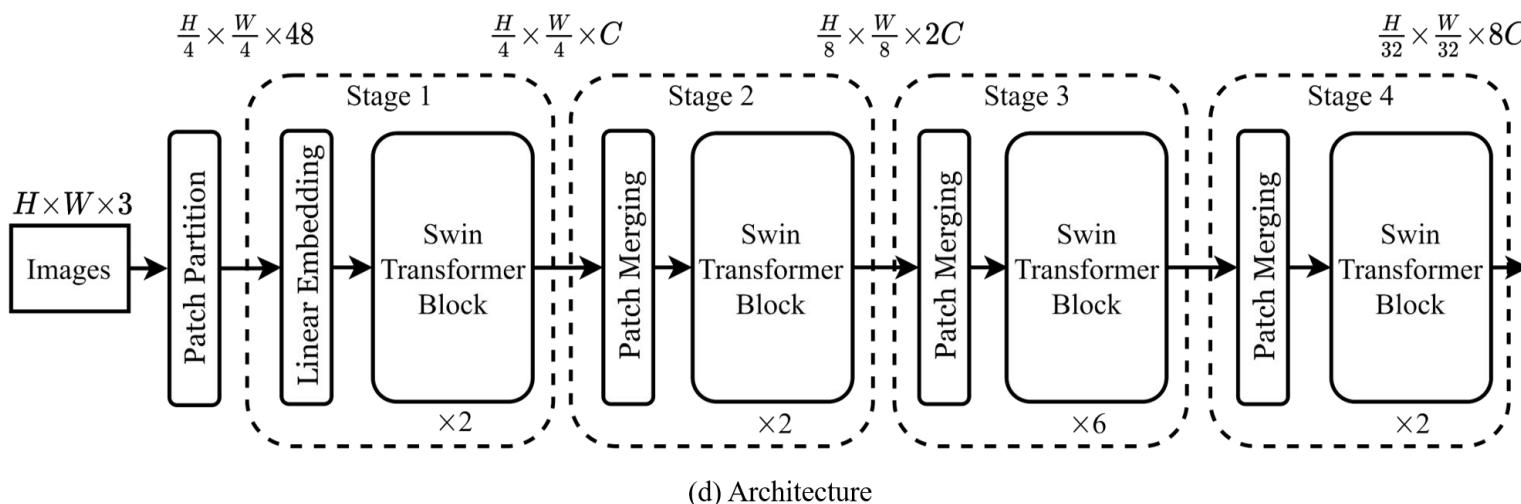
(a) Swin Transformer



(b) Shifted Window

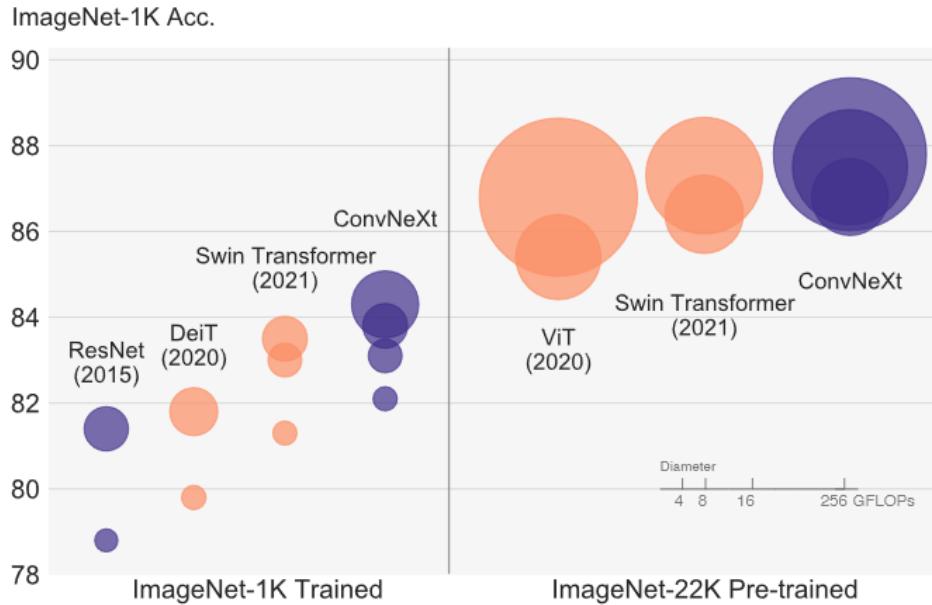


(c) Two Successive Swin Transformer Blocks



(d) Architecture

# Revisiting Convnet: ConvNeXt



**Figure 1. ImageNet-1K classification** results for • ConvNets and ○ vision Transformers. Each bubble's area is proportional to FLOPs of a variant in a model family. ImageNet-1K/22K models here take  $224^2/384^2$  images respectively. ResNet and ViT results were obtained with improved training procedures over the original papers. We demonstrate that a standard ConvNet model can achieve the same level of scalability as hierarchical vision Transformers while being much simpler in design.

## Pure conv net:

- Very few non-linear functions
- LN instead of BN
- 7x7 conv
- Gelu instead of Relu

<https://arxiv.org/pdf/2201.03545.pdf>

# And the revenge of ViT!

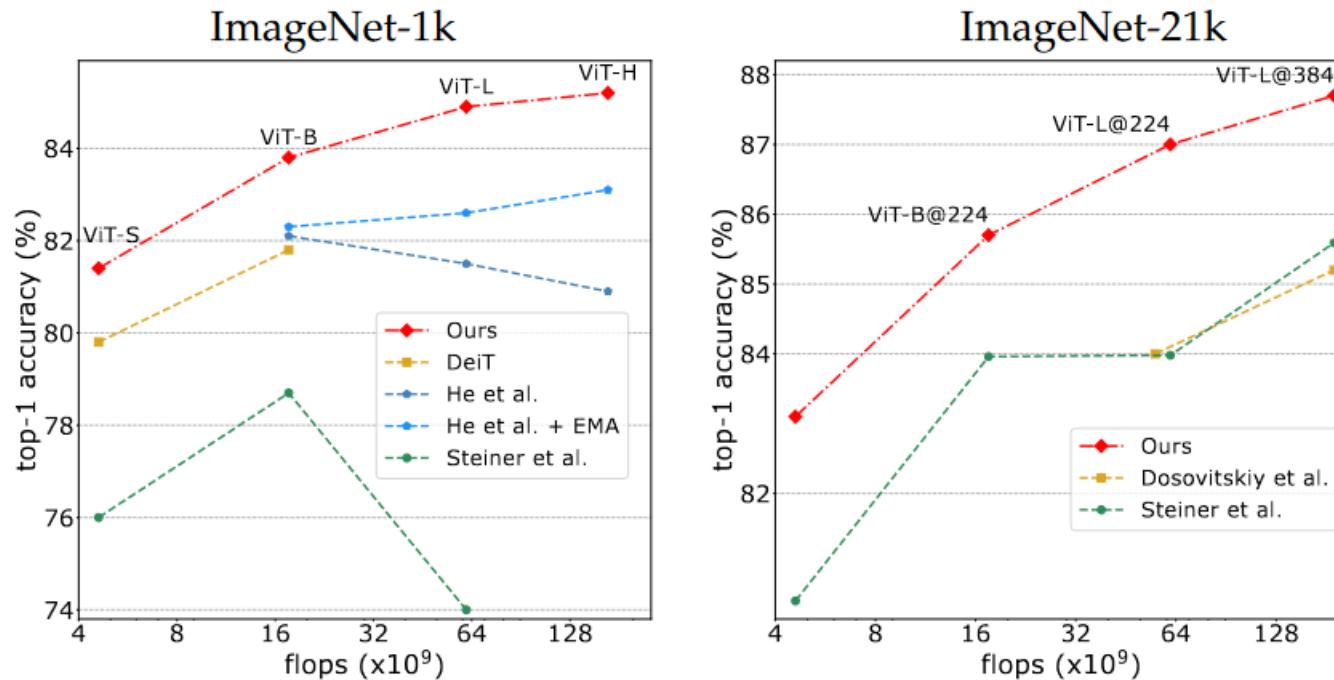


Figure 1: Comparison of training recipes for (left) vanilla vision transformers trained on ImageNet-1k and evaluated at resolution  $224 \times 224$ , and (right) pre-trained on ImageNet-21k at  $224 \times 224$  and fine-tuned on ImageNet-1k at resolution  $224 \times 224$  or  $384 \times 384$ .

# Image Captioning

# Flickr8k



A child in a pink dress is climbing up a set of stairs in an entry way .

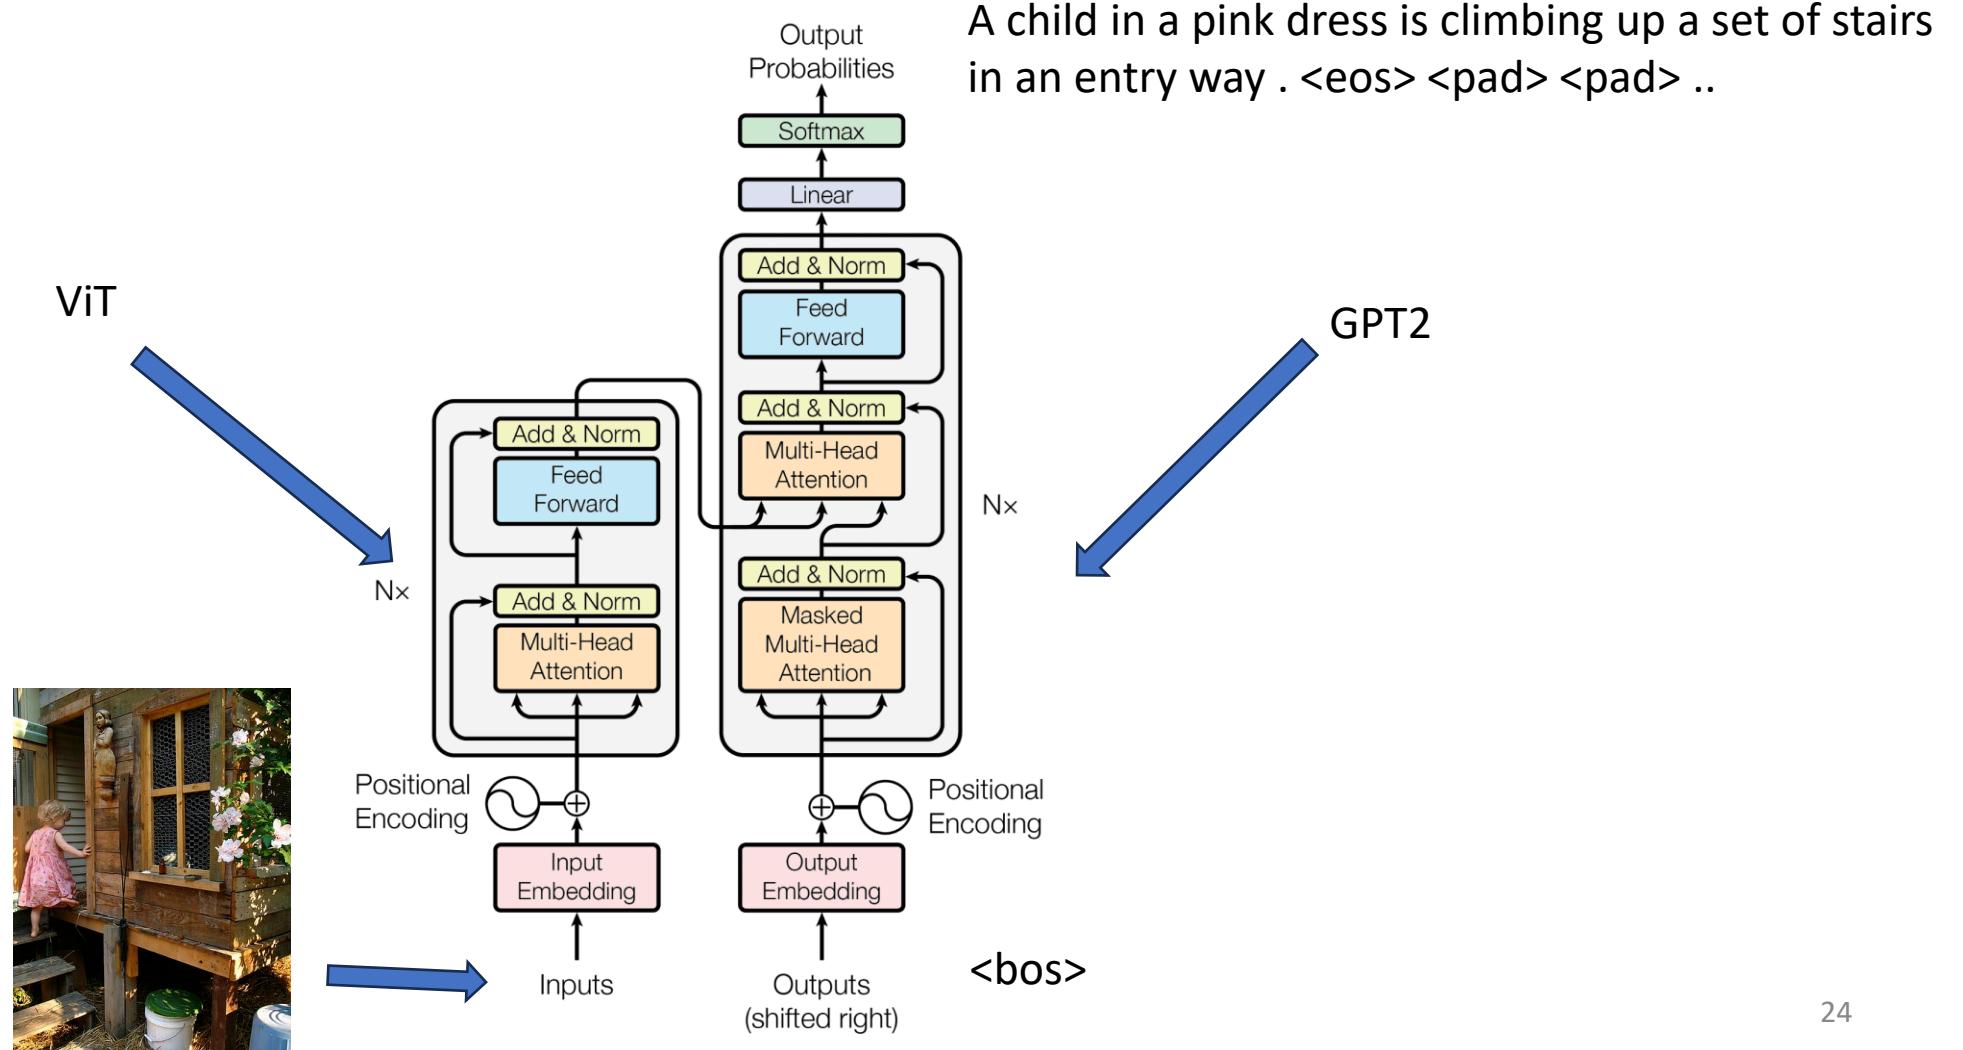
A girl going into a wooden building .

A little girl climbing into a wooden playhouse .

A little girl climbing the stairs to her playhouse .

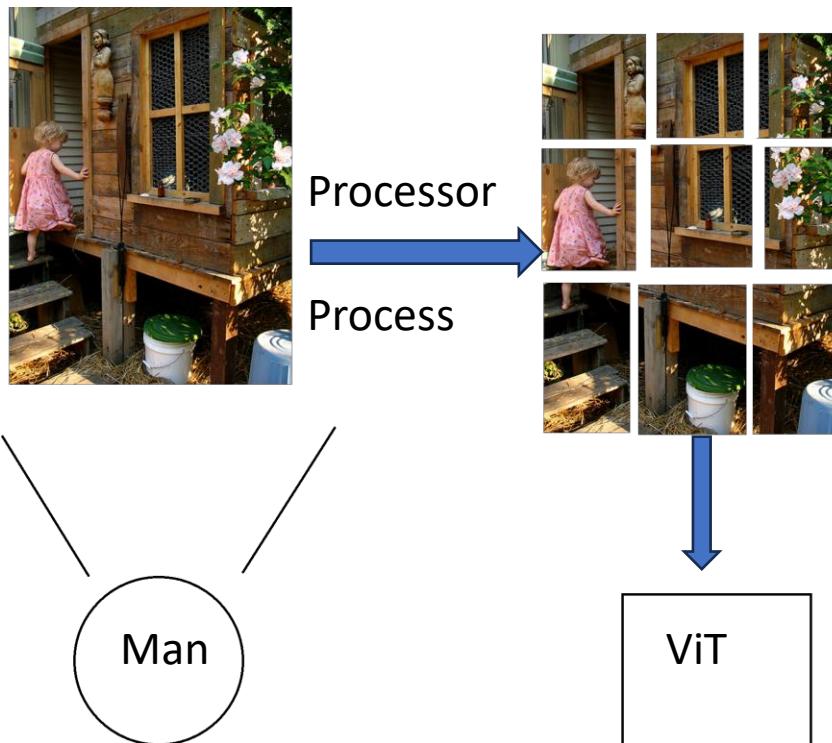
A little girl in a pink dress going into a wooden cabin .

# Transformer Encoder-Decoder (Generation)

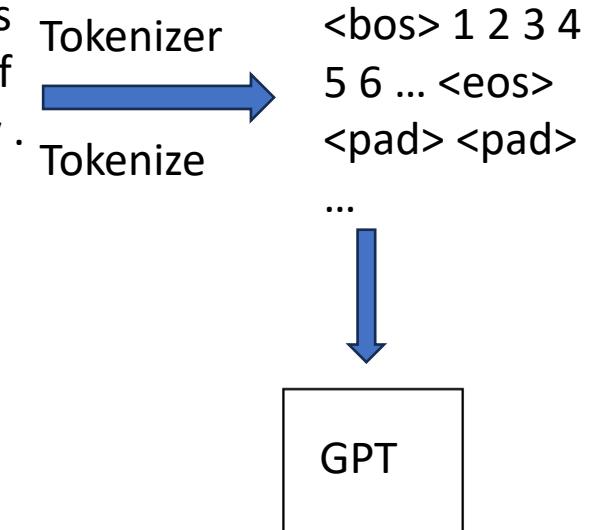


# Tokens

- ViT reads image in batch, GPT reads words (and punctuations) in tokens



A child in a pink dress  
is climbing up a set of  
stairs in an entry way .



# Tokens

- Usually positive numbers
- <cls>
- <pad> : Pad to unify the size of (maybe truncated) input output.  
Usually, this token should not be learned by the model. So, during training, this should be replaced with a random token that does not exist in the vocab.
- <eos>
- <sep>
- <bos>
- 1 sentence is around 30-40 tokens (English)

# Masking

- [Attention in transformers, visually explained | Chapter 6, Deep Learning \(youtube.com\)](#)

# Seq2SeqTrainer

- [Trainer \(huggingface.co\)](#)
- [google/vit-base-patch16-224 · Hugging Face](#)
- [openai-community/gpt2 · Hugging Face](#)
- [OpenAI GPT2 \(huggingface.co\)](#)
- [Vision Encoder Decoder Models \(huggingface.co\)](#)