

Autoencoders and Image Generation

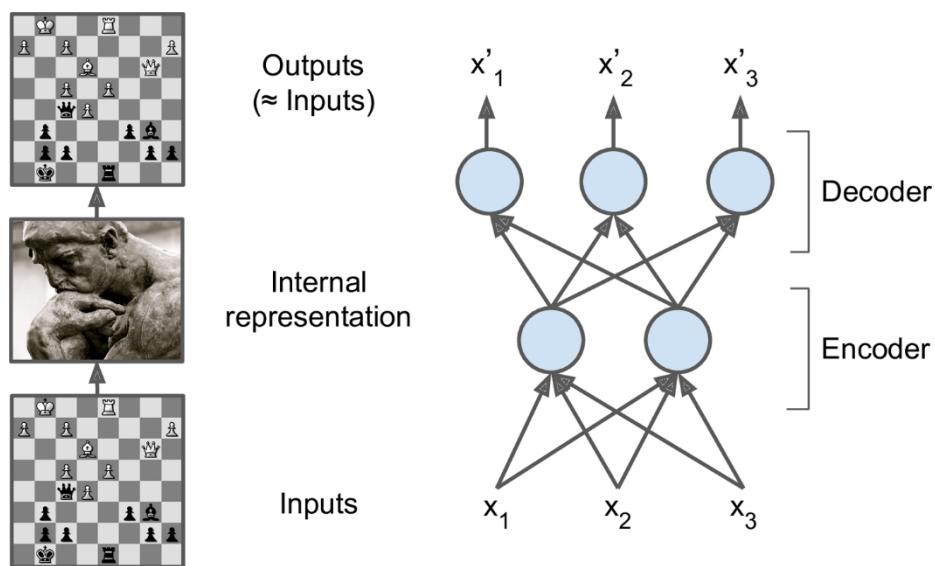
CMPUT 328

Nilanjan Ray

Picture source: "Hands on Machine Learning with Scikit-Learn & TensorFlow" by Aurelien Geron


Part I: Autoencoders

Representation learning by autoencoders




Can we learn interesting, hidden representations from **unlabeled** data?

Extract underlying (low) dimensionality



Even though the data points are 3D, they more or less lie on a 2D plane.

An example autoencoder for MNIST

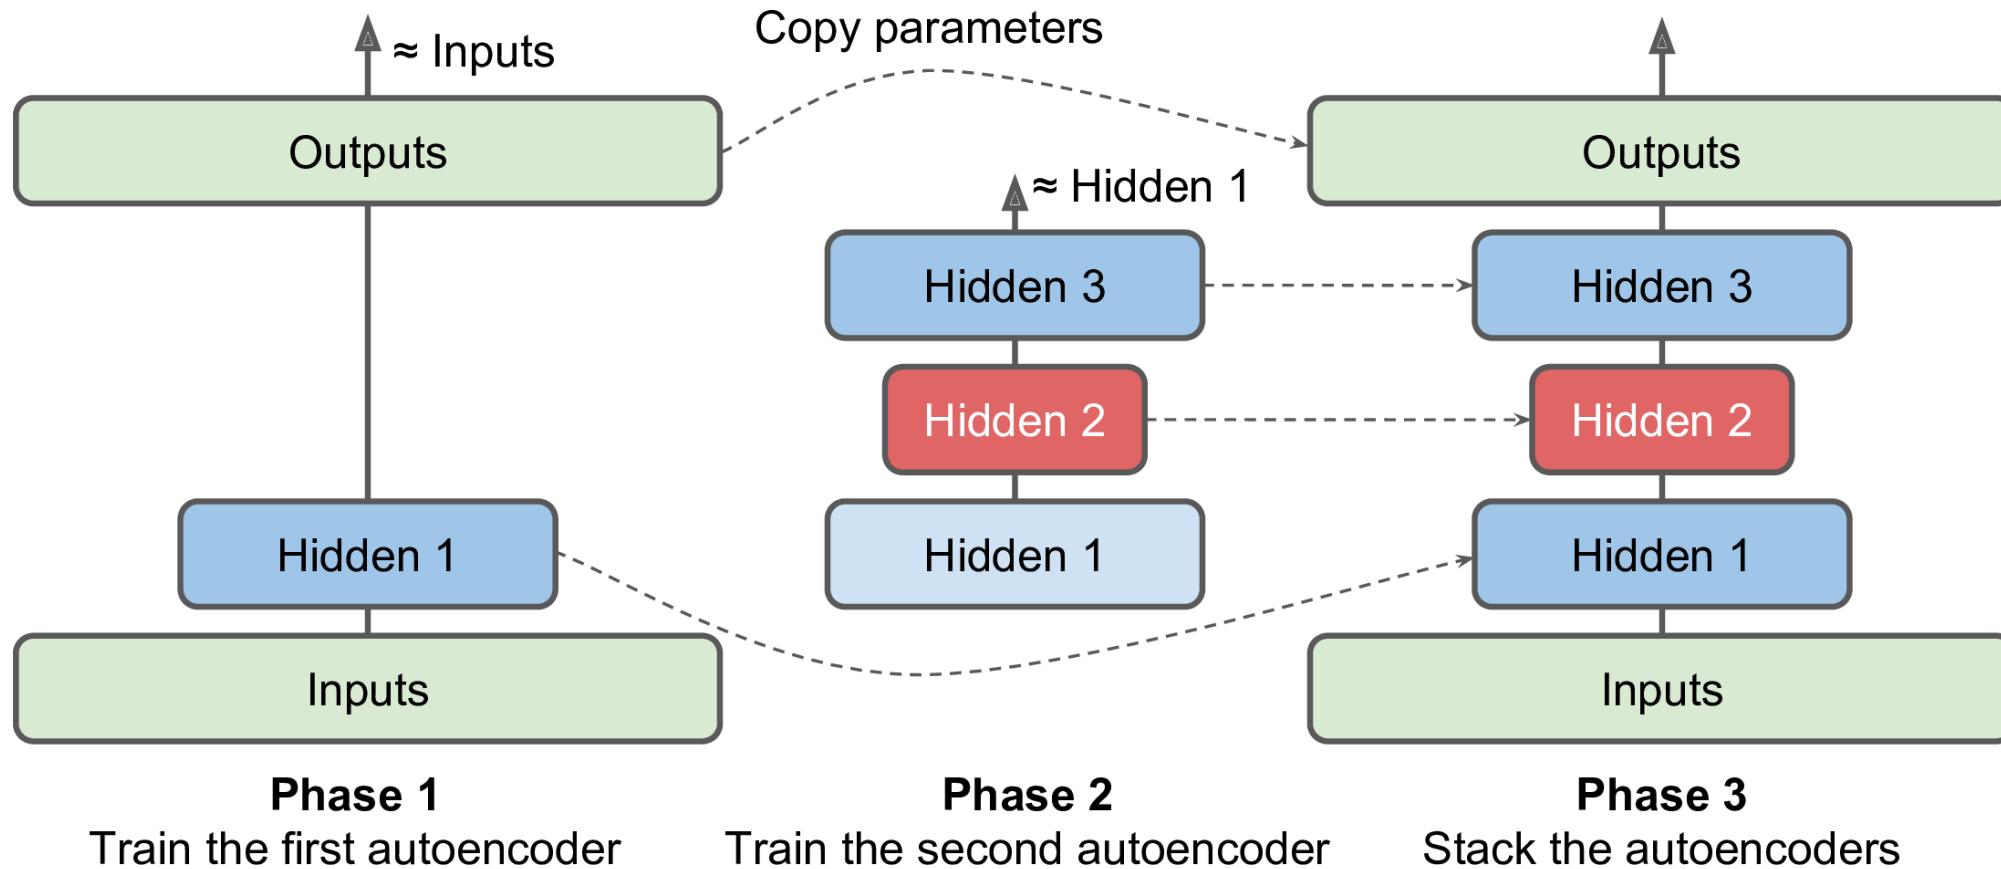


MNIST_AE.ipnyb implements this architecture


Convolutional AE

- We can easily replace fully connected layers by all convolutional layers with proper padding
- Look at MNIST_AE.ipnyb

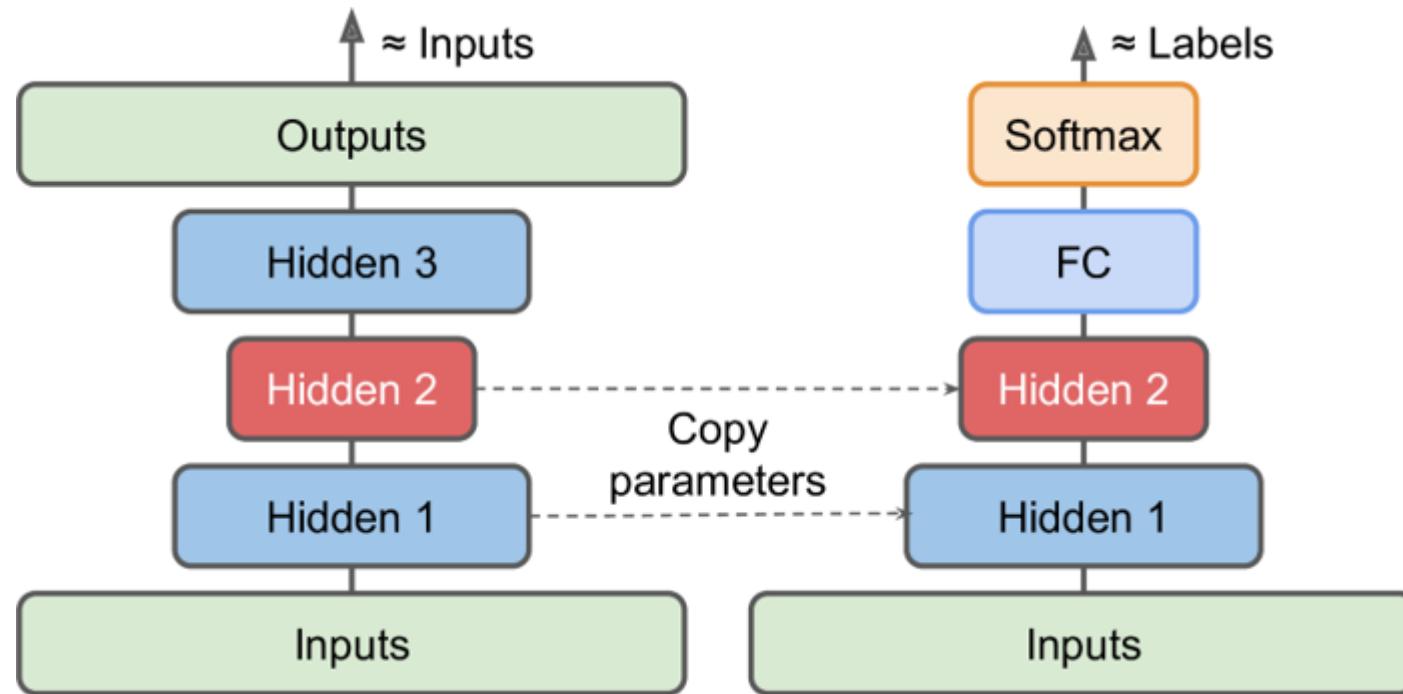
Phased training in AE



Classification from AE




Semi-supervised learning



Often, labeled data is partially available: maybe 80% training data is unlabeled, only 20% is labeled. How do we make use of unlabeled data during training a classifier?


Denoising AE



Add noise or insert a dropout layer


Part II: Image Generation

Generating images by neural networks



This is a typical setup: generate images with random noise inputs

Types of image generators





Picture: <https://lilianweng.github.io/posts/2021-07-11-diffusion-models/>

Probability primer

- We need to refresh our understanding of probability and random variables
- Gaussian probability density function
 - Scalar
 - Vector
- Generating random variables from a pdf
- Conditional density function


Probability mass function (pmf)

A familiar example: a grayscale image histogram



Random variables: Pixel values


PMF = mass attached
to random variables



Sampling random variables from a pmf

- There are various methods
- Here's an example method:
- Generate random variable from a **familiar distribution**, such as a uniform distribution or a Gaussian distribution– these in-built functions are available to us
- Then transform this r.v. so that the transformed r.v. follows our target pmf
 - How?
- At the end of the day, it will become a function call (in numpy or pytorch)

Transformation of r. v. visually




Uniform histogram

Cameraman histogram

Uniformly distributed r.v. (uniformly spaced points) are mapped (aka **transformed**) to r.v. that follows cameraman pmf

The original space is being **stretched and compressed** to mimic the mass distribution in cameraman pmf.


Gaussian density function



A Gaussian r.v. with mean m and Standard deviation s , is generated as:

$g = m + s * \text{torch.randn}(...)$

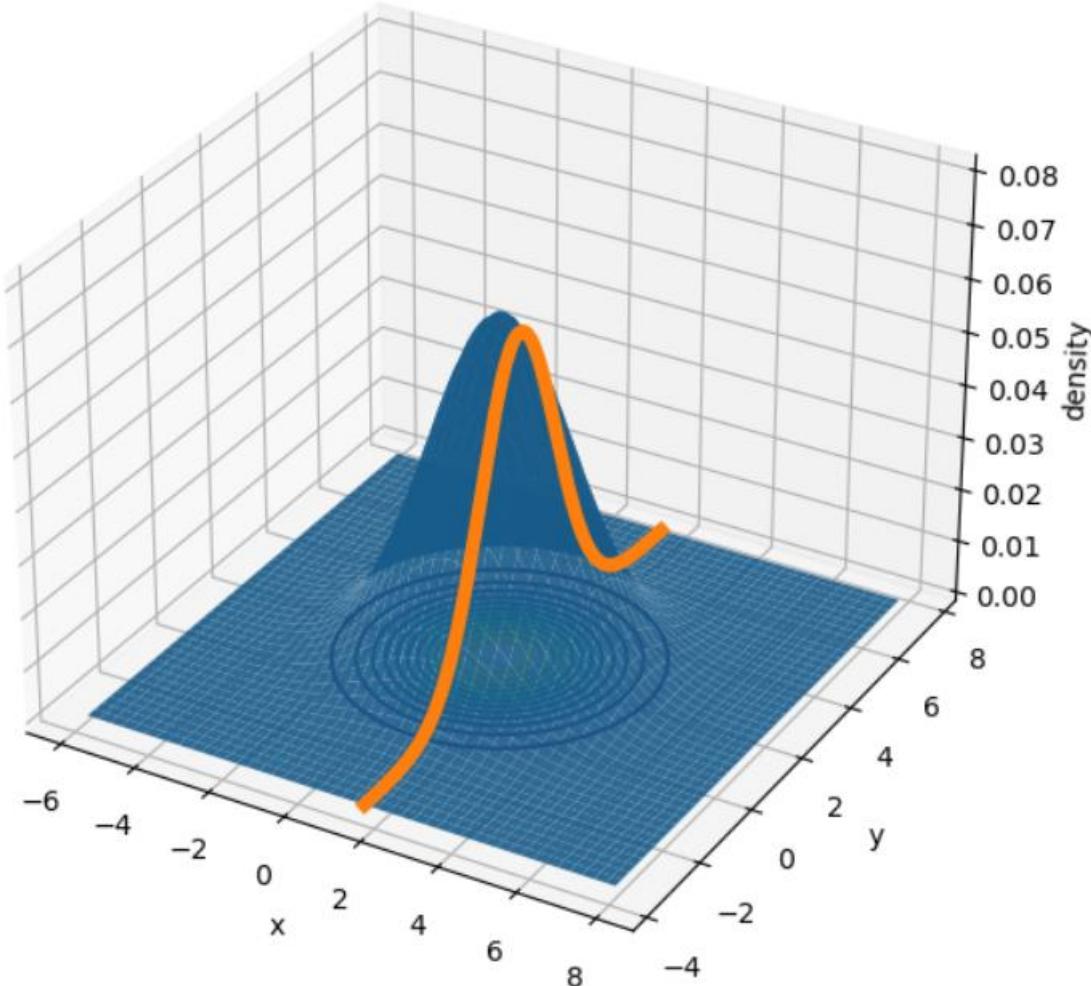
2D Gaussian PDF



One more
element:
Conditional
density


2D Gaussian Surface ($\mu=(1,2)$, $\Sigma=\text{diag}(3,2)$) with Conditional at $x = 1.5$

■ $p(y | x = 1.5)$ on surface

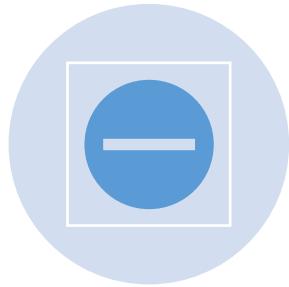


Conditions
can be
complicated...

$$p(x, y \mid \text{condition}) = \frac{p(x, y)}{p(\text{condition})},$$



Just an aside: Can we generate the cameraman picture from its histogram?




No



We lost all spatial organization of pixels in the histogram

Variational autoencoder: Probabilistic encoder and decoder



Variational autoencoder...


We want: $q_\phi(z|x) \approx p_\theta(z|x)$

Probability of latent code
from the **encoder** side
given image x

Probability of latent code
from the **decoder** side
given image x

After a ton of math (<https://lilianweng.github.io/posts/2018-08-12-vae/>), our loss function:

Image Reconstruction + KL-divergence
between $q_\phi(z|x)$ and $N(0,1)$



KL-divergence: a type of distance between density functions

If you imagine encoding outcomes optimized for Q , but nature draws from P : $D_{\text{KL}}(P \parallel Q)$ tells you the **expected number of extra bits (or nats)** per sample you waste.

Definition

$$D_{\text{KL}}(P \parallel Q) = \sum_i P(i) \ln \frac{P(i)}{Q(i)}$$

Example

Outcome	$P(i)$	$Q(i)$	$P(i) \ln \frac{P(i)}{Q(i)}$
a	0.5	0.4	0.1116
b	0.3	0.4	-0.0863
c	0.2	0.2	0
Sum			0.0253 nats (≈ 0.036 bits)

Interpretation

- $D_{\text{KL}}(P \parallel Q) \geq 0$, zero only if $P = Q$.
- Asymmetric: $D_{\text{KL}}(P \parallel Q) \neq D_{\text{KL}}(Q \parallel P)$.
- Measures *information loss* when approximating P by Q .

KL-divergence between two univariate Gaussians

KL divergence between two univariate Gaussians

Let

$$P = \mathcal{N}(\mu_p, \sigma_p^2), \quad Q = \mathcal{N}(\mu_q, \sigma_q^2).$$

Then the KL divergence $D_{\text{KL}}(P\|Q)$ is:

$$D_{\text{KL}}(P\|Q) = \log \frac{\sigma_q}{\sigma_p} + \frac{\sigma_p^2 + (\mu_p - \mu_q)^2}{2\sigma_q^2} - \frac{1}{2}.$$

Example

Let

$$\mu_p = 0, \sigma_p = 1, \mu_q = 1, \sigma_q = 2.$$

$$\begin{aligned} D_{\text{KL}}(P\|Q) &= \log \frac{2}{1} + \frac{1^2 + (0-1)^2}{2(2^2)} - \frac{1}{2} \\ &= \log 2 + \frac{1+1}{8} - 0.5 = 0.6931 + 0.25 - 0.5 = 0.4431. \end{aligned}$$

So $D_{\text{KL}}(P\|Q) \approx 0.443$ nats (≈ 0.64 bits).

KL-divergence term in VAE

Setup:

$$q_\phi(z|x) = \mathcal{N}(z; \mu_\phi(x), \text{diag}(\sigma_\phi(x)^2)),$$
$$p(z) = \mathcal{N}(0, I)$$

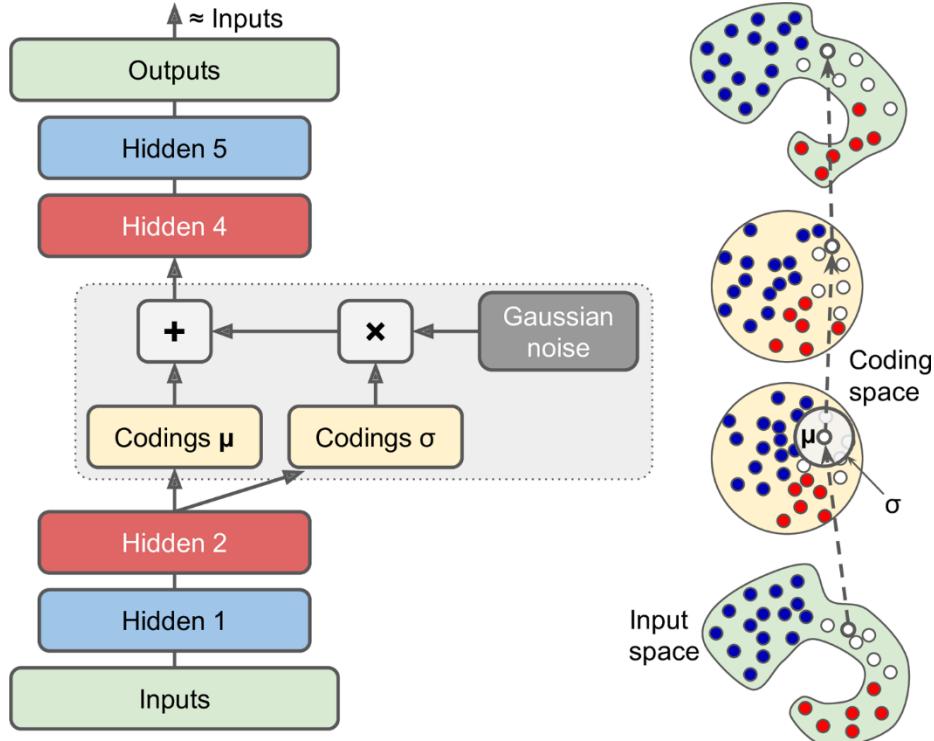
KL-divergence term:

$$D_{KL}(q_\phi(z|x) \mid\mid p(z)) = \frac{1}{2} \sum_i (\mu_i^2 + \sigma_i^2 - \log \sigma_i^2 - 1)$$

PyTorch implementation:

```
kl = -0.5 * torch.sum(1 + logvar - mu.pow(2) - logvar.exp())
```

Variational autoencoder...



Cost function has two components:

Image Reconstruction +
KL-divergence (constraints on μ and σ)

The constraint: the encoded distribution should look like a zero-mean, unit variance Gaussian.

<https://lilianweng.github.io/posts/2018-08-12-vae/>

<https://www.youtube.com/watch?v=OM95kDPAW-M>

MNIST_Variational_AE.ipynb

Issues with VAE

1. Blurry Reconstructions

- VAEs average over many possible versions of an image.
 - This produces smooth, washed-out results instead of sharp details.
-

2. Posterior Collapse

- The encoder sometimes stops using the input image.
 - All latent codes become nearly identical, so the model ignores the latent space entirely.
-

3. Reconstruction vs. Regularization Trade-off

- Balancing image quality and latent structure is difficult.
- Too much regularization = poor reconstructions;
too little = unstable or meaningless latent space.

4. Overly Smooth Latent Space

- The continuous Gaussian latent space blends different concepts together.
 - Interpolations between unrelated images create unrealistic "hybrid" results.
-

5. Hard to Combine with Discrete Generative Models

- PixelCNNs and Transformers work best with discrete symbols.
 - VAEs use continuous latents, so they can't take advantage of these powerful priors.
-

6. Loss of Fine Details

- The decoder models pixels independently, ignoring local texture patterns.
- This leads to loss of crisp edges, fine textures, and high-frequency detail.

VQ (Vector Quantized)-VAE

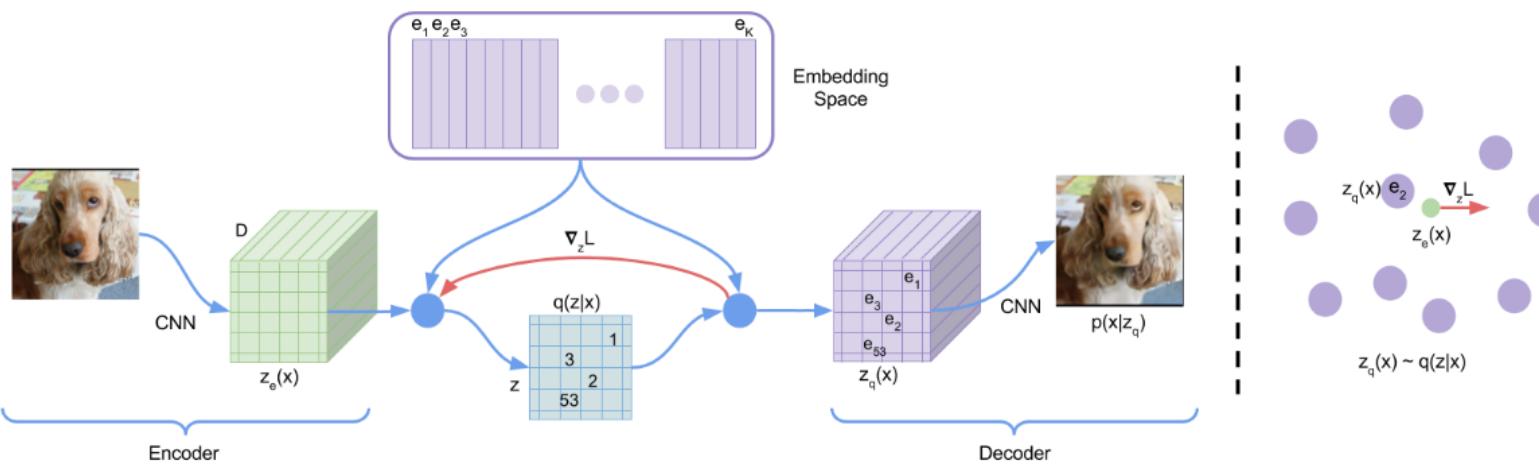


Figure 1: Left: A figure describing the VQ-VAE. Right: Visualisation of the embedding space. The output of the encoder $z(x)$ is mapped to the nearest point e_2 . The gradient $\nabla_z L$ (in red) will push the encoder to change its output, which could alter the configuration in the next forward pass.

<https://arxiv.org/pdf/1711.00937>

Vector quantization: toy example

Vector Quantization – Simple Example

Codebook (learned prototypes)

$$e_0 = (1.0, 1.0) \quad e_1 = (3.0, 1.0) \quad e_2 = (2.0, 3.0)$$

Encoder outputs (continuous latents)

$$A = (0.8, 1.2) \quad B = (2.9, 0.8) \quad C = (2.1, 2.7)$$

Input	Nearest code	Quantized vector z_q
A	e_0	(1.0, 1.0)
B	e_1	(3.0, 1.0)
C	e_2	(2.0, 3.0)

→ Each continuous latent z_q is “snapped” to its nearest codebook vector.

VQ-VAE...

1. The codebook

- Shape: $(K \times D)$
 - K = number of code entries (e.g., 256)
 - D = dimension of each embedding vector (e.g., 64)

Codebook ($K=4$, $D=2$):

Index	Embedding vector
0	[1.2, -0.7]
1	[-0.3, 0.9]
2	[2.1, 1.5]
3	[0.0, -1.1]

So each row in the codebook is a **prototype vector** of length D .

2. The encoder output

When the encoder processes an image, it doesn't give *just one* D -dimensional vector — it gives a **spatial grid** of them.

- Shape: $(B \times D \times H \times W)$
 - B = batch size
 - D = feature dimension (same as codebook embedding size)
 - $H \times W$ = latent grid size (e.g., 7×7 for 28×28 MNIST)

So each position in that grid (each of the $H \times W$ locations) is a **D -dimensional vector**.

That's the one that gets quantized.

3. Quantization step

For every one of those $H \times W$ vectors:

1. Compute distance to all K codebook vectors (each D -dim).
2. Pick the nearest one.
3. Replace the encoder vector with that codebook embedding.

Result: a quantized latent map

- shape $(B \times D \times H \times W)$ again (same shape as input)
- but with values *snapped* to codebook vectors.

4. Discrete view

You can also represent the quantized grid as:

- **indices** of shape $(B \times H \times W)$ — each element $\in [0, K-1]$
(which codebook row was chosen for that pixel in latent space)

VQ-VAE Image generation

Stage 1 — Train VQ-VAE

- Train an encoder–quantizer–decoder model on many images.
- The encoder converts each image to a latent grid of vectors (size $H \times W \times D$).
- Each latent vector is replaced by its nearest **codebook embedding** (from K entries).

Stage 2 — Build the Discrete Dataset

- After training, encode all images.
- Each image is now represented by an **$H \times W$ grid of integers** in $[0, K-1]$.
- These integer grids form a **new dataset** in discrete latent space.

Stage 3 — Train a Prior (PixelRNN / PixelCNN / Transformer)

- Train a generative model to **predict code indices** pixel-by-pixel (or token-by-token).
- The prior learns the **distribution of valid index patterns** that correspond to real images.

Stage 4 — Image Generation

1. Use the trained prior to **sample a new $H \times W$ grid of indices**.
2. Look up the corresponding **codebook embeddings** → get a tensor of shape $D \times H \times W$.
3. Feed this tensor to the **VQ-VAE decoder** to reconstruct the final image.

✓ Key Idea

A VQ-VAE turns images into discrete latent “tokens.”
A PixelRNN (or Transformer) learns to generate those tokens.
The decoder turns the generated tokens back into a realistic image.

Why VQ-VAE Works Better Than a VAE

1. Continuous vs. Discrete Latents

- A standard VAE uses a **continuous Gaussian latent space**.
 - Pros: smooth interpolation
 - Cons: tends to produce **blurry reconstructions** because it averages over many possible latent codes.
- A VQ-VAE uses **discrete codebook entries** (nearest prototypes).
 - Enforces *crisp decisions* → sharper reconstructions.
 - Reduces posterior variance — no sampling noise from Gaussians.

2. Sharper Reconstructions

- VAEs reconstruct images from *probabilistic* latents → decoder must handle uncertainty → blurred results.
- VQ-VAEs replace that uncertainty with **hard assignments** → decoder sees *clean, consistent codes* → produces sharper, more detailed images.

3. Discrete “Vocabulary” Enables Language-like Modeling

- The codebook gives a **discrete latent vocabulary** (like visual words).
- This allows powerful **autoregressive models** (PixelCNN, Transformer) to model image structure effectively — something that's hard in a continuous latent space.

4. No KL-collapse problem

- VAEs often suffer from *posterior collapse* (the encoder ignores the latent).
- VQ-VAE sidesteps this entirely — the quantization and commitment losses naturally keep encoder and codebook aligned.

5. Better compression & interpretability

- Discrete latents are easier to store, transmit, and interpret.
- The learned codebook can be seen as a “palette” of visual patterns or texture prototypes.